Cisco Nexus 3132Q-V

Cluster and storage switches

NetApp
May 08, 2024

Table of Contents

Cisco Nexus 3132Q-V

Overview

Overview of installation and configuration for Cisco Nexus 3132Q-V switches

Cisco Nexus 3132Q-V switches can be used as cluster switches in your AFF or FAS cluster. Cluster switches allow you to build ONTAP clusters with more than two nodes.

Initial configuration overview

To initially configure a Cisco Nexus 3132Q-V switch on systems running ONTAP, follow these steps:

1. Complete Cisco Nexus 3132Q-V cabling worksheet. The sample cabling worksheet provides examples of recommended port assignments from the switches to the controllers. The blank worksheet provides a template that you can use in setting up your cluster.
2. Install a Cisco Nexus 3132Q-V cluster switch in a NetApp cabinet. install the Cisco Nexus 3132Q-V switch and pass-through panel in a NetApp cabinet with the standard brackets that are included with the switch.
3. Configure the Cisco Nexus 3132Q-V switch. Set up and configure the Cisco Nexus 3132Q-V switch.
4. Prepare to install NX-OS software and Reference Configuration File. Prepare to install the NX-OS software and the Reference Configuration File (RCF).
5. Install the NX-OS software. Follow this procedure to install the NX-OS software on the Nexus 3132Q-V cluster switch.
6. Install the Reference Configuration File (RCF). Follow this procedure to install the RCF after setting up the Nexus 3132Q-V switch for the first time. You can also use this procedure to upgrade your RCF version.

Additional information

Before you begin installation or maintenance, be sure to review the following:

- Configuration requirements
- Required documentation
- Smart Call Home requirements

Configuration requirements for Cisco Nexus 3132Q-V switches

For Cisco Nexus 3132Q-V switch installation and maintenance, be sure to review network and configuration requirements.

Configuration requirements

To configure your cluster, you need the appropriate number and type of cables and cable connectors for your switches. Depending on the type of switch you are initially configuring, you need to connect to the switch console port with the included console cable; you also need to provide specific network information.

Network requirements

You need the following network information for all switch configurations:

- IP subnet for management network traffic.

- Host names and IP addresses for each of the storage system controllers and all applicable switches.

- Most storage system controllers are managed through the e0M interface by connecting to the Ethernet service port (wrench icon). On AFF A800 and AFF A700 systems, the e0M interface uses a dedicated Ethernet port.

Refer to the Hardware Universe for latest information.

Documentation requirements for Cisco Nexus 3132Q-V switches

For Cisco Nexus 3132Q-V switch installation and maintenance, be sure to review all the recommended documentation.

Switch documentation

To set up the Cisco Nexus 3132Q-V switches, you need the following documentation from the Cisco Nexus 3000 Series Switches Support page.

Document title	Description
Nexus 3000 Series Hardware Installation Guide	Provides detailed information about site requirements, switch hardware details, and installation options.
Cisco Nexus 3000 Series Switch Software Configuration Guides (choose the guide for the NX-OS release installed on your switches)	Provides initial switch configuration information that you need before you can configure the switch for ONTAP operation.
Cisco Nexus 3000 Series NX-OS Software Upgrade and Downgrade Guide (choose the guide for the NX-OS release installed on your switches)	Provides information on how to downgrade the switch to ONTAP supported switch software, if necessary.
Cisco Nexus 3000 Series NX-OS Command Reference Master Index	Provides links to the various command references provided by Cisco.
Cisco Nexus 3000 MIBs Reference	Describes the Management Information Base (MIB) files for the Nexus 3000 switches.
Nexus 3000 Series NX-OS System Message Reference	Describes the system messages for Cisco Nexus 3000 series switches, those that are informational, and others that might help diagnose problems with links, internal hardware, or the system software.
Cisco Nexus 3000 Series NX-OS Release Notes (choose the notes for the NX-OS release installed on your switches)	Describes the features, bugs, and limitations for the Cisco Nexus 3000 Series.

Document title	Description
Regulatory, Compliance, and Safety Information for the Cisco Nexus 6000, Cisco Nexus 5000 Series, Cisco Nexus 3000 Series, and Cisco Nexus 2000 Series	Provides international agency compliance, safety, and statutory information for the Nexus 3000 series switches.

ONTAP systems documentation

To set up an ONTAP system, you need the following documents for your version of the operating system from the ONTAP 9 Documentation Center.

Name	Description
Controller-specific *Installation and Setup Instructions*	Describes how to install NetApp hardware.
ONTAP documentation	Provides detailed information about all aspects of the ONTAP releases.
Hardware Universe	Provides NetApp hardware configuration and compatibility information.

Rail kit and cabinet documentation

To install a 3132Q-V Cisco switch in a NetApp cabinet, see the following hardware documentation.

Name	Description
42U System Cabinet, Deep Guide	Describes the FRUs associated with the 42U system cabinet, and provides maintenance and FRU replacement instructions.
Install Cisco Nexus 3132Q-V switch in a NetApp Cabinet	Describes how to install a Cisco Nexus 3132Q-V switch in a four-post NetApp cabinet.

Smart Call Home requirements

To use Smart Call Home feature, review the following guidelines.

Smart Call Home monitors the hardware and software components on your network. When a critical system configuration occurs, it generates an email-based notification and raises an alert to all the recipients that are configured in your destination profile. To use Smart Call Home, you must configure a cluster network switch to communicate using email with the Smart Call Home system. In addition, you can optionally set up your cluster network switch to take advantage of Cisco's embedded Smart Call Home support feature.

Before you can use Smart Call Home, be aware of the following considerations:

- An email server must be in place.
- The switch must have IP connectivity to the email server.
- The contact name (SNMP server contact), phone number, and street address information must be

configured. This is required to determine the origin of messages received.

- A CCO ID must be associated with an appropriate Cisco SMARTnet Service contract for your company.
- Cisco SMARTnet Service must be in place for the device to be registered.

The Cisco support site contains information about the commands to configure Smart Call Home.

Install hardware

Complete Cisco Nexus 3132Q-V cabling worksheet

If you want to document the supported platforms, download a PDF of this page and complete the cabling worksheet.

The sample cabling worksheet provides examples of recommended port assignments from the switches to the controllers. The blank worksheet provides a template that you can use in setting up your cluster.

Each switch can be configured as a single 40GbE port or 4 x 10GbE ports.

Sample cabling worksheet

The sample port definition on each pair of switches is as follows:

Cluster switch A		Cluster switch B	
Switch port	Node and port usage	Switch port	Node and port usage
1	4x10G/40G node	1	4x10G/40G node
2	4x10G/40G node	2	4x10G/40G node
3	4x10G/40G node	3	4x10G/40G node
4	4x10G/40G node	4	4x10G/40G node
5	4x10G/40G node	5	4x10G/40G node
6	4x10G/40G node	6	4x10G/40G node
7	4x10G/40G node	7	4x10G/40G node
8	4x10G/40G node	8	4x10G/40G node
9	4x10G/40G node	9	4x10G/40G node
10	4x10G/40G node	10	4x10G/40G node
11	4x10G/40G node	11	4x10G/40G node

Cluster switch A		Cluster switch B	
12	4x10G/40G node	12	4x10G/40G node
13	4x10G/40G node	13	4x10G/40G node
14	4x10G/40G node	14	4x10G/40G node
15	4x10G/40G node	15	4x10G/40G node
16	4x10G/40G node	16	4x10G/40G node
17	4x10G/40G node	17	4x10G/40G node
18	4x10G/40G node	18	4x10G/40G node
19	40G node 19	19	40G node 19
20	40G node 20	20	40G node 20
21	40G node 21	21	40G node 21
22	40G node 22	22	40G node 22
23	40G node 23	23	40G node 23
24	40G node 24	24	40G node 24
25 through 30	Reserved	25 through 30	Reserved
31	40G ISL to switch B port 31	31	40G ISL to switch A port 31
32	40G ISL to switch B port 32	32	40G ISL to switch A port 32

Blank cabling worksheet

You can use the blank cabling worksheet to document the platforms that are supported as nodes in a cluster. The *Supported Cluster Connections* section of the Hardware Universe defines the cluster ports used by the platform.

Cluster switch A		Cluster switch B	
Switch port	Node/port usage	Switch port	Node/port usage
1		1	

Cluster switch A		Cluster switch B	
2		2	
3		3	
4		4	
5		5	
6		6	
7		7	
8		8	
9		9	
10		10	
11		11	
12		12	
13		13	
14		14	
15		15	
16		16	
17		17	
18		18	
19		19	
20		20	
21		21	
22		22	
23		23	

Cluster switch A		Cluster switch B	
24		24	
25 through 30	Reserved	25 through 30	Reserved
31	40G ISL to switch B port 31	31	40G ISL to switch A port 31
32	40G ISL to switch B port 32	32	40G ISL to switch A port 32

Configure the Cisco Nexus 3132Q-V switch

Follow this procedure to configure the Cisco Nexus 3132Q-V switch.

What you'll need

- Access to an HTTP, FTP or TFTP server at the installation site to download the applicable NX-OS and reference configuration file (RCF) releases.
- Applicable NX-OS version, downloaded from the Cisco software download page.
- Required network switch documentation, controller documentation, and ONTAP documentation. For more information, see Required documentation.
- Applicable licenses, network and configuration information, and cables.
- Completed cabling worksheets. See Complete Cisco Nexus 3132Q-V cabling worksheet.
- Applicable NetApp cluster network and management network RCFs, downloaded from the NetApp Support Site at mysupport.netapp.com for the switches that you receive. All Cisco cluster network and management network switches arrive with the standard Cisco factory-default configuration. These switches also have the current version of the NX-OS software, but do not have the RCFs loaded.

Steps

1. Rack the cluster network and management network switches and controllers.

If you are installing your…	Then…
Cisco Nexus 3132Q-V in a NetApp system cabinet	See the *Installing a Cisco Nexus 3132Q-V cluster switch and pass-through panel in a NetApp cabinet* guide for instructions to install the switch in a NetApp cabinet.
Equipment in a Telco rack	See the procedures provided in the switch hardware installation guides and the NetApp installation and setup instructions.

2. Cable the cluster network and management network switches to the controllers using the completed cabling worksheet, as described in Complete Cisco Nexus 3132Q-V cabling worksheet.

3. Power on the cluster network and management network switches and controllers.

4. Perform an initial configuration of the cluster network switches.

   Provide applicable responses to the following initial setup questions when you first boot the switch. Your

site's security policy defines the responses and services to enable.

Prompt	Response
Abort Auto Provisioning and continue with normal setup? (yes/no)	Respond with **yes**. The default is no.
Do you want to enforce secure password standard? (yes/no)	Respond with **yes**. The default is yes.
Enter the password for admin:	The default password is "admin"; you must create a new, strong password. A weak password can be rejected.
Would you like to enter the basic configuration dialog? (yes/no)	Respond with **yes** at the initial configuration of the switch.
Create another login account? (yes/no)	Your answer depends on your site's policies on alternate administrators. The default is **no**.
Configure read-only SNMP community string? (yes/no)	Respond with **no**. The default is no.
Configure read-write SNMP community string? (yes/no)	Respond with **no**. The default is no.
Enter the switch name.	The switch name is limited to 63 alphanumeric characters.
Continue with Out-of-band (mgmt0) management configuration? (yes/no)	Respond with **yes** (the default) at that prompt. At the mgmt0 IPv4 address: prompt, enter your IP address: ip_address.
Configure the default-gateway? (yes/no)	Respond with **yes**. At the IPv4 address of the default-gateway: prompt, enter your default_gateway.
Configure advanced IP options? (yes/no)	Respond with **no**. The default is no.
Enable the telnet service? (yes/no)	Respond with **no**. The default is no.
Enabled SSH service? (yes/no)	Respond with **yes**. The default is yes. ⓘ SSH is recommended when using Cluster Switch Health Monitor (CSHM) for its log collection features. SSHv2 is also recommended for enhanced security.

Prompt	Response
Enter the type of SSH key you want to generate (dsa/rsa/rsa1).	The default is **rsa**.
Enter the number of key bits (1024-2048).	Enter the key bits from 1024-2048.
Configure the NTP server? (yes/no)	Respond with **no**. The default is no.
Configure default interface layer (L3/L2):	Respond with **L2**. The default is L2.
Configure default switch port interface state (shut/noshut):	Respond with **noshut**. The default is noshut.
Configure CoPP system profile (strict/moderate/lenient/dense):	Respond with **strict**. The default is strict.
Would you like to edit the configuration? (yes/no)	You should see the new configuration at this point. Review and make any necessary changes to the configuration you just entered. Respond with **no** at the prompt if you are satisfied with the configuration. Respond with **yes** if you want to edit your configuration settings.
Use this configuration and save it? (yes/no)	Respond with **yes** to save the configuration. This automatically updates the kickstart and system images. ⓘ If you do not save the configuration at this stage, none of the changes will be in effect the next time you reboot the switch.

5. Verify the configuration choices you made in the display that appears at the end of the setup, and make sure that you save the configuration.

6. Check the version on the cluster network switches, and if necessary, download the NetApp-supported version of the software to the switches from the Cisco software download page.

What's next?

Prepare to install NX-OS and RCF.

Install a Cisco Nexus 3132Q-V cluster switch in a NetApp cabinet

Depending on your configuration, you might need to install the Cisco Nexus 3132Q-V switch and pass-through panel in a NetApp cabinet with the standard brackets that are included with the switch.

What you'll need

- The initial preparation requirements, kit contents, and safety precautions in the Cisco Nexus 3000 Series Hardware Installation Guide. Review these documents before you begin the procedure.

- The pass-through panel kit, available from NetApp (part number X8784-R6). The NetApp pass-through panel kit contains the following hardware:

   - One pass-through blanking panel

   - Four 10-32 x .75 screws

   - Four 10-32 clip nuts

- Eight 10-32 or 12-24 screws and clip nuts to mount the brackets and slider rails to the front and rear cabinet posts.

- Cisco standard rail kit to install the switch in a NetApp cabinet.

> ⓘ The jumper cords are not included with the pass-through kit and should be included with your switches. If they were not shipped with the switches, you can order them from NetApp (part number X1558A-R6).

Steps

1. Install the pass-through blanking panel in the NetApp cabinet.

   a. Determine the vertical location of the switches and blanking panel in the cabinet.

      In this procedure, the blanking panel will be installed in U40.

   b. Install two clip nuts on each side in the appropriate square holes for front cabinet rails.

   c. Center the panel vertically to prevent intrusion into adjacent rack space, and then tighten the screws.

   d. Insert the female connectors of both 48-inch jumper cords from the rear of the panel and through the brush assembly.

   (1) Female connector of the jumper cord.

2. Install the rack-mount brackets on the Nexus 3132Q-V switch chassis.

   a. Position a front rack-mount bracket on one side of the switch chassis so that the mounting ear is aligned with the chassis faceplate (on the PSU or fan side), and then use four M4 screws to attach the bracket to the chassis.

b. Repeat step 2a with the other front rack-mount bracket on the other side of the switch.

c. Install the rear rack-mount bracket on the switch chassis.

d. Repeat step 2c with the other rear rack-mount bracket on the other side of the switch.

3. Install the clip nuts in the square hole locations for all four IEA posts.

The two 3132Q-V switches will always be mounted in the top 2U of the cabinet RU41 and 42.

4. Install the slider rails in the cabinet.

a. Position the first slider rail at the RU42 mark on the back side of the rear left post, insert screws with the matching thread type, and then tighten the screws with your fingers.

(1) As you gently slide the slider rail, align it to the screw holes in the rack.

(2) Tighten the screws of the slider rails to the cabinet posts.

    b. Repeat step 4a for the right side rear post.

    c. Repeat steps 4a and 4b at the RU41 locations on the cabinet.

5. Install the switch in the cabinet.

> ⓘ This step requires two people: one person to support the switch from the front and another to guide the switch into the rear slider rails.

    a. Position the back of the switch at RU41.

Position switch and rails at U41 and 42

(1) As the chassis is pushed toward the rear posts, align the two rear rack-mount guides with the slider rails.

(2) Gently slide the switch until the front rack-mount brackets are flush with the front posts.

    b. Attach the switch to the cabinet.

(1) With one person holding the front of the chassis level, the other person should fully tighten the four rear screws to the cabinet posts.

c. With the chassis now supported without assistance, fully tighten the front screws to the posts.

d. Repeat steps 5a through 5c for the second switch at the RU42 location.

> ⓘ By using the fully installed switch as a support, you do not need to hold the front of the second switch during the installation process.

6. When the switches are installed, connect the jumper cords to the switch power inlets.

7. Connect the male plugs of both jumper cords to the closest available PDU outlets.

> ⓘ To maintain redundancy, the two cords must be connected to different PDUs.

8. Connect the management port on each 3132Q-V switch to either of the management switches (if ordered) or connect them directly to your management network.

   The management port is the upper-right port located on the PSU side of the switch. The CAT6 cable for each switch needs to be routed through the pass-through panel after the switches are installed to connect to the management switches or management network.

Review cabling and configuration considerations

Before configuring your Cisco 3132Q-V switch, review the following considerations.

Support for NVIDIA CX6, CX6-DX, and CX7 Ethernet ports

If connecting a switch port to an ONTAP controller using NVIDIA ConnectX-6 (CX6), ConnectX-6 Dx (CX6-DX), or ConnectX-7 (CX7) NIC ports, you must hard-code the switch port speed.

```
(cs1)(config)# interface Ethernet1/19
For 100GbE speed:
(cs1)(config-if)# speed 100000
For 40GbE speed:
(cs1)(config-if)# speed 40000
(cs1)(config-if)# no negotiate auto
(cs1)(config-if)# exit
(cs1)(config)# exit
Save the changes:
(cs1)# copy running-config startup-config
```

See the Hardware Universe for more information on switch ports.

Configure software

Prepare to install NX-OS software and Reference Configuration File

Before you install the NX-OS software and the Reference Configuration File (RCF), follow this procedure.

About the examples

The examples in this procedure use two nodes. These nodes use two 10GbE cluster interconnect ports `e0a` and `e0b`.

See the Hardware Universe to verify the correct cluster ports on your platforms.

> ⓘ   The command outputs might vary depending on different releases of ONTAP.

The examples in this procedure use the following switch and node nomenclature:

* The names of the two Cisco switches are `cs1` and `cs2`.

* The node names are `cluster1-01` and `cluster1-02`.

* The cluster LIF names are `cluster1-01_clus1` and `cluster1-01_clus2` for cluster1-01 and `cluster1-02_clus1` and `cluster1-02_clus2` for cluster1-02.

* The `cluster1::*>` prompt indicates the name of the cluster.

About this task

The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

Steps

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

   ```
   system node autosupport invoke -node * -type all -message MAINT=xh
   ```

where *x* is the duration of the maintenance window in hours.

> (i) The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Change the privilege level to advanced, entering **y** when prompted to continue:

```
set -privilege advanced
```

The advanced prompt (`*>`) appears.

3. Display how many cluster interconnect interfaces are configured in each node for each cluster interconnect switch:

```
network device-discovery show -protocol cdp
```

Show example

```
cluster1::*> network device-discovery show -protocol cdp

Node/        Local  Discovered
Protocol     Port   Device (LLDP: ChassisID)  Interface
Platform
-----------  ------ ------------------------  ------------------
--------
cluster1-02/cdp
             e0a    cs1                        Eth1/2           N3K-
C3132Q-V
             e0b    cs2                        Eth1/2           N3K-
C3132Q-V
cluster1-01/cdp
             e0a    cs1                        Eth1/1           N3K-
C3132Q-V
             e0b    cs2                        Eth1/1           N3K-
C3132Q-V
```

4. Check the administrative or operational status of each cluster interface.

   a. Display the network port attributes:

   ```
   network port show -ipspace Cluster
   ```

Show example

```
cluster1::*> network port show -ipspace Cluster

Node: cluster1-02
                                                Speed(Mbps)
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status
--------- ------------ ---------------- ---- ---- -----------
------
e0a       Cluster      Cluster          up   9000 auto/10000
healthy
e0b       Cluster      Cluster          up   9000 auto/10000
healthy

Node: cluster1-01
                                                Speed(Mbps)
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status
--------- ------------ ---------------- ---- ---- -----------
------
e0a       Cluster      Cluster          up   9000 auto/10000
healthy
e0b       Cluster      Cluster          up   9000 auto/10000
healthy
```

b. Display information about the LIFs:

```
network interface show -vserver Cluster
```

Show example

```
cluster1::*> network interface show -vserver Cluster

          Logical               Status      Network
Current        Current Is
Vserver     Interface           Admin/Oper Address/Mask       Node
Port     Home
----------- ------------------ ---------- ------------------
------------- ------- ----
Cluster
          cluster1-01_clus1  up/up       169.254.209.69/16
cluster1-01   e0a      true
          cluster1-01_clus2  up/up       169.254.49.125/16
cluster1-01   e0b      true
          cluster1-02_clus1  up/up       169.254.47.194/16
cluster1-02   e0a      true
          cluster1-02_clus2  up/up       169.254.19.183/16
cluster1-02   e0b      true
```

5. Ping the remote cluster LIFs:

```
cluster ping-cluster -node local
```

Show example

```
cluster1::*> cluster ping-cluster -node local
Host is cluster1-02
Getting addresses from network interface table...
Cluster cluster1-01_clus1 169.254.209.69 cluster1-01     e0a
Cluster cluster1-01_clus2 169.254.49.125 cluster1-01     e0b
Cluster cluster1-02_clus1 169.254.47.194 cluster1-02     e0a
Cluster cluster1-02_clus2 169.254.19.183 cluster1-02     e0b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 9000 byte MTU on 4 path(s):
    Local 169.254.19.183 to Remote 169.254.209.69
    Local 169.254.19.183 to Remote 169.254.49.125
    Local 169.254.47.194 to Remote 169.254.209.69
    Local 169.254.47.194 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)
```

6. Verify that the `auto-revert` command is enabled on all cluster LIFs:

```
network interface show -vserver Cluster -fields auto-revert
```

Show example

```
cluster1::*> network interface show -vserver Cluster -fields auto-
revert

          Logical
Vserver   Interface           Auto-revert
--------- ------------------- ------------
Cluster
          cluster1-01_clus1   true
          cluster1-01_clus2   true
          cluster1-02_clus1   true
          cluster1-02_clus2   true
```

What's next?

Install NX-OS software.

Install the NX-OS software

Follow this procedure to install the NX-OS software on the Nexus 3132Q-V cluster switch.

Review requirements

What you'll need

- A current backup of the switch configuration.
- A fully functioning cluster (no errors in the logs or similar issues).

Suggested documentation

- Cisco Ethernet switch. Consult the switch compatibility table for the supported ONTAP and NX-OS versions.
- Cisco Nexus 3000 Series Switches. Consult the appropriate software and upgrade guides available on the Cisco web site for complete documentation on the Cisco switch upgrade and downgrade procedures.

Install the software

About this task

The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

Be sure to complete the procedure in Prepare to install NX-OS software and Reference Configuration File, and then follow the steps below.

Steps

1. Connect the cluster switch to the management network.
2. Use the `ping` command to verify connectivity to the server hosting the NX-OS software and the RCF.

Show example

```
cs2# ping 172.19.2.1 vrf management
Pinging 172.19.2.1 with 0 bytes of data:

Reply From 172.19.2.1: icmp_seq = 0. time= 5910 usec.
```

3. Copy the NX-OS software to the Nexus 3132Q-V switch using one of the following transfer protocols: FTP, TFTP, SFTP, or SCP. For more information on Cisco commands, see the appropriate guide in Cisco Nexus 3000 Series NX-OS Command Reference guides.

Show example

```
cs2# copy sftp: bootflash: vrf management
Enter source filename: /code/nxos.9.3.4.bin
Enter hostname for the sftp server: 172.19.2.1
Enter username: user1

Outbound-ReKey for 172.19.2.1:22
Inbound-ReKey for 172.19.2.1:22
user1@172.19.2.1's password: xxxxxxxx
sftp> progress
Progress meter enabled
sftp> get   /code/nxos.9.3.4.bin  /bootflash/nxos.9.3.4.bin
/code/nxos.9.3.4.bin  100% 1261MB   9.3MB/s   02:15
sftp> exit
Copy complete, now saving to disk (please wait)...
Copy complete.
```

4. Verify the running version of the NX-OS software:

```
show version
```

Show example

```
cs2# show version
Cisco Nexus Operating System (NX-OS) Software
TAC support: http://www.cisco.com/tac
Copyright (C) 2002-2020, Cisco and/or its affiliates.
All rights reserved.
The copyrights to certain works contained in this software are
owned by other third parties and used and distributed under their
own
licenses, such as open source.  This software is provided "as is,"
and unless
otherwise stated, there is no warranty, express or implied,
including but not
limited to warranties of merchantability and fitness for a
particular purpose.
Certain components of this software are licensed under
the GNU General Public License (GPL) version 2.0 or
GNU General Public License (GPL) version 3.0  or the GNU
Lesser General Public License (LGPL) Version 2.1 or
Lesser General Public License (LGPL) Version 2.0.
A copy of each such license is available at
http://www.opensource.org/licenses/gpl-2.0.php and
http://opensource.org/licenses/gpl-3.0.html and
http://www.opensource.org/licenses/lgpl-2.1.php and
http://www.gnu.org/licenses/old-licenses/library.txt.

Software
  BIOS: version 04.25
NXOS: version 9.3(3)
  BIOS compile time:  01/28/2020
  NXOS image file is: bootflash:///nxos.9.3.3.bin
                 NXOS compile time: 12/22/2019 2:00:00 [12/22/2019
14:00:37]

Hardware
  cisco Nexus 3132QV Chassis (Nexus 9000 Series)
  Intel(R) Core(TM) i3- CPU @ 2.50GHz with 16399900 kB of memory.
  Processor Board ID FOxxxxxxx23

  Device name: cs2
  bootflash:   15137792 kB
  usb1:               0 kB (expansion flash)

Kernel uptime is 79 day(s), 10 hour(s), 23 minute(s), 53 second(s)
```

```
Last reset at 663500 usecs after Mon Nov  2 10:50:33 2020
  Reason: Reset Requested by CLI command reload
  System version: 9.3(3)
  Service:

plugin
  Core Plugin, Ethernet Plugin

Active Package(s):
cs2#
```

5. Install the NX-OS image.

   Installing the image file causes it to be loaded every time the switch is rebooted.

Show example

```
cs2# install all nxos bootflash:nxos.9.3.4.bin
Installer will perform compatibility check first. Please wait.
Installer is forced disruptive

Verifying image bootflash:/nxos.9.3.4.bin for boot variable "nxos".
[] 100% -- SUCCESS

Verifying image type.
[] 100% -- SUCCESS

Preparing "nxos" version info using image bootflash:/nxos.9.3.4.bin.
[] 100% -- SUCCESS

Preparing "bios" version info using image bootflash:/nxos.9.3.4.bin.
[] 100% -- SUCCESS

Performing module support checks.
[] 100% -- SUCCESS

Notifying services about system upgrade.
[] 100% -- SUCCESS


Compatibility check is done:
Module  bootable          Impact              Install-type Reason
------- ---------------- ------------------ ------------- ------
    1     yes             disruptive          reset       default
upgrade is not hitless


Images will be upgraded according to following table:
Module      Image       Running-Version(pri:alt)
New-Version        Upg-Required
------------ ---------- ---------------------------------------
-------------------- ------------
    1      nxos        9.3(3)
9.3(4)             yes
    1      bios        v04.25(01/28/2020):v04.25(10/18/2016)
v04.25(01/28/2020)   no


Switch will be reloaded for disruptive upgrade.
Do you want to continue with the installation (y/n)?  [n] y
```

```
Install is in progress, please wait.

Performing runtime checks.
[] 100% -- SUCCESS

Setting boot variables.
[] 100% -- SUCCESS

Performing configuration copy.
[] 100% -- SUCCESS

Module 1: Refreshing compact flash and upgrading
bios/loader/bootrom.
Warning: please do not remove or power off the module at this time.
[] 100% -- SUCCESS


Finishing the upgrade, switch will reboot in 10 seconds.
cs2#
```

6. Verify the new version of NX-OS software after the switch has rebooted:

```
show version
```

Show example

```
cs2# show version
Cisco Nexus Operating System (NX-OS) Software
TAC support: http://www.cisco.com/tac
Copyright (C) 2002-2020, Cisco and/or its affiliates.
All rights reserved.
The copyrights to certain works contained in this software are
owned by other third parties and used and distributed under their
own
licenses, such as open source.  This software is provided "as is,"
and unless
otherwise stated, there is no warranty, express or implied,
including but not
limited to warranties of merchantability and fitness for a
particular purpose.
Certain components of this software are licensed under
the GNU General Public License (GPL) version 2.0 or
GNU General Public License (GPL) version 3.0  or the GNU
Lesser General Public License (LGPL) Version 2.1 or
Lesser General Public License (LGPL) Version 2.0.
A copy of each such license is available at
http://www.opensource.org/licenses/gpl-2.0.php and
http://opensource.org/licenses/gpl-3.0.html and
http://www.opensource.org/licenses/lgpl-2.1.php and
http://www.gnu.org/licenses/old-licenses/library.txt.

Software
  BIOS: version 04.25
NXOS: version 9.3(4)
  BIOS compile time:  05/22/2019
  NXOS image file is: bootflash:///nxos.9.3.4.bin
  NXOS compile time:  4/28/2020 21:00:00 [04/29/2020 06:28:31]


Hardware
  cisco Nexus 3132QV Chassis (Nexus 9000 Series)
  Intel(R) Core(TM) i3- CPU @ 2.50GHz with 16399900 kB of memory.
  Processor Board ID FOxxxxxxx23

  Device name: cs2
  bootflash:    15137792 kB
  usb1:                0 kB (expansion flash)

Kernel uptime is 79 day(s), 10 hour(s), 23 minute(s), 53 second(s)
```

```
Last reset at 663500 usecs after Mon Nov  2 10:50:33 2020
  Reason: Reset Requested by CLI command reload
  System version: 9.3(4)
  Service:

plugin
  Core Plugin, Ethernet Plugin

Active Package(s):

cs2#
```

What's next?

Install the Reference Configuration File (RCF).

Install the Reference Configuration File (RCF)

Follow this procedure to install the RCF after setting up the Nexus 3132Q-V switch for the first time. You can also use this procedure to upgrade your RCF version.

Review requirements

What you'll need

- A current backup of the switch configuration.

- A fully functioning cluster (no errors in the logs or similar issues).

- The current Reference Configuration File (RCF).

- A console connection to the switch, required when installing the RCF.

- Cisco Ethernet switch. Consult the switch compatibility table for the supported ONTAP and RCF versions. Note that there can be command dependencies between the command syntax in the RCF and that found in versions of NX-OS.

- Cisco Nexus 3000 Series Switches. Consult the appropriate software and upgrade guides available on the Cisco web site for complete documentation on the Cisco switch upgrade and downgrade procedures.

Install the file

About the examples

The examples in this procedure use the following switch and node nomenclature:

- The names of the two Cisco switches are `cs1` and `cs2`.

- The node names are `cluster1-01`, `cluster1-02`, `cluster1-03`, and `cluster1-04`.

- The cluster LIF names are `cluster1-01_clus1`, `cluster1-01_clus2`, `cluster1-02_clus1`, `cluster1-02_clus2`, `cluster1-03_clus1`, `cluster1-03_clus2`, `cluster1-04_clus1`, and `cluster1-04_clus2`.

- The `cluster1::*>` prompt indicates the name of the cluster.

About this task

The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

No operational inter-switch link (ISL) is needed during this procedure. This is by design because RCF version changes can affect ISL connectivity temporarily. To ensure non-disruptive cluster operations, the following procedure migrates all of the cluster LIFs to the operational partner switch while performing the steps on the target switch.

Be sure to complete the procedure in Prepare to install NX-OS software and Reference Configuration File, and then follow the steps below.

Step 1: Check port status

1. Display the cluster ports on each node that are connected to the cluster switches:

   `network device-discovery show`

   Show example

   ```
   cluster1::*> network device-discovery show
   Node/       Local  Discovered
   Protocol    Port   Device (LLDP: ChassisID)  Interface
   Platform
   ----------- ------ ------------------------- ----------------
   ------------
   cluster1-01/cdp
               e0a    cs1                        Ethernet1/7       N3K-
   C3132Q-V
               e0d    cs2                        Ethernet1/7       N3K-
   C3132Q-V
   cluster1-02/cdp
               e0a    cs1                        Ethernet1/8       N3K-
   C3132Q-V
               e0d    cs2                        Ethernet1/8       N3K-
   C3132Q-V
   cluster1-03/cdp
               e0a    cs1                        Ethernet1/1/1     N3K-
   C3132Q-V
               e0b    cs2                        Ethernet1/1/1     N3K-
   C3132Q-V
   cluster1-04/cdp
               e0a    cs1                        Ethernet1/1/2     N3K-
   C3132Q-V
               e0b    cs2                        Ethernet1/1/2     N3K-
   C3132Q-V
   cluster1::*>
   ```

2. Check the administrative and operational status of each cluster port.

    a. Verify that all the cluster ports are up with a healthy status:

    `network port show -ipspace Cluster`

```
cluster1::*> network port show -ipspace Cluster


Node: cluster1-01


Ignore
                                                          Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000  auto/100000
healthy false
e0d       Cluster      Cluster          up   9000  auto/100000
healthy false


Node: cluster1-02


Ignore
                                                          Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000  auto/100000
healthy false
e0d       Cluster      Cluster          up   9000  auto/100000
healthy false
8 entries were displayed.


Node: cluster1-03


   Ignore
                                                          Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000  auto/10000
healthy  false
e0b       Cluster      Cluster          up   9000  auto/10000
healthy  false
```

```
Node: cluster1-04

Ignore
                                                      Speed(Mbps)
Health    Health
Port        IPspace        Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a        Cluster        Cluster              up    9000  auto/10000
healthy  false
e0b        Cluster        Cluster              up    9000  auto/10000
healthy  false
cluster1::*>
```

b. Verify that all the cluster interfaces (LIFs) are on the home port:

```
network interface show -vserver Cluster
```

Show example

```
cluster1::*> network interface show -vserver Cluster
            Logical             Status      Network
Current        Current Is
Vserver      Interface           Admin/Oper Address/Mask      Node
Port    Home
----------- ------------------ ---------- -----------------
------------ ------- ----
Cluster
            cluster1-01_clus1   up/up      169.254.3.4/23
cluster1-01  e0a      true
            cluster1-01_clus2   up/up      169.254.3.5/23
cluster1-01  e0d      true
            cluster1-02_clus1   up/up      169.254.3.8/23
cluster1-02  e0a      true
            cluster1-02_clus2   up/up      169.254.3.9/23
cluster1-02  e0d      true
            cluster1-03_clus1   up/up      169.254.1.3/23
cluster1-03  e0a      true
            cluster1-03_clus2   up/up      169.254.1.1/23
cluster1-03  e0b      true
            cluster1-04_clus1   up/up      169.254.1.6/23
cluster1-04  e0a      true
            cluster1-04_clus2   up/up      169.254.1.7/23
cluster1-04  e0b      true
cluster1::*>
```

c. Verify that the cluster displays information for both cluster switches:

```
system cluster-switch show -is-monitoring-enabled-operational true
```

Show example

```
cluster1::*> system cluster-switch show -is-monitoring-enabled
-operational true
Switch                         Type             Address
Model
--------------------------- ------------------ ----------------
--------------
cs1                            cluster-network    10.0.0.1
NX3132QV
        Serial Number: FOXXXXXXXGS
        Is Monitored: true
              Reason: None
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                   9.3(4)
     Version Source: CDP

cs2                            cluster-network    10.0.0.2
NX3132QV
        Serial Number: FOXXXXXXXGD
        Is Monitored: true
              Reason: None
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                   9.3(4)
     Version Source: CDP

2 entries were displayed.
```

> (i) For ONTAP 9.8 and later, use the command `system switch ethernet show -is
> -monitoring-enabled-operational true`.

3. Disable auto-revert on the cluster LIFs.

```
cluster1::*> network interface modify -vserver Cluster -lif * -auto
-revert false
```

Make sure that auto-revert is disabled after running this command.

4. On cluster switch cs2, shut down the ports connected to the cluster ports of the nodes.

```
cs2(config)# interface eth1/1/1-2,eth1/7-8
cs2(config-if-range)# shutdown
```

5. Verify that the cluster ports have migrated to the ports hosted on cluster switch cs1. This might take a few seconds.

```
network interface show -vserver Cluster
```

Show example

```
cluster1::*> network interface show -vserver Cluster
          Logical              Status     Network                Current
Current Is
Vserver     Interface          Admin/Oper Address/Mask          Node
Port    Home
----------- ----------------- ---------- ------------------
------------ ------- ----
Cluster
          cluster1-01_clus1 up/up       169.254.3.4/23
cluster1-01   e0a      true
          cluster1-01_clus2 up/up       169.254.3.5/23
cluster1-01   e0a      false
          cluster1-02_clus1 up/up       169.254.3.8/23
cluster1-02   e0a      true
          cluster1-02_clus2 up/up       169.254.3.9/23
cluster1-02   e0a      false
          cluster1-03_clus1 up/up       169.254.1.3/23
cluster1-03   e0a      true
          cluster1-03_clus2 up/up       169.254.1.1/23
cluster1-03   e0a      false
          cluster1-04_clus1 up/up       169.254.1.6/23
cluster1-04   e0a      true
          cluster1-04_clus2 up/up       169.254.1.7/23
cluster1-04   e0a      false
cluster1::*>
```

6. Verify that the cluster is healthy:

```
cluster show
```

Show example

```
cluster1::*> cluster show
Node                    Health  Eligibility   Epsilon
--------------------    -------  ------------   -------
cluster1-01             true    true           false
cluster1-02             true    true           false
cluster1-03             true    true           true
cluster1-04             true    true           false
cluster1::*>
```

Step 2: Configure and verify the setup

1. If you have not already done so, save a copy of the current switch configuration by copying the output of the following command to a text file:

```
show running-config
```

2. Clean the configuration on switch cs2 and perform a basic setup.

> ⚠️ When updating or applying a new RCF, you must erase the switch settings and perform basic configuration. You must be connected to the switch serial console port to set up the switch again.

   a. Clean the configuration:

   Show example

```
(cs2)# write erase

Warning: This command will erase the startup-configuration.

Do you wish to proceed anyway? (y/n)  [n]  y
```

   b. Perform a reboot of the switch:

   Show example

```
(cs2)# reload

Are you sure you would like to reset the system? (y/n) y
```

3. Copy the RCF to the bootflash of switch cs2 using one of the following transfer protocols: FTP, TFTP, SFTP, or SCP. For more information on Cisco commands, see the appropriate guide in the Cisco Nexus 3000 Series NX-OS Command Reference guides.

Show example

```
cs2# copy tftp: bootflash: vrf management
Enter source filename: Nexus_3132QV_RCF_v1.6-Cluster-HA-Breakout.txt
Enter hostname for the tftp server: 172.22.201.50
Trying to connect to tftp server......Connection to Server
Established.
TFTP get operation was successful
Copy complete, now saving to disk (please wait)...
```

4. Apply the RCF previously downloaded to the bootflash.

   For more information on Cisco commands, see the appropriate guide in the Cisco Nexus 3000 Series NX-OS Command Reference guides.

Show example

```
cs2# copy Nexus_3132QV_RCF_v1.6-Cluster-HA-Breakout.txt running-
config echo-commands
```

5. Examine the banner output from the `show banner motd` command. You must read and follow the instructions under **Important Notes** to ensure the proper configuration and operation of the switch.

Show example

```
cs2# show banner motd

*********************************************************************
**********
* NetApp Reference Configuration File (RCF)
*
* Switch   : Cisco Nexus 3132Q-V
* Filename : Nexus_3132QV_RCF_v1.6-Cluster-HA-Breakout.txt
* Date     : Nov-02-2020
* Version  : v1.6
*
* Port Usage : Breakout configuration
* Ports  1- 6: Breakout mode (4x10GbE) Intra-Cluster Ports, int
e1/1/1-4,
* e1/2/1-4, e1/3/1-4,int e1/4/1-4, e1/5/1-4, e1/6/1-4
* Ports  7-30: 40GbE Intra-Cluster/HA Ports, int e1/7-30
* Ports 31-32: Intra-Cluster ISL Ports, int e1/31-32
*
* IMPORTANT NOTES
* - Load Nexus_3132QV_RCF_v1.6-Cluster-HA.txt for non breakout
config
*
* - This RCF utilizes QoS and requires specific TCAM configuration,
requiring
*   cluster switch to be rebooted before the cluster becomes
operational.
*
* - Perform the following steps to ensure proper RCF installation:
*
*    (1) Apply RCF, expect following messages:
*        - Please save config and reload the system...
*        - Edge port type (portfast) should only be enabled on
ports...
*        - TCAM region is not configured for feature QoS class
IPv4...
*
*    (2) Save running-configuration and reboot Cluster Switch
*
*    (3) After reboot, apply same RCF second time and expect
following messages:
*        - % Invalid command at '^' marker
*
*    (4) Save running-configuration again
*
```

```
*  - If running NX-OS versions 9.3(5) 9.3(6), 9.3(7), or 9.3(8)
*     - Downgrade the NX-OS firmware to version 9.3(5) or earlier if
*        NX-OS using a version later than 9.3(5).
*     - Do not upgrade NX-OS prior to applying v1.9 RCF file.
*     - After the RCF is applied and switch rebooted, then proceed to
upgrade
*        NX-OS to version 9.3(5) or later.
*
* - If running 9.3(9) 10.2(2) or later the RCF can be applied to the
switch
*        after the upgrade.
*
* - Port 1 multiplexed H/W configuration options:
*     hardware profile front portmode qsfp       (40G H/W port 1/1 is
active - default)
*     hardware profile front portmode sfp-plus   (10G H/W ports 1/1/1
- 1/1/4 are active)
*     hardware profile front portmode qsfp       (To reset to QSFP)
*
********************************************************************
**********
```

6. Verify that the RCF file is the correct newer version:

   ```
   show running-config
   ```

   When you check the output to verify you have the correct RCF, make sure that the following information is correct:

   ◦ The RCF banner

   ◦ The node and port settings

   ◦ Customizations

   The output varies according to your site configuration. Check the port settings and refer to the release notes for any changes specific to the RCF that you have installed.

   > (i) For steps on how to bring your 10GbE ports online after an upgrade of the RCF, see the Knowledge Base article 10GbE ports on a Cisco 3132Q cluster switch do not come online.

7. After you verify the RCF versions and switch settings are correct, copy the running-config file to the startup-config file.

   For more information on Cisco commands, see the appropriate guide in the Cisco Nexus 3000 Series NX-OS Command Reference guides.

Show example

```
cs2# copy running-config startup-config
[##################################] 100% Copy complete
```

8. Reboot switch cs2. You can ignore the "cluster ports down" events reported on the nodes while the switch reboots.

Show example

```
cs2# reload
This command will reboot the system. (y/n)?  [n] y
```

9. Apply the same RCF and save the running configuration for a second time.

Show example

```
cs2# copy Nexus_3132QV_RCF_v1.6-Cluster-HA-Breakout.txt running-
config echo-commands
cs2# copy running-config startup-config
[##################################] 100% Copy complete
```

10. Verify the health of cluster ports on the cluster.

    a. Verify that cluster ports are up and healthy across all nodes in the cluster:

    ```
    network port show -ipspace Cluster
    ```

Show example

```
cluster1::*> network port show -ipspace Cluster

Node: cluster1-01

Ignore
                                                  Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a      Cluster      Cluster          up   9000  auto/10000
healthy  false
e0b      Cluster      Cluster          up   9000  auto/10000
healthy  false

Node: cluster1-02

Ignore
                                                  Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a      Cluster      Cluster          up   9000  auto/10000
healthy  false
e0b      Cluster      Cluster          up   9000  auto/10000
healthy  false

Node: cluster1-03

Ignore
                                                  Speed(Mbps)
Health   Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a      Cluster      Cluster          up   9000  auto/100000
healthy false
e0d      Cluster      Cluster          up   9000  auto/100000
healthy false
```

```
  Node: cluster1-04

  Ignore

                                                Speed(Mbps)
  Health    Health
  Port       IPspace       Broadcast Domain Link MTU  Admin/Oper
  Status    Status
  --------- ------------ ---------------- ---- ---- -----------
  -------- ------
  e0a       Cluster      Cluster             up   9000  auto/100000
  healthy false
  e0d       Cluster      Cluster             up   9000  auto/100000
  healthy false
```

b. Verify the switch health from the cluster.

```
network device-discovery show -protocol cdp
```

Show example

```
cluster1::*> network device-discovery show -protocol cdp
Node/      Local  Discovered
Protocol   Port   Device (LLDP: ChassisID)  Interface
Platform
---------- ------ ------------------------ -----------------
--------
cluster1-01/cdp
           e0a    cs1                       Ethernet1/7
N3K-C3132Q-V
           e0d    cs2                       Ethernet1/7
N3K-C3132Q-V
cluster01-2/cdp
           e0a    cs1                       Ethernet1/8
N3K-C3132Q-V
           e0d    cs2                       Ethernet1/8
N3K-C3132Q-V
cluster01-3/cdp
           e0a    cs1                       Ethernet1/1/1
N3K-C3132Q-V
           e0b    cs2                       Ethernet1/1/1
N3K-C3132Q-V
cluster1-04/cdp
           e0a    cs1                       Ethernet1/1/2
N3K-C3132Q-V
           e0b    cs2                       Ethernet1/1/2
N3K-C3132Q-V

cluster1::*> system cluster-switch show -is-monitoring-enabled
-operational true
Switch                         Type              Address
Model
-------------------------- ----------------- ----------------
-----
cs1                            cluster-network   10.233.205.90
N3K-C3132Q-V
     Serial Number: FOXXXXXXXGD
      Is Monitored: true
            Reason: None
  Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                  9.3(4)
    Version Source: CDP

cs2                            cluster-network   10.233.205.91
```

```
N3K-C3132Q-V
        Serial Number: FOXXXXXXXGS
          Is Monitored: true
                Reason: None
      Software Version: Cisco Nexus Operating System (NX-OS)
    Software, Version
                       9.3(4)
        Version Source: CDP

    2 entries were displayed.
```

> (i) For ONTAP 9.8 and later, use the command `system switch ethernet show -is -monitoring-enabled-operational true`.

> (i) You might observe the following output on the cs1 switch console depending on the RCF version previously loaded on the switch:
>
> ```
> 2020 Nov 17 16:07:18 cs1 %$ VDC-1 %$ %STP-2-
> UNBLOCK_CONSIST_PORT: Unblocking port port-channel1 on
> VLAN0092. Port consistency restored.
> 2020 Nov 17 16:07:23 cs1 %$ VDC-1 %$ %STP-2-
> BLOCK_PVID_PEER: Blocking port-channel1 on VLAN0001.
> Inconsistent peer vlan.
> 2020 Nov 17 16:07:23 cs1 %$ VDC-1 %$ %STP-2-
> BLOCK_PVID_LOCAL: Blocking port-channel1 on VLAN0092.
> Inconsistent local vlan.
> ```

> (i) It can take up to 5 minutes for the cluster nodes to report as healthy.

11. On cluster switch cs1, shut down the ports connected to the cluster ports of the nodes.

Show example

```
cs1(config)# interface eth1/1/1-2,eth1/7-8
cs1(config-if-range)# shutdown
```

12. Verify that the cluster LIFs have migrated to the ports hosted on switch cs2. This might take a few seconds.

```
network interface show -vserver Cluster
```

Show example

```
cluster1::*> network interface show -vserver Cluster
          Logical                Status     Network                 Current
Current Is
Vserver     Interface          Admin/Oper Address/Mask         Node
Port    Home
----------- ----------------- ---------- ------------------
------------------ ------- ----
Cluster
          cluster1-01_clus1  up/up      169.254.3.4/23
cluster1-01        e0d    false
          cluster1-01_clus2  up/up      169.254.3.5/23
cluster1-01        e0d    true
          cluster1-02_clus1  up/up      169.254.3.8/23
cluster1-02        e0d    false
          cluster1-02_clus2  up/up      169.254.3.9/23
cluster1-02        e0d    true
          cluster1-03_clus1  up/up      169.254.1.3/23
cluster1-03        e0b    false
          cluster1-03_clus2  up/up      169.254.1.1/23
cluster1-03        e0b    true
          cluster1-04_clus1  up/up      169.254.1.6/23
cluster1-04        e0b    false
          cluster1-04_clus2  up/up      169.254.1.7/23
cluster1-04        e0b    true
cluster1::*>
```

13. Verify that the cluster is healthy:

```
cluster show
```

Show example

```
cluster1::*> cluster show
Node                 Health   Eligibility   Epsilon
-------------------- -------- ------------- -------
cluster1-01          true     true          false
cluster1-02          true     true          false
cluster1-03          true     true          true
cluster1-04          true     true          false
4 entries were displayed.
cluster1::*>
```

14. Repeat Steps 1 to 10 on switch cs1.

15. Enable auto-revert on the cluster LIFs.

Show example

```
cluster1::*> network interface modify -vserver Cluster -lif * -auto
-revert True
```

16. Reboot switch cs1. You do this to trigger the cluster LIFs to revert to their home ports. You can ignore the "cluster ports down" events reported on the nodes while the switch reboots.

```
cs1# reload
This command will reboot the system. (y/n)?  [n] y
```

Step 3: Verify the configuration

1. Verify that the switch ports connected to the cluster ports are up.

```
show interface brief | grep up
```

Show example

```
cs1# show interface brief | grep up
.
.
Eth1/1/1      1       eth  access up      none
10G(D) --
Eth1/1/2      1       eth  access up      none
10G(D) --
Eth1/7        1       eth  trunk  up      none
100G(D) --
Eth1/8        1       eth  trunk  up      none
100G(D) --
.
.
```

2. Verify that the ISL between cs1 and cs2 is functional:

```
show port-channel summary
```

Show example

```
cs1# show port-channel summary
Flags:  D - Down         P - Up in port-channel (members)
        I - Individual  H - Hot-standby (LACP only)
        s - Suspended   r - Module-removed
        b - BFD Session Wait
        S - Switched     R - Routed
        U - Up (port-channel)
        p - Up in delay-lacp mode (member)
        M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-       Type     Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1    Po1(SU)      Eth      LACP      Eth1/31(P)   Eth1/32(P)
cs1#
```

3. Verify that the cluster LIFs have reverted to their home port:

```
network interface show -vserver Cluster
```

Show example

```
cluster1::*> network interface show -vserver Cluster
            Logical               Status      Network                    Current
Current Is
Vserver     Interface             Admin/Oper Address/Mask        Node
Port    Home
----------- ----------------- ---------- ------------------
------------------ ------- ----
Cluster
            cluster1-01_clus1  up/up       169.254.3.4/23
cluster1-01         e0d      true
            cluster1-01_clus2  up/up       169.254.3.5/23
cluster1-01         e0d      true
            cluster1-02_clus1  up/up       169.254.3.8/23
cluster1-02         e0d      true
            cluster1-02_clus2  up/up       169.254.3.9/23
cluster1-02         e0d      true
            cluster1-03_clus1  up/up       169.254.1.3/23
cluster1-03         e0b      true
            cluster1-03_clus2  up/up       169.254.1.1/23
cluster1-03         e0b      true
            cluster1-04_clus1  up/up       169.254.1.6/23
cluster1-04         e0b      true
            cluster1-04_clus2  up/up       169.254.1.7/23
cluster1-04         e0b      true
cluster1::*>
```

4. Verify that the cluster is healthy:

```
cluster show
```

Show example

```
cluster1::*> cluster show
Node                 Health  Eligibility   Epsilon
-------------------- ------- ------------- -------
cluster1-01          true    true          false
cluster1-02          true    true          false
cluster1-03          true    true          true
cluster1-04          true    true          false
cluster1::*>
```

5. Ping the remote cluster interfaces to verify connectivity:

`cluster ping-cluster -node local`

Show example

```
cluster1::*> cluster ping-cluster -node local
Host is cluster1-03
Getting addresses from network interface table...
Cluster cluster1-03_clus1 169.254.1.3 cluster1-03 e0a
Cluster cluster1-03_clus2 169.254.1.1 cluster1-03 e0b
Cluster cluster1-04_clus1 169.254.1.6 cluster1-04 e0a
Cluster cluster1-04_clus2 169.254.1.7 cluster1-04 e0b
Cluster cluster1-01_clus1 169.254.3.4 cluster1-01 e0a
Cluster cluster1-01_clus2 169.254.3.5 cluster1-01 e0d
Cluster cluster1-02_clus1 169.254.3.8 cluster1-02 e0a
Cluster cluster1-02_clus2 169.254.3.9 cluster1-02 e0d
Local = 169.254.1.3 169.254.1.1
Remote = 169.254.1.6 169.254.1.7 169.254.3.4 169.254.3.5 169.254.3.8
169.254.3.9
Cluster Vserver Id = 4294967293
Ping status:
............
Basic connectivity succeeds on 12 path(s)
Basic connectivity fails on 0 path(s)
................................................
Detected 9000 byte MTU on 12 path(s):
    Local 169.254.1.3 to Remote 169.254.1.6
    Local 169.254.1.3 to Remote 169.254.1.7
    Local 169.254.1.3 to Remote 169.254.3.4
    Local 169.254.1.3 to Remote 169.254.3.5
    Local 169.254.1.3 to Remote 169.254.3.8
    Local 169.254.1.3 to Remote 169.254.3.9
    Local 169.254.1.1 to Remote 169.254.1.6
    Local 169.254.1.1 to Remote 169.254.1.7
    Local 169.254.1.1 to Remote 169.254.3.4
    Local 169.254.1.1 to Remote 169.254.3.5
    Local 169.254.1.1 to Remote 169.254.3.8
    Local 169.254.1.1 to Remote 169.254.3.9
Larger than PMTU communication succeeds on 12 path(s)
RPC status:
6 paths up, 0 paths down (tcp check)
6 paths up, 0 paths down (udp check)
```

6. For ONTAP 9.8 and later, enable the Ethernet switch health monitor log collection feature for collecting

switch-related log files by using the commands:

`system switch ethernet log setup-password` and

`system switch ethernet log enable-collection`

a. Enter: `system switch ethernet log setup-password`

Show example

```
cluster1::*> system switch ethernet log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs1
RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>
```

b. Enter: `system switch ethernet log enable-collection`

Show example

```
cluster1::*> system  switch ethernet log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>
```

ⓘ If any of these commands return an error, contact NetApp support.

7. For ONTAP releases 9.5P16, 9.6P12, and 9.7P10 and later patch releases, enable the Ethernet switch health monitor log collection feature for collecting switch-related log files by using the commands:

`system cluster-switch log setup-password` and

`system cluster-switch log enable-collection`

  a. Enter: `system cluster-switch log setup-password`

Show example

```
cluster1::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs1
RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system cluster-switch log setup-password

Enter the switch name: cs2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>
```

b. Enter: `system cluster-switch log enable-collection`

Show example

```
cluster1::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*>
```

Ethernet Switch Health Monitoring log collection

You can use the log collection feature to collect switch-related log files in ONTAP.

The Ethernet switch health monitor (CSHM) is responsible for ensuring the operational health of Cluster and Storage network switches and collecting switch logs for debugging purposes. This procedure guides you through the process of setting up and starting the collection of detailed **Support** logs from the switch and starts an hourly collection of **Periodic** data that is collected by AutoSupport.

Before you begin

- Verify that you have set up your environment using the Cisco 3132Q-V cluster switch **CLI**.

- Switch health monitoring must be enabled for the switch. Verify this by ensuring the `Is Monitored:` field is set to **true** in the output of the `system switch ethernet show` command.

Steps

1. Create a password for the Ethernet switch health monitor log collection feature:

   `system switch ethernet log setup-password`

Show example

```
cluster1::*> system switch ethernet log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
cs1
cs2

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs1
Would you like to specify a user other than admin for log
collection? {y|n}: n

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster1::*> system switch ethernet log setup-password

Enter the switch name: cs2
Would you like to specify a user other than admin for log
collection? {y|n}: n

Enter the password: <enter switch password>
Enter the password again: <enter switch password>
```

2. To start log collection, run the following command, replacing DEVICE with the switch used in the previous command. This starts both types of log collection: the detailed **Support** logs and an hourly collection of **Periodic** data.

```
system switch ethernet log modify -device <switch-name> -log-request true
```

Show example

```
cluster1::*> system switch ethernet log modify -device cs1 -log
-request true

Do you want to modify the cluster switch log collection
configuration? {y|n}: [n] y

Enabling cluster switch log collection.

cluster1::*> system switch ethernet log modify -device cs2 -log
-request true

Do you want to modify the cluster switch log collection
configuration? {y|n}: [n] y

Enabling cluster switch log collection.
```

Wait for 10 minutes, and then check that the log collection completes:

```
system switch ethernet log show
```

> ⚠️ If any of these commands return an error or if the log collection does not complete, contact NetApp support.

Troubleshooting

If you encounter any of the following error statuses reported by the log collection feature (visible in the output of `system switch ethernet log show`), try the corresponding debug steps:

Log collection error status	Resolution
RSA keys not present	Regenerate ONTAP SSH keys. Contact NetApp support.
switch password error	Verify credentials, test SSH connectivity, and regenerate ONTAP SSH keys. Review the switch documentation or contact NetApp support for instructions.
ECDSA keys not present for FIPS	If FIPS mode is enabled, ECDSA keys need to be generated on the switch before retrying.
pre-existing log found	Remove the previous log collection file on the switch.

switch dump log error	Ensure the switch user has log collection permissions. Refer to the prerequisites above.

Configure SNMPv3

Follow this procedure to configure SNMPv3, which supports Ethernet switch health monitoring (CSHM).

About this task

The following commands configure an SNMPv3 username on Cisco 3132Q-V switches:

* For **no authentication**: `snmp-server user` *SNMPv3_USER* `NoAuth`

* For **MD5/SHA authentication**: `snmp-server user` *SNMPv3_USER* `auth [md5|sha]` *AUTH-PASSWORD*

* For **MD5/SHA authentication with AES/DES encryption**: `snmp-server user` *SNMPv3_USER* `AuthEncrypt auth [md5|sha]` *AUTH-PASSWORD* `priv aes-128` *PRIV-PASSWORD*

The following command configures an SNMPv3 username on the ONTAP side: `cluster1::*> security login create -user-or-group-name` *SNMPv3_USER* `-application snmp -authentication -method usm -remote-switch-ipaddress` *ADDRESS*

The following command establishes the SNMPv3 username with CSHM: `cluster1::*> system switch ethernet modify -device` *DEVICE* `-snmp-version SNMPv3 -community-or-username` *SNMPv3_USER*

Steps

1. Set up the SNMPv3 user on the switch to use authentication and encryption:

   ```
   show snmp user
   ```

Show example

```
(sw1)(Config)# snmp-server user SNMPv3User auth md5 <auth_password>
priv aes-128 <priv_password>

(sw1)(Config)# show snmp user

-----------------------------------------------------------------------
---------
                                SNMP USERS
-----------------------------------------------------------------------
---------

User             Auth            Priv(enforce)    Groups
acl_filter
---------------- -------------- -------------- ---------------
-----------
admin            md5             des(no)          network-admin
SNMPv3User       md5             aes-128(no)      network-operator


-----------------------------------------------------------------------
---------
     NOTIFICATION TARGET USERS (configured  for sending V3 Inform)
-----------------------------------------------------------------------
---------

User             Auth             Priv
---------------- ----------------- ------------

(sw1)(Config)#
```

2. Set up the SNMPv3 user on the ONTAP side:

```
security login create -user-or-group-name <username> -application snmp
-authentication-method usm -remote-switch-ipaddress 10.231.80.212
```

Show example

```
cluster1::*> system switch ethernet modify -device "sw1
(b8:59:9f:09:7c:22)" -is-monitoring-enabled-admin true

cluster1::*> security login create -user-or-group-name <username>
-application snmp -authentication-method usm -remote-switch
-ipaddress 10.231.80.212


Enter the authoritative entity's EngineID [remote EngineID]:

Which authentication protocol do you want to choose (none, md5, sha,
sha2-256)
[none]: md5

Enter the authentication protocol password (minimum 8 characters
long):

Enter the authentication protocol password again:

Which privacy protocol do you want to choose (none, des, aes128)
[none]: aes128

Enter privacy protocol password (minimum 8 characters long):
Enter privacy protocol password again:
```

3. Configure CSHM to monitor with the new SNMPv3 user:

```
system switch ethernet show-all -device "sw1" -instance
```

Show example

```
cluster1::*> system switch ethernet show-all -device "sw1" -instance

                                Device Name: sw1
                                  IP Address: 10.231.80.212
                               SNMP Version: SNMPv2c
                              Is Discovered: true
    SNMPv2c Community String or SNMPv3 Username: cshm1!
                               Model Number: N3K-C3132Q-V
                             Switch Network: cluster-network
                           Software Version: Cisco Nexus
Operating System (NX-OS) Software, Version 9.3(7)
                     Reason For Not Monitoring: None  <---- displays
when SNMP settings are valid
                      Source Of Switch Version: CDP/ISDP
                              Is Monitored ?: true
                  Serial Number of the Device: QTFCU3826001C
                                 RCF Version: v1.8X2 for
Cluster/HA/RDMA

cluster1::*>
cluster1::*> system switch ethernet modify -device "sw1" -snmp
-version SNMPv3 -community-or-username <username>
cluster1::*>
```

4. Verify that the serial number to be queried with the newly created SNMPv3 user is the same as detailed in the previous step after the CSHM polling period has completed.

```
system switch ethernet polling-interval show
```

Show example

```
cluster1::*> system switch ethernet polling-interval show
          Polling Interval (in minutes): 5

cluster1::*> system switch ethernet show-all -device "sw1" -instance

                                  Device Name: sw1
                                   IP Address: 10.231.80.212
                                 SNMP Version: SNMPv3
                                Is Discovered: true
   SNMPv2c Community String or SNMPv3 Username: SNMPv3User
                                 Model Number: N3K-C3132Q-V
                               Switch Network: cluster-network
                             Software Version: Cisco Nexus
Operating System (NX-OS) Software, Version 9.3(7)
                      Reason For Not Monitoring: None  <---- displays
when SNMP settings are valid
                        Source Of Switch Version: CDP/ISDP
                                Is Monitored ?: true
                    Serial Number of the Device: QTFCU3826001C
                                  RCF Version: v1.8X2 for
Cluster/HA/RDMA

cluster1::*>
```

Migrate switches

Migrate a Cisco Nexus 5596 cluster switch to a Cisco Nexus 3132Q-V cluster switch

Follow this procedure to replace an existing Nexus 5596 cluster switch with a Nexus 3132Q-V cluster switch.

Review requirements

Review the Cisco Nexus 5596 requirements in Requirements for replacing Cisco Nexus 3132Q-V cluster switches.

For more information, see:

- Cisco Ethernet Switch description page
- Hardware Universe

Replace the switch

About the examples

The examples in this procedure describe replacing Nexus 5596 switches with Nexus 3132Q-V switches. You can use these steps (with modifications) to replace other older Cisco switches.

The procedure uses the following switch and node nomenclature:

- The command outputs might vary depending on different releases of ONTAP.
- The Nexus 5596 switches to be replaced are CL1 and CL2.
- The Nexus 3132Q-V switches to replace the Nexus 5596 switches are C1 and C2.
- n1_clus1 is the first cluster logical interface (LIF) connected to cluster switch 1 (CL1 or C1) for node n1.
- n1_clus2 is the first cluster LIF connected to cluster switch 2 (CL2 or C2) for node n1.
- n1_clus3 is the second LIF connected to cluster switch 2 (CL2 or C2) for node n1.
- n1_clus4 is the second LIF connected to cluster switch 1 (CL1 or C1) for node n1.
- The nodes are n1, n2, n3, and n4.
- The examples in this procedure use four nodes: Two nodes use four 10 GbE cluster interconnect ports: e0a, e0b, e0c, and e0d. The other two nodes use two 40/100 GbE cluster interconnect ports: e4a, e4e. The Hardware Universe lists the actual cluster ports on your platforms.
- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco® Cluster Network Switch Reference Configuration File Download page.

> ⓘ The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

About this task

This procedure covers the following scenarios:

- The cluster starts with two nodes connected and functioning in a 2 Nexus 5596 cluster switches.
- The cluster switch CL2 to be replaced by C2 (Steps 1 - 19)
  - Traffic on all cluster ports and LIFs on all nodes connected to CL2 are migrated onto the first cluster ports and LIFs connected to CL1.
  - Disconnect cabling from all cluster ports on all nodes connected to CL2, and then use supported break-out cabling to reconnect the ports to new cluster switch C2.
  - Disconnect cabling between ISL ports between CL1 and CL2, and then use supported break-out cabling to reconnect the ports from CL1 to C2.
  - Traffic on all cluster ports and LIFs connected to C2 on all nodes is reverted.
- The cluster switch CL2 to be replaced by C2
  - Traffic on all cluster ports or LIFs on all nodes connected to CL1 are migrated onto the second cluster ports or LIFs connected to C2.
  - Disconnect cabling from all cluster port on all nodes connected to CL1 and reconnect, using supported break-out cabling, to new cluster switch C1.
  - Disconnect cabling between ISL ports between CL1 and C2, and reconnect using supported cabling, from C1 to C2.

- ◦ Traffic on all cluster ports or LIFs connected to C1 on all nodes is reverted.
- Two FAS9000 nodes have been added to cluster with examples showing cluster details.

Step 1: Prepare for replacement

To replace an existing Nexus 5596 cluster switch with a Nexus 3132Q-V cluster switch, you must perform a specific sequence of tasks.

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message: `system node autosupport invoke -node * -type all -message MAINT=xh`

   x is the duration of the maintenance window in hours.

   > (i) The message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Display information about the devices in your configuration:

   `network device-discovery show`

   Show example

   The following example shows how many cluster interconnect interfaces have been configured in each node for each cluster interconnect switch:

   ```
   cluster::> network device-discovery show
               Local  Discovered
   Node         Port  Device                Interface         Platform
   ----------- ------ -------------------- ----------------
   ----------------
   n1          /cdp
               e0a    CL1                   Ethernet1/1       N5K-C5596UP
               e0b    CL2                   Ethernet1/1       N5K-C5596UP
               e0c    CL2                   Ethernet1/2       N5K-C5596UP
               e0d    CL1                   Ethernet1/2       N5K-C5596UP
   n2          /cdp
               e0a    CL1                   Ethernet1/3       N5K-C5596UP
               e0b    CL2                   Ethernet1/3       N5K-C5596UP
               e0c    CL2                   Ethernet1/4       N5K-C5596UP
               e0d    CL1                   Ethernet1/4       N5K-C5596UP
   8 entries were displayed.
   ```

3. Determine the administrative or operational status for each cluster interface:

   a. Display the network port attributes:

      `network port show`

Show example

The following example displays the network port attributes on a system:

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1

Ignore
                                                   Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000 auto/10000  -
-
e0b       Cluster      Cluster          up   9000 auto/10000  -
-
e0c       Cluster      Cluster          up   9000 auto/10000  -
-
e0d       Cluster      Cluster          up   9000 auto/10000  -
-


Node: n2

Ignore
                                                   Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000  auto/10000 -
-
e0b       Cluster      Cluster          up   9000  auto/10000 -
-
e0c       Cluster      Cluster          up   9000  auto/10000 -
-
e0d       Cluster      Cluster          up   9000  auto/10000 -
-
8 entries were displayed.
```

b. Display information about the logical interfaces:
   `network interface show`

Show example

The following example displays the general information about all of the LIFs on your system:

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status      Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node
Port    Home
----------- ---------- ---------- ------------------
------------- ------- ----
Cluster
            n1_clus1   up/up       10.10.0.1/24         n1
e0a     true
            n1_clus2   up/up       10.10.0.2/24         n1
e0b     true
            n1_clus3   up/up       10.10.0.3/24         n1
e0c     true
            n1_clus4   up/up       10.10.0.4/24         n1
e0d     true
            n2_clus1   up/up       10.10.0.5/24         n2
e0a     true
            n2_clus2   up/up       10.10.0.6/24         n2
e0b     true
            n2_clus3   up/up       10.10.0.7/24         n2
e0c     true
            n2_clus4   up/up       10.10.0.8/24         n2
e0d     true
8 entries were displayed.
```

c. Display information about the discovered cluster switches:
`system cluster-switch show`

Show example

The following example displays the cluster switches that are known to the cluster, along with their management IP addresses:

```
cluster::*> system cluster-switch show

Switch                           Type              Address
Model
---------------------------- ----------------- ----------------
---------------
CL1                              cluster-network    10.10.1.101
NX5596
     Serial Number: 01234567
      Is Monitored: true
            Reason:
  Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.1(1)N1(1)
    Version Source: CDP
CL2                              cluster-network    10.10.1.102
NX5596
     Serial Number: 01234568
      Is Monitored: true
            Reason:
  Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.1(1)N1(1)
    Version Source: CDP

2 entries were displayed.
```

4. Set the `-auto-revert` parameter to `false` on cluster LIFs clus1 and clus2 on both nodes:

   `network interface modify`

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node1 -lif clus2 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus2 -auto
-revert false
```

5. Verify that the appropriate RCF and image are installed on the new 3132Q-V switches as necessary for your requirements, and make the essential site customizations, such as users and passwords, network addresses, and so on.

   You must prepare both switches at this time. If you need to upgrade the RCF and image, follow these steps:

   a. Go to the Cisco Ethernet Switches page on the NetApp Support Site.

   b. Note your switch and the required software versions in the table on that page.

   c. Download the appropriate version of the RCF.

   d. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.

   e. Download the appropriate version of the image software.

   See the *ONTAP 8.x or later Cluster and Management Network Switch Reference Configuration Files*Download page, and then click the appropriate version.

   To find the correct version, see the *ONTAP 8.x or later Cluster Network Switch Download page*.

6. Migrate the LIFs associated with the second Nexus 5596 switch to be replaced:

```
network interface migrate
```

Show example

The following example shows n1 and n2, but LIF migration must be done on all of the nodes:

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus2
-source-node n1 –
destination-node n1 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n1_clus3
-source-node n1 –
destination-node n1 -destination-port e0d
cluster::*> network interface migrate -vserver Cluster -lif n2_clus2
-source-node n2 –
destination-node n2 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n2_clus3
-source-node n2 –
destination-node n2 -destination-port e0d
```

7. Verify the cluster's health:

```
network interface show
```

Show example

The following example shows the result of the previous `network interface migrate` command:

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status     Network                Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask           Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24           n1
e0a     true
            n1_clus2   up/up      10.10.0.2/24           n1
e0a     false
            n1_clus3   up/up      10.10.0.3/24           n1
e0d     false
            n1_clus4   up/up      10.10.0.4/24           n1
e0d     true
            n2_clus1   up/up      10.10.0.5/24           n2
e0a     true
            n2_clus2   up/up      10.10.0.6/24           n2
e0a     false
            n2_clus3   up/up      10.10.0.7/24           n2
e0d     false
            n2_clus4   up/up      10.10.0.8/24           n2
e0d     true
8 entries were displayed.
```

8. Shut down the cluster interconnect ports that are physically connected to switch CL2:

`network port modify`

Show example

The following commands shut down the specified ports on n1 and n2, but the ports must be shut down on all nodes:

```
cluster::*> network port modify -node n1 -port e0b -up-admin false
cluster::*> network port modify -node n1 -port e0c -up-admin false
cluster::*> network port modify -node n2 -port e0b -up-admin false
cluster::*> network port modify -node n2 -port e0c -up-admin false
```

9. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8
Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check
```

10. Shut down the ISL ports 41 through 48 on the active Nexus 5596 switch CL1:

Show example

The following example shows how to shut down ISL ports 41 through 48 on the Nexus 5596 switch CL1:

```
(CL1)# configure
(CL1)(Config)# interface e1/41-48
(CL1)(config-if-range)# shutdown
(CL1)(config-if-range)# exit
(CL1)(Config)# exit
(CL1)#
```

If you are replacing a Nexus 5010 or 5020, specify the appropriate port numbers for ISL.

11. Build a temporary ISL between CL1 and C2.

Show example

The following example shows a temporary ISL being set up between CL1 and C2:

```
C2# configure
C2(config)# interface port-channel 2
C2(config-if)# switchport mode trunk
C2(config-if)# spanning-tree port type network
C2(config-if)# mtu 9216
C2(config-if)# interface breakout module 1 port 24 map 10g-4x
C2(config)# interface e1/24/1-4
C2(config-if-range)# switchport mode trunk
C2(config-if-range)# mtu 9216
C2(config-if-range)# channel-group 2 mode active
C2(config-if-range)# exit
C2(config-if)# exit
```

Step 2: Configure ports

1. On all nodes, remove all cables attached to the Nexus 5596 switch CL2.

   With supported cabling, reconnect disconnected ports on all nodes to the Nexus 3132Q-V switch C2.

2. Remove all the cables from the Nexus 5596 switch CL2.

   Attach the appropriate Cisco QSFP to SFP+ break-out cables connecting port 1/24 on the new Cisco 3132Q-V switch, C2, to ports 45 to 48 on existing Nexus 5596, CL1.

3. Verify that interfaces eth1/45-48 already have `channel-group 1 mode active` in their running configuration.

4. Bring up ISLs ports 45 through 48 on the active Nexus 5596 switch CL1.

Show example

The following example shows ISLs ports 45 through 48 being brought up:

```
(CL1)# configure
(CL1)(Config)# interface e1/45-48
(CL1)(config-if-range)# no shutdown
(CL1)(config-if-range)# exit
(CL1)(Config)# exit
(CL1)#
```

5. Verify that the ISLs are `up` on the Nexus 5596 switch CL1:

`show port-channel summary`

Show example

Ports eth1/45 through eth1/48 should indicate (P) meaning that the ISL ports are `up` in the port-channel:

```
Example
CL1# show port-channel summary
Flags: D - Down         P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------------
------------
Group Port-       Type    Protocol  Member Ports
      Channel
--------------------------------------------------------------------------
------------
1     Po1(SU)     Eth     LACP      Eth1/41(D)   Eth1/42(D)
Eth1/43(D)
                                    Eth1/44(D)   Eth1/45(P)
Eth1/46(P)
                                    Eth1/47(P)   Eth1/48(P)
```

6. Verify that the ISLs are `up` on the 3132Q-V switch C2:

```
show port-channel summary
```

Show example

Ports eth1/24/1, eth1/24/2, eth1/24/3, and eth1/24/4 should indicate (P) meaning that the ISL ports are `up` in the port-channel:

```
C2# show port-channel summary
Flags: D - Down          P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-        Type   Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1     Po1(SU)      Eth    LACP      Eth1/31(D)   Eth1/32(D)
2     Po2(SU)      Eth    LACP      Eth1/24/1(P)  Eth1/24/2(P)
Eth1/24/3(P)
                                    Eth1/24/4(P)
```

7. On all nodes, bring up all the cluster interconnect ports connected to the 3132Q-V switch C2:

```
network port modify
```

Show example

The following example shows the specified ports being brought up on nodes n1 and n2:

```
cluster::*> network port modify -node n1 -port e0b -up-admin true
cluster::*> network port modify -node n1 -port e0c -up-admin true
cluster::*> network port modify -node n2 -port e0b -up-admin true
cluster::*> network port modify -node n2 -port e0c -up-admin true
```

8. On all nodes, revert all of the migrated cluster interconnect LIFs connected to C2:

```
network interface revert
```

Show example

The following example shows the migrated cluster LIFs being reverted to their home ports on nodes n1 and n2:

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus2
cluster::*> network interface revert -vserver Cluster -lif n1_clus3
cluster::*> network interface revert -vserver Cluster -lif n2_clus2
cluster::*> network interface revert -vserver Cluster -lif n2_clus3
```

9. Verify all the cluster interconnect ports are now reverted to their home:

```
network interface show
```

Show example

The following example shows that the LIFs on clus2 reverted to their home ports and shows that the LIFs are successfully reverted if the ports in the Current Port column have a status of `true` in the `Is Home` column. If the `Is Home` value is `false`, the LIF has not been reverted.

```
cluster::*> network interface show -role cluster
(network interface show)
          Logical     Status      Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
          n1_clus1   up/up       10.10.0.1/24          n1
e0a     true
          n1_clus2   up/up       10.10.0.2/24          n1
e0b     true
          n1_clus3   up/up       10.10.0.3/24          n1
e0c     true
          n1_clus4   up/up       10.10.0.4/24          n1
e0d     true
          n2_clus1   up/up       10.10.0.5/24          n2
e0a     true
          n2_clus2   up/up       10.10.0.6/24          n2
e0b     true
          n2_clus3   up/up       10.10.0.7/24          n2
e0c     true
          n2_clus4   up/up       10.10.0.8/24          n2
e0d     true
8 entries were displayed.
```

10. Verify that the clustered ports are connected:

```
network port show
```

Show example

The following example shows the result of the previous `network port modify` command, verifying that all the cluster interconnects are `up`:

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1

Ignore
                                                      Speed(Mbps) Health
Health
Port       IPspace       Broadcast Domain Link MTU   Admin/Oper  Status
Status
---------  ------------  ---------------- ---- ----  -----------
--------  ------
e0a        Cluster       Cluster          up   9000  auto/10000  -
-
e0b        Cluster       Cluster          up   9000  auto/10000  -
-
e0c        Cluster       Cluster          up   9000  auto/10000  -
-
e0d        Cluster       Cluster          up   9000  auto/10000  -
-

Node: n2

Ignore
                                                      Speed(Mbps) Health
Health
Port       IPspace       Broadcast Domain Link MTU   Admin/Oper  Status
Status
---------  ------------  ---------------- ---- ----  -----------
--------  ------
e0a        Cluster       Cluster          up   9000   auto/10000 -
-
e0b        Cluster       Cluster          up   9000   auto/10000 -
-
e0c        Cluster       Cluster          up   9000   auto/10000 -
-
e0d        Cluster       Cluster          up   9000   auto/10000 -
-
8 entries were displayed.
```

11. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8
Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check)
```

12. On each node in the cluster, migrate the interfaces associated with the first Nexus 5596 switch, CL1, to be replaced:

```
network interface migrate
```

Show example

The following example shows the ports or LIFs being migrated on nodes n1 and n2:

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus1
-source-node n1 -
destination-node n1 -destination-port e0b
cluster::*> network interface migrate -vserver Cluster -lif n1_clus4
-source-node n1 -
destination-node n1 -destination-port e0c
cluster::*> network interface migrate -vserver Cluster -lif n2_clus1
-source-node n2 -
destination-node n2 -destination-port e0b
cluster::*> network interface migrate -vserver Cluster -lif n2_clus4
-source-node n2 -
destination-node n2 -destination-port e0c
```

13. Verify the cluster status:

```
network interface show
```

Show example

The following example shows that the required cluster LIFs have been migrated to appropriate cluster ports hosted on cluster switch C2:

```
 (network interface show)
          Logical     Status     Network              Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask       Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
          n1_clus1    up/up      10.10.0.1/24        n1
e0b     false
          n1_clus2    up/up      10.10.0.2/24        n1
e0b     true
          n1_clus3    up/up      10.10.0.3/24        n1
e0c     true
          n1_clus4    up/up      10.10.0.4/24        n1
e0c     false
          n2_clus1    up/up      10.10.0.5/24        n2
e0b     false
          n2_clus2    up/up      10.10.0.6/24        n2
e0b     true
          n2_clus3    up/up      10.10.0.7/24        n2
e0c     true
          n2_clus4    up/up      10.10.0.8/24        n2
e0c     false
8 entries were displayed.

----- ------- ----
```

14. On all the nodes, shut down the node ports that are connected to CL1:

```
network port modify
```

Show example

> The following example shows the specified ports being shut down on nodes n1 and n2:
>
> ```
> cluster::*> network port modify -node n1 -port e0a -up-admin false
> cluster::*> network port modify -node n1 -port e0d -up-admin false
> cluster::*> network port modify -node n2 -port e0a -up-admin false
> cluster::*> network port modify -node n2 -port e0d -up-admin false
> ```

15. Shut down the ISL ports 24, 31, and 32 on the active 3132Q-V switch C2:

`shutdown`

Show example

> The following example shows how to shut down ISLs 24, 31, and 32:
>
> ```
> C2# configure
> C2(Config)# interface e1/24/1-4
> C2(config-if-range)# shutdown
> C2(config-if-range)# exit
> C2(config)# interface 1/31-32
> C2(config-if-range)# shutdown
> C2(config-if-range)# exit
> C2(config-if)# exit
> C2#
> ```

16. On all nodes, remove all cables attached to the Nexus 5596 switch CL1.

    With supported cabling, reconnect disconnected ports on all nodes to the Nexus 3132Q-V switch C1.

17. Remove the QSFP breakout cable from Nexus 3132Q-V C2 ports e1/24.

    Connect ports e1/31 and e1/32 on C1 to ports e1/31 and e1/32 on C2 using supported Cisco QSFP optical fiber or direct-attach cables.

18. Restore the configuration on port 24 and remove the temporary Port Channel 2 on C2:

```
C2# configure
C2(config)# no interface breakout module 1 port 24 map 10g-4x
C2(config)# no interface port-channel 2
C2(config-if)# int e1/24
C2(config-if)# description 40GbE Node Port
C2(config-if)# spanning-tree port type edge
C2(config-if)# spanning-tree bpduguard enable
C2(config-if)# mtu 9216
C2(config-if-range)# exit
C2(config)# exit
C2# copy running-config startup-config
[########################################] 100%
Copy Complete.
```

19. Bring up ISL ports 31 and 32 on C2, the active 3132Q-V switch: `no shutdown`

Show example

The following example shows how to bring up ISLs 31 and 32 on the 3132Q-V switch C2:

```
C2# configure
C2(config)# interface ethernet 1/31-32
C2(config-if-range)# no shutdown
C2(config-if-range)# exit
C2(config)# exit
C2# copy running-config startup-config
[########################################] 100%
Copy Complete.
```

Step 3: Verify the configuration

1. Verify that the ISL connections are `up` on the 3132Q-V switch C2:

   `show port-channel summary`

Show example

Ports Eth1/31 and Eth1/32 should indicate `(P)`, meaning that both the ISL ports are `up` in the port-channel:

```
C1# show port-channel summary
Flags: D - Down         P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-        Type    Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1     Po1(SU)      Eth     LACP      Eth1/31(P)   Eth1/32(P)
```

2. On all nodes, bring up all the cluster interconnect ports connected to the new 3132Q-V switch C1:

`network port modify`

Show example

The following example shows all the cluster interconnect ports being brought up for n1 and n2 on the 3132Q-V switch C1:

```
cluster::*> network port modify -node n1 -port e0a -up-admin true
cluster::*> network port modify -node n1 -port e0d -up-admin true
cluster::*> network port modify -node n2 -port e0a -up-admin true
cluster::*> network port modify -node n2 -port e0d -up-admin true
```

3. Verify the status of the cluster node port:

`network port show`

Show example

The following example verifies that all cluster interconnect ports on all nodes on the new 3132Q-V switch C1 are `up`:

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1

Ignore
                                              Speed(Mbps) Health
Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster       Cluster          up   9000 auto/10000  -
-
e0b       Cluster       Cluster          up   9000 auto/10000  -
-
e0c       Cluster       Cluster          up   9000 auto/10000  -
-
e0d       Cluster       Cluster          up   9000 auto/10000  -
-

Node: n2

Ignore
                                              Speed(Mbps) Health
Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster       Cluster          up   9000  auto/10000 -
-
e0b       Cluster       Cluster          up   9000  auto/10000 -
-
e0c       Cluster       Cluster          up   9000  auto/10000 -
-
e0d       Cluster       Cluster          up   9000  auto/10000 -
-
8 entries were displayed.
```

4. On all nodes, revert the specific cluster LIFs to their home ports:

```
network interface revert
```

Show example

The following example shows the specific cluster LIFs being reverted to their home ports on nodes n1 and n2:

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus1
cluster::*> network interface revert -vserver Cluster -lif n1_clus4
cluster::*> network interface revert -vserver Cluster -lif n2_clus1
cluster::*> network interface revert -vserver Cluster -lif n2_clus4
```

5. Verify that the interface is home:

```
network interface show
```

Show example

The following example shows the status of cluster interconnect interfaces is `up` and `Is home` for n1 and n2:

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status     Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24         n1
e0a     true
            n1_clus2   up/up      10.10.0.2/24         n1
e0b     true
            n1_clus3   up/up      10.10.0.3/24         n1
e0c     true
            n1_clus4   up/up      10.10.0.4/24         n1
e0d     true
            n2_clus1   up/up      10.10.0.5/24         n2
e0a     true
            n2_clus2   up/up      10.10.0.6/24         n2
e0b     true
            n2_clus3   up/up      10.10.0.7/24         n2
e0c     true
            n2_clus4   up/up      10.10.0.8/24         n2
e0d     true
8 entries were displayed.
```

6. Ping the remote cluster interfaces and then perform a remote procedure call server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8
Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check)
```

7. Expand the cluster by adding nodes to the Nexus 3132Q-V cluster switches.

8. Display the information about the devices in your configuration:
    - ◦ `network device-discovery show`
    - ◦ `network port show -role cluster`
    - ◦ `network interface show -role cluster`
    - ◦ `system cluster-switch show`

Show example

The following examples show nodes n3 and n4 with 40 GbE cluster ports connected to ports e1/7 and e1/8, respectively on both the Nexus 3132Q-V cluster switches, and both nodes have joined the cluster. The 40 GbE cluster interconnect ports used are e4a and e4e.

```
cluster::> network device-discovery show
            Local  Discovered
Node        Port   Device                 Interface       Platform
----------- ------ ------------------- ----------------
----------------
n1          /cdp
            e0a    C1                     Ethernet1/1/1   N3K-
C3132Q-V
            e0b    C2                     Ethernet1/1/1   N3K-
C3132Q-V
            e0c    C2                     Ethernet1/1/2   N3K-
C3132Q-V
            e0d    C1                     Ethernet1/1/2   N3K-
C3132Q-V
n2          /cdp
            e0a    C1                     Ethernet1/1/3   N3K-
C3132Q-V
            e0b    C2                     Ethernet1/1/3   N3K-
C3132Q-V
            e0c    C2                     Ethernet1/1/4   N3K-
C3132Q-V
            e0d    C1                     Ethernet1/1/4   N3K-
C3132Q-V
n3          /cdp
            e4a    C1                     Ethernet1/7     N3K-
C3132Q-V
            e4e    C2                     Ethernet1/7     N3K-
C3132Q-V
n4          /cdp
            e4a    C1                     Ethernet1/8     N3K-
C3132Q-V
            e4e    C2                     Ethernet1/8     N3K-
C3132Q-V
12 entries were displayed.
```

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1
```

```
Ignore
                                                 Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000 auto/10000  -
-
e0b       Cluster      Cluster          up   9000 auto/10000  -
-
e0c       Cluster      Cluster          up   9000 auto/10000  -
-
e0d       Cluster      Cluster          up   9000 auto/10000  -
-

Node: n2

Ignore
                                                 Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster          up   9000  auto/10000 -
-
e0b       Cluster      Cluster          up   9000  auto/10000 -
-
e0c       Cluster      Cluster          up   9000  auto/10000 -
-
e0d       Cluster      Cluster          up   9000  auto/10000 -
-

Node: n3

Ignore
                                                 Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
```

```
-

Node: n4

Ignore
                                              Speed(Mbps)
Health    Health
Port       IPspace        Broadcast Domain Link MTU   Admin/Oper
Status    Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
-
12 entries were displayed.
```

```
cluster::*> network interface show -role cluster
  (network interface show)
          Logical      Status       Network               Current
Current Is
Vserver     Interface  Admin/Oper  Address/Mask          Node
Port     Home
----------- ---------- ---------- ------------------
------------- ------- ----
Cluster
          n1_clus1   up/up        10.10.0.1/24          n1
e0a      true
          n1_clus2   up/up        10.10.0.2/24          n1
e0b      true
          n1_clus3   up/up        10.10.0.3/24          n1
e0c      true
          n1_clus4   up/up        10.10.0.4/24          n1
e0d      true
          n2_clus1   up/up        10.10.0.5/24          n2
e0a      true
          n2_clus2   up/up        10.10.0.6/24          n2
e0b      true
          n2_clus3   up/up        10.10.0.7/24          n2
e0c      true
          n2_clus4   up/up        10.10.0.8/24          n2
e0d      true
          n3_clus1   up/up        10.10.0.9/24          n3
e4a      true
          n3_clus2   up/up        10.10.0.10/24         n3
e4e      true
          n4_clus1   up/up        10.10.0.11/24         n4
e4a      true
          n4_clus2   up/up        10.10.0.12/24         n4
e4e      true
12 entries were displayed.
```

```
cluster::*> system cluster-switch show

Switch                          Type              Address
Model
--------------------------- ----------------- ----------------
--------------
C1                              cluster-network   10.10.1.103
NX3132V
      Serial Number: FOX000001
       Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.0(3)I4(1)
     Version Source: CDP

C2                              cluster-network   10.10.1.104
NX3132V
      Serial Number: FOX000002
       Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.0(3)I4(1)
     Version Source: CDP

CL1                             cluster-network   10.10.1.101
NX5596
      Serial Number: 01234567
       Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.1(1)N1(1)
     Version Source: CDP
CL2                             cluster-network   10.10.1.102
NX5596
      Serial Number: 01234568
       Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                    7.1(1)N1(1)
     Version Source: CDP

4 entries were displayed.
```

9. Remove the replaced Nexus 5596 if they are not automatically removed:

```
system cluster-switch delete
```

Show example

> The following example shows how to remove the Nexus 5596:
>
> ```
> cluster::> system cluster-switch delete –device CL1
> cluster::> system cluster-switch delete –device CL2
> ```

10. Configure clusters clus1 and clus2 to auto revert on each node and confirm.

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node1 -lif clus2 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus2 -auto
-revert true
```

11. Verify that the proper cluster switches are monitored:

```
system cluster-switch show
```

Show example

```
cluster::> system cluster-switch show

Switch                          Type              Address
Model
-------------------------- ----------------- -----------------
--------------
C1                              cluster-network   10.10.1.103
NX3132V
      Serial Number: FOX000001
       Is Monitored: true
            Reason:
  Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                   7.0(3)I4(1)
    Version Source: CDP

C2                              cluster-network    10.10.1.104
NX3132V
      Serial Number: FOX000002
       Is Monitored: true
            Reason:
  Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                   7.0(3)I4(1)
    Version Source: CDP

2 entries were displayed.
```

12. Enable the cluster switch health monitor log collection feature for collecting switch-related log files:

```
system cluster-switch log setup-password
```

```
system cluster-switch log enable-collection
```

Show example

```
cluster::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
C1
C2

cluster::*> system cluster-switch log setup-password

Enter the switch name: C1
**RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log setup-password

Enter the switch name: C2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster::*>
```

(i)   If any of these commands return an error, contact NetApp support.

13. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Migrate from CN1610 cluster switches to Cisco Nexus 3132Q-V cluster switches

Follow this procedure to replace the existing CN1610 cluster switches with Cisco Nexus 3132Q-V cluster switches.

Review requirements

Review the NetApp CN1610 requirements requirements in Requirements for replacing Cisco Nexus 3132Q-V cluster switches.

For more information, see:

- NetApp CN1601 and CN1610 description page
- Cisco Ethernet Switch description page
- Hardware Universe

Replace the switch

Switch and node nomenclature

The examples in this procedure use the following switch and node nomenclature:

- The command outputs might vary depending on different releases of ONTAP software.
- The CN1610 switches to be replaced are CL1 and CL2.
- The Nexus 3132Q-V switches to replace the CN1610 switches are C1 and C2.
- n1_clus1 is the first cluster logical interface (LIF) that is connected to cluster switch 1 (CL1 or C1) for node n1.
- n1_clus2 is the first cluster LIF that is connected to cluster switch 2 (CL2 or C2) for node n1.
- n1_clus3 is the second LIF that is connected to cluster switch 2 (CL2 or C2) for node n1.
- n1_clus4 is the second LIF that is connected to cluster switch 1 (CL1 or C1) for node n1.
- The nodes are n1, n2, n3, and n4.
- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco® Cluster Network Switch Reference Configuration File Download page.

About the examples

The examples in this procedure use four nodes:

- Two nodes use four 10 GbE cluster interconnect ports: e0a, e0b, e0c, and e0d.
- The other two nodes use two 40/100 GbE cluster interconnect fiber cables: e4a and e4e.

The Hardware Universe has information about the cluster fiber cables on your platforms.

About this task

This procedure covers the following scenario:

- The cluster starts with two nodes connected to two CN1610 cluster switches.
- Cluster switch CL2 to be replaced by C2
  - Traffic on all cluster ports and LIFs on all nodes connected to CL2 are migrated onto the first cluster ports and LIFs connected to CL1.

- Disconnect cabling from all cluster ports on all nodes connected to CL2, and then use supported breakout cabling to reconnect the ports to new cluster switch C2.

- Disconnect cabling between ISL ports CL1 and CL2, and then use supported breakout cabling to reconnect the ports from CL1 to C2.

- Traffic on all cluster ports and LIFs connected to C2 on all nodes is reverted.

- Cluster switch CL1 to be replaced by C1

  - Traffic on all cluster ports and LIFs on all nodes connected to CL1 are migrated onto the second cluster ports and LIFs connected to C2.

  - Disconnect cabling from all cluster ports on all nodes connected to CL1, and then use supported breakout cabling to reconnect the ports to new cluster switch C1.

  - Disconnect cabling between ISL ports CL1 and C2, and then use supported breakout cabling to reconnect the ports from C1 to C2.

  - Traffic on all migrated cluster ports and LIFs connected to C1 on all nodes is reverted.

> The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

Step 1: Prepare for replacement

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

   ```
   system node autosupport invoke -node * -type all - message MAINT=xh
   ```

   x is the duration of the maintenance window in hours.

   > The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Display information about the devices in your configuration:

   ```
   network device-discovery show
   ```

Show example

The following example displays how many cluster interconnect interfaces have been configured in each node for each cluster interconnect switch:

```
cluster::> network device-discovery show

        Local  Discovered
Node    Port   Device        Interface   Platform
------  ------ ------------ ----------- ----------
n1      /cdp
         e0a   CL1           0/1         CN1610
         e0b   CL2           0/1         CN1610
         e0c   CL2           0/2         CN1610
         e0d   CL1           0/2         CN1610
n2      /cdp
         e0a   CL1           0/3         CN1610
         e0b   CL2           0/3         CN1610
         e0c   CL2           0/4         CN1610
         e0d   CL1           0/4         CN1610

8 entries were displayed.
```

3. Determine the administrative or operational status for each cluster interface.

   a. Display the cluster network port attributes:

   ```
   network port show
   ```

The following example displays the network port attributes on a system:

```
cluster::*> network port show -role Cluster
       (network port show)

Node: n1
                Broadcast                Speed (Mbps) Health Ignore
Port  IPspace   Domain     Link  MTU    Admin/Open   Status Health
Status
----- --------- ---------- ----- ----- ------------ ------
-------------
e0a   cluster   cluster    up    9000   auto/10000      -        -
e0b   cluster   cluster    up    9000   auto/10000      -        -
e0c   cluster   cluster    up    9000   auto/10000      -        -
e0d   cluster   cluster    up    9000   auto/10000      -        -

Node: n2
                Broadcast                Speed (Mbps) Health Ignore
Port  IPspace   Domain     Link  MTU    Admin/Open   Status Health
Status
----- --------- ---------- ----- ----- ------------ ------
-------------
e0a   cluster   cluster    up    9000   auto/10000      -        -
e0b   cluster   cluster    up    9000   auto/10000      -        -
e0c   cluster   cluster    up    9000   auto/10000      -        -
e0d   cluster   cluster    up    9000   auto/10000      -        -

8 entries were displayed.
```

b. Display information about the logical interfaces: + `network interface show`

Show example

The following example displays the general information about all of the LIFs on your system:

```
cluster::*> network interface show -role Cluster
        (network interface show)

          Logical     Status       Network         Current  Current
Is
Vserver   Interface   Admin/Oper   Address/Mask    Node     Port
Home
--------  ----------  -----------  --------------  -------- --------
-----
Cluster
          n1_clus1    up/up        10.10.0.1/24    n1       e0a
true
          n1_clus2    up/up        10.10.0.2/24    n1       e0b
true
          n1_clus3    up/up        10.10.0.3/24    n1       e0c
true
          n1_clus4    up/up        10.10.0.4/24    n1       e0d
true
          n2_clus1    up/up        10.10.0.5/24    n2       e0a
true
          n2_clus2    up/up        10.10.0.6/24    n2       e0b
true
          n2_clus3    up/up        10.10.0.7/24    n2       e0c
true
          n2_clus4    up/up        10.10.0.8/24    n2       e0d
true

8 entries were displayed.
```

c. Display information about the discovered cluster switches:

```
system cluster-switch show
```

Show example

The following example displays the cluster switches that are known to the cluster, along with their management IP addresses:

```
cluster::> system cluster-switch show

Switch                              Type            Address
Model
---------------------------- --------------- -------------
--------
CL1                                 cluster-network  10.10.1.101
CN1610
        Serial Number: 01234567
         Is Monitored: true
               Reason:
     Software Version: 1.2.0.7
       Version Source: ISDP

CL2                                 cluster-network  10.10.1.102
CN1610
        Serial Number: 01234568
         Is Monitored: true
               Reason:
     Software Version: 1.2.0.7
       Version Source: ISDP

2 entries were displayed.
```

4. Set the `-auto-revert` parameter to false on cluster LIFs clus1 and clus4 on both nodes:

```
network interface modify
```

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node1 -lif clus4 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus4 -auto
-revert false
```

5. Verify that the appropriate RCF and image are installed on the new 3132Q-V switches as necessary for your requirements, and make any essential site customizations, such as users and passwords, network addresses, and so on.

   You must prepare both switches at this time. If you need to upgrade the RCF and image, follow these steps:

   a. See the Cisco Ethernet Switches page on NetApp Support Site.

   b. Note your switch and the required software versions in the table on that page.

   c. Download the appropriate version of the RCF.

   d. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.

   e. Download the appropriate version of the image software.

   Cisco® Cluster and Management Network Switch Reference Configuration File Download

6. Migrate the LIFs associated with the second CN1610 switch to be replaced:

   ```
   network interface migrate
   ```

   > (i) You must migrate the cluster LIFs from a connection to the node, either through the service processor or node management interface, which owns the cluster LIF being migrated.

   Show example

   The following example shows n1 and n2, but LIF migration must be done on all the nodes:

   ```
   cluster::*> network interface migrate -vserver Cluster -lif n1_clus2
   -destination-node  n1  -destination-port  e0a
   cluster::*> network interface migrate -vserver Cluster -lif n1_clus3
   -destination-node  n1  -destination-port  e0d
   cluster::*> network interface migrate -vserver Cluster -lif n2_clus2
   -destination-node  n2  -destination-port  e0a
   cluster::*> network interface migrate -vserver Cluster -lif n2_clus3
   -destination-node  n2  -destination-port  e0d
   ```

7. Verify the cluster's health:

   ```
   network interface show
   ```

Show example

The following example shows the result of the previous `network interface migrate` command:

```
cluster::*> network interface show -role Cluster
        (network interface show)

        Logical     Status      Network          Current  Current  Is
Vserver  Interface  Admin/Oper  Address/Mask     Node     Port
Home
-------- ---------- ----------- --------------- -------- --------
-----
Cluster
        n1_clus1   up/up        10.10.0.1/24     n1       e0a
true
        n1_clus2   up/up        10.10.0.2/24     n1       e0a
false
        n1_clus3   up/up        10.10.0.3/24     n1       e0d
false
        n1_clus4   up/up        10.10.0.4/24     n1       e0d
true
        n2_clus1   up/up        10.10.0.5/24     n2       e0a
true
        n2_clus2   up/up        10.10.0.6/24     n2       e0a
false
        n2_clus3   up/up        10.10.0.7/24     n2       e0d
false
        n2_clus4   up/up        10.10.0.8/24     n2       e0d
true

8 entries were displayed.
```

8. Shut down the cluster interconnect ports that are physically connected to switch CL2:

```
network port modify
```

Show example

The following commands shut down the specified ports on n1 and n2, but the ports must be shut down on all nodes:

```
cluster::*> network port modify -node n1 -port e0b -up-admin false
cluster::*> network port modify -node n1 -port e0c -up-admin false
cluster::*> network port modify -node n2 -port e0b -up-admin false
cluster::*> network port modify -node n2 -port e0c -up-admin false
```

9. Ping the remote cluster interfaces, and then perform a remote procedure call server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1        e0a    10.10.0.1
Cluster n1_clus2 n1        e0b    10.10.0.2
Cluster n1_clus3 n1        e0c    10.10.0.3
Cluster n1_clus4 n1        e0d    10.10.0.4
Cluster n2_clus1 n2        e0a    10.10.0.5
Cluster n2_clus2 n2        e0b    10.10.0.6
Cluster n2_clus3 n2        e0c    10.10.0.7
Cluster n2_clus4 n2        e0d    10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8

Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check)
```

10. Shut down the ISL ports 13 through 16 on the active CN1610 switch CL1:

```
shutdown
```

Show example

The following example shows how to shut down ISL ports 13 through 16 on the CN1610 switch CL1:

```
(CL1)# configure
(CL1)(Config)# interface 0/13-0/16
(CL1)(Interface 0/13-0/16)# shutdown
(CL1)(Interface 0/13-0/16)# exit
(CL1)(Config)# exit
(CL1)#
```

11. Build a temporary ISL between CL1 and C2:

Show example

The following example builds a temporary ISL between CL1 (ports 13-16) and C2 (ports e1/24/1-4):

```
C2# configure
C2(config)# interface port-channel 2
C2(config-if)# switchport mode trunk
C2(config-if)# spanning-tree port type network
C2(config-if)# mtu 9216
C2(config-if)# interface breakout module 1 port 24 map 10g-4x
C2(config)# interface e1/24/1-4
C2(config-if-range)# switchport mode trunk
C2(config-if-range)# mtu 9216
C2(config-if-range)# channel-group 2 mode active
C2(config-if-range)# exit
C2(config-if)# exit
```

Step 2: Configure ports

1. On all nodes, remove the cables that are attached to the CN1610 switch CL2.

   With supported cabling, you must reconnect the disconnected ports on all of the nodes to the Nexus 3132Q-V switch C2.

2. Remove four ISL cables from ports 13 to 16 on the CN1610 switch CL1.

   You must attach appropriate Cisco QSFP to SFP+ breakout cables connecting port 1/24 on the new Cisco 3132Q-V switch C2, to ports 13 to 16 on existing CN1610 switch CL1.

> ⓘ When reconnecting any cables to the new Cisco 3132Q-V switch, you must use either optical fiber or Cisco twinax cables.

3. To make the ISL dynamic, configure the ISL interface 3/1 on the active CN1610 switch to disable the static mode: `no port-channel static`

   This configuration matches with the ISL configuration on the 3132Q-V switch C2 when the ISLs are brought up on both switches in step 11

   Show example

   The following example shows the configuration of the ISL interface 3/1 using the `no port-channel static` command to make the ISL dynamic:

   ```
   (CL1)# configure
   (CL1)(Config)# interface 3/1
   (CL1)(Interface 3/1)# no port-channel static
   (CL1)(Interface 3/1)# exit
   (CL1)(Config)# exit
   (CL1)#
   ```

4. Bring up ISLs 13 through 16 on the active CN1610 switch CL1.

   Show example

   The following example illustrates the process of bringing up ISL ports 13 through 16 on the port-channel interface 3/1:

   ```
   (CL1)# configure
   (CL1)(Config)# interface 0/13-0/16,3/1
   (CL1)(Interface 0/13-0/16,3/1)# no shutdown
   (CL1)(Interface 0/13-0/16,3/1)# exit
   (CL1)(Config)# exit
   (CL1)#
   ```

5. Verify that the ISLs are `up` on the CN1610 switch CL1:

   `show port-channel`

   The "Link State" should be `Up`, "Type" should be `Dynamic`, and the "Port Active" column should be `True` for ports 0/13 to 0/16:

Show example

```
(CL1)# show port-channel 3/1
Local Interface................................ 3/1
Channel Name................................... ISL-LAG
Link State..................................... Up
Admin Mode..................................... Enabled
Type........................................... Dynamic
Load Balance Option............................ 7
(Enhanced hashing mode)

Mbr     Device/         Port        Port
Ports   Timeout         Speed       Active
------  -------------   ----------  -------
0/13    actor/long      10 Gb Full  True
        partner/long
0/14    actor/long      10 Gb Full  True
        partner/long
0/15    actor/long      10 Gb Full  True
        partner/long
0/16    actor/long      10 Gb Full  True
        partner/long
```

6. Verify that the ISLs are `up` on the 3132Q-V switch C2:

   `show port-channel summary`

Show example

Ports Eth1/24/1 through Eth1/24/4 should indicate `(P)`, meaning that all four ISL ports are up in the port-channel. Eth1/31 and Eth1/32 should indicate `(D)` as they are not connected:

```
C2# show port-channel summary

Flags:  D - Down         P - Up in port-channel (members)
        I - Individual   H - Hot-standby (LACP only)
        s - Suspended    r - Module-removed
        S - Switched     R - Routed
        U - Up (port-channel)
        M - Not in use. Min-links not met
--------------------------------------------------------------------
----------
Group Port-       Type      Protocol  Member Ports
      Channel
--------------------------------------------------------------------
----------
1     Po1(SU)     Eth       LACP      Eth1/31(D)   Eth1/32(D)
2     Po2(SU)     Eth       LACP      Eth1/24/1(P) Eth1/24/2(P)
Eth1/24/3(P)
                                      Eth1/24/4(P)
```

7. Bring up all of the cluster interconnect ports that are connected to the 3132Q-V switch C2 on all of the nodes:

```
network port modify
```

Show example

The following example shows how to bring up the cluster interconnect ports connected to the 3132Q-V switch C2:

```
cluster::*> network port modify -node n1 -port e0b -up-admin true
cluster::*> network port modify -node n1 -port e0c -up-admin true
cluster::*> network port modify -node n2 -port e0b -up-admin true
cluster::*> network port modify -node n2 -port e0c -up-admin true
```

8. Revert all of the migrated cluster interconnect LIFs that are connected to C2 on all of the nodes:

```
network interface revert
```

Show example

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus2
cluster::*> network interface revert -vserver Cluster -lif n1_clus3
cluster::*> network interface revert -vserver Cluster -lif n2_clus2
cluster::*> network interface revert -vserver Cluster -lif n2_clus3
```

9. Verify that all of the cluster interconnect ports are reverted to their home ports:

```
network interface show
```

Show example

The following example shows that the LIFs on clus2 are reverted to their home ports, and shows that the LIFs are successfully reverted if the ports in the "Current Port" column have a status of `true` in the "Is Home" column. If the Is Home value is `false`, then the LIF is not reverted.

```
cluster::*> network interface show -role cluster
        (network interface show)

        Logical    Status      Network         Current  Current  Is
Vserver Interface  Admin/Oper  Address/Mask    Node     Port
Home
------- ---------- ----------- --------------- -------- --------
-----
Cluster
        n1_clus1   up/up       10.10.0.1/24    n1       e0a
true
        n1_clus2   up/up       10.10.0.2/24    n1       e0b
true
        n1_clus3   up/up       10.10.0.3/24    n1       e0c
true
        n1_clus4   up/up       10.10.0.4/24    n1       e0d
true
        n2_clus1   up/up       10.10.0.5/24    n2       e0a
true
        n2_clus2   up/up       10.10.0.6/24    n2       e0b
true
        n2_clus3   up/up       10.10.0.7/24    n2       e0c
true
        n2_clus4   up/up       10.10.0.8/24    n2       e0d
true

8 entries were displayed.
```

10. Verify that all of the cluster ports are connected:

`network port show`

Show example

The following example shows the result of the previous `network port modify` command, verifying that all of the cluster interconnects are `up`:

```
cluster::*> network port show -role Cluster
        (network port show)

Node: n1
                Broadcast                   Speed (Mbps) Health  Ignore
Port  IPspace   Domain      Link  MTU  Admin/Open   Status   Health
Status
----- --------- ----------- ----- ----- ------------ --------
-------------
e0a   cluster   cluster     up    9000  auto/10000     -        -
e0b   cluster   cluster     up    9000  auto/10000     -        -
e0c   cluster   cluster     up    9000  auto/10000     -        -
e0d   cluster   cluster     up    9000  auto/10000     -        -

Node: n2
                Broadcast                   Speed (Mbps) Health  Ignore
Port  IPspace   Domain      Link  MTU  Admin/Open   Status   Health
Status
----- --------- ----------- ----- ----- ------------ --------
-------------
e0a   cluster   cluster     up    9000  auto/10000     -        -
e0b   cluster   cluster     up    9000  auto/10000     -        -
e0c   cluster   cluster     up    9000  auto/10000     -        -
e0d   cluster   cluster     up    9000  auto/10000     -        -

8 entries were displayed.
```

11. Ping the remote cluster interfaces and then perform a remote procedure call server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1        e0a    10.10.0.1
Cluster n1_clus2 n1        e0b    10.10.0.2
Cluster n1_clus3 n1        e0c    10.10.0.3
Cluster n1_clus4 n1        e0d    10.10.0.4
Cluster n2_clus1 n2        e0a    10.10.0.5
Cluster n2_clus2 n2        e0b    10.10.0.6
Cluster n2_clus3 n2        e0c    10.10.0.7
Cluster n2_clus4 n2        e0d    10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8

Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check)
```

12. On each node in the cluster, migrate the interfaces that are associated with the first CN1610 switch CL1, to

be replaced:

```
network interface migrate
```

Show example

The following example shows the ports or LIFs being migrated on nodes n1 and n2:

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus1
-destination-node n1 -destination-port e0b
cluster::*> network interface migrate -vserver Cluster -lif n1_clus4
-destination-node n1 -destination-port e0c
cluster::*> network interface migrate -vserver Cluster -lif n2_clus1
-destination-node n2 -destination-port e0b
cluster::*> network interface migrate -vserver Cluster -lif n2_clus4
-destination-node n2 -destination-port e0c
```

13. Verify the cluster status:

```
network interface show
```

Show example

The following example shows that the required cluster LIFs have been migrated to the appropriate cluster ports hosted on cluster switch C2:

```
cluster::*> network interface show -role Cluster
        (network interface show)

         Logical    Status      Network         Current  Current  Is
Vserver  Interface  Admin/Oper  Address/Mask    Node     Port
Home
-------- ---------- ----------- --------------- -------- --------
-----
Cluster
         n1_clus1   up/up       10.10.0.1/24    n1       e0b
false
         n1_clus2   up/up       10.10.0.2/24    n1       e0b
true
         n1_clus3   up/up       10.10.0.3/24    n1       e0c
true
         n1_clus4   up/up       10.10.0.4/24    n1       e0c
false
         n2_clus1   up/up       10.10.0.5/24    n2       e0b
false
         n2_clus2   up/up       10.10.0.6/24    n2       e0b
true
         n2_clus3   up/up       10.10.0.7/24    n2       e0c
true
         n2_clus4   up/up       10.10.0.8/24    n2       e0c
false

8 entries were displayed.
```

14. Shut down the node ports that are connected to CL1 on all of the nodes:

```
network port modify
```

Show example

The following example shows how to shut down the specified ports on nodes n1 and n2:

```
cluster::*> network port modify -node n1 -port e0a -up-admin false
cluster::*> network port modify -node n1 -port e0d -up-admin false
cluster::*> network port modify -node n2 -port e0a -up-admin false
cluster::*> network port modify -node n2 -port e0d -up-admin false
```

15. Shut down the ISL ports 24, 31, and 32 on the active 3132Q-V switch C2:

    `shutdown`

    Show example

    The following example shows how to shut down ISLs 24, 31, and 32 on the active 3132Q-V switch C2:

    ```
    C2# configure
    C2(config)# interface ethernet 1/24/1-4
    C2(config-if-range)# shutdown
    C2(config-if-range)# exit
    C2(config)# interface ethernet 1/31-32
    C2(config-if-range)# shutdown
    C2(config-if-range)# exit
    C2(config)# exit
    C2#
    ```

16. Remove the cables that are attached to the CN1610 switch CL1 on all of the nodes.

    With supported cabling, you must reconnect the disconnected ports on all of the nodes to the Nexus 3132Q-V switch C1.

17. Remove the QSFP cables from Nexus 3132Q-V C2 port e1/24.

    You must connect ports e1/31 and e1/32 on C1 to ports e1/31 and e1/32 on C2 using supported Cisco QSFP optical fiber or direct-attach cables.

18. Restore the configuration on port 24 and remove the temporary port-channel 2 on C2, by copying the `running-configuration` file to the `startup-configuration` file.

The following example copies the `running-configuration` file to the `startup-configuration` file:

```
C2# configure
C2(config)# no interface breakout module 1 port 24 map 10g-4x
C2(config)# no interface port-channel 2
C2(config-if)# interface e1/24
C2(config-if)# description 40GbE Node Port
C2(config-if)# spanning-tree port type edge
C2(config-if)# spanning-tree bpduguard enable
C2(config-if)# mtu 9216
C2(config-if-range)# exit
C2(config)# exit
C2# copy running-config startup-config
[########################################] 100%
Copy Complete.
```

19. Bring up ISL ports 31 and 32 on C2, the active 3132Q-V switch:

```
no shutdown
```

The following example shows how to bring up ISLs 31 and 32 on the 3132Q-V switch C2:

```
C2# configure
C2(config)# interface ethernet 1/31-32
C2(config-if-range)# no shutdown
C2(config-if-range)# exit
C2(config)# exit
C2# copy running-config startup-config
[########################################] 100%
Copy Complete.
```

Step 3: Verify the configuration

1. Verify that the ISL connections are `up` on the 3132Q-V switch C2:

```
show port-channel summary
```

Ports Eth1/31 and Eth1/32 should indicate `(P)`, meaning that both the ISL ports are `up` in the port-channel.

Show example

```
C1# show port-channel summary
Flags:  D - Down         P - Up in port-channel (members)
        I - Individual  H - Hot-standby (LACP only)
        s - Suspended   r - Module-removed
        S - Switched    R - Routed
        U - Up (port-channel)
        M - Not in use. Min-links not met
--------------------------------------------------------------------
----------
Group Port-       Type     Protocol  Member Ports
      Channel
--------------------------------------------------------------------
----------
1     Po1(SU)     Eth      LACP      Eth1/31(P)   Eth1/32(P)
```

2. Bring up all of the cluster interconnect ports connected to the new 3132Q-V switch C1 on all of the nodes:

`network port modify`

Show example

The following example shows how to bring up all of the cluster interconnect ports connected to the new 3132Q-V switch C1:

```
cluster::*> network port modify -node n1 -port e0a -up-admin true
cluster::*> network port modify -node n1 -port e0d -up-admin true
cluster::*> network port modify -node n2 -port e0a -up-admin true
cluster::*> network port modify -node n2 -port e0d -up-admin true
```

3. Verify the status of the cluster node port:

`network port show`

Show example

The following example verifies that all of the cluster interconnect ports on n1 and n2 on the new 3132Q-V switch C1 are `up`:

```
cluster::*> network port show -role Cluster
        (network port show)

Node: n1
                Broadcast                Speed (Mbps) Health   Ignore
Port  IPspace   Domain      Link  MTU    Admin/Open   Status   Health
Status
----- --------- ---------- ----- ----- ------------ --------
-------------
e0a   cluster   cluster     up    9000   auto/10000   -        -
e0b   cluster   cluster     up    9000   auto/10000   -        -
e0c   cluster   cluster     up    9000   auto/10000   -        -
e0d   cluster   cluster     up    9000   auto/10000   -        -

Node: n2
                Broadcast                Speed (Mbps) Health   Ignore
Port  IPspace   Domain      Link  MTU    Admin/Open   Status   Health
Status
----- --------- ---------- ----- ----- ------------ --------
-------------
e0a   cluster   cluster     up    9000   auto/10000   -        -
e0b   cluster   cluster     up    9000   auto/10000   -        -
e0c   cluster   cluster     up    9000   auto/10000   -        -
e0d   cluster   cluster     up    9000   auto/10000   -        -

8 entries were displayed.
```

4. Revert all of the migrated cluster interconnect LIFs that were originally connected to C1 on all of the nodes:

```
network interface revert
```

Show example

The following example shows how to revert the migrated cluster LIFs to their home ports:

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus1
cluster::*> network interface revert -vserver Cluster -lif n1_clus4
cluster::*> network interface revert -vserver Cluster -lif n2_clus1
cluster::*> network interface revert -vserver Cluster -lif n2_clus4
```

5. Verify that the interface is now home:

```
network interface show
```

Show example

The following example shows the status of cluster interconnect interfaces is `up` and `Is home` for n1 and n2:

```
cluster::*> network interface show -role Cluster
        (network interface show)

          Logical    Status      Network        Current  Current  Is
  Vserver Interface  Admin/Oper  Address/Mask   Node     Port
  Home
  -------- ---------- ----------- -------------- -------- --------
  -----
  Cluster
          n1_clus1   up/up       10.10.0.1/24   n1       e0a
  true
          n1_clus2   up/up       10.10.0.2/24   n1       e0b
  true
          n1_clus3   up/up       10.10.0.3/24   n1       e0c
  true
          n1_clus4   up/up       10.10.0.4/24   n1       e0d
  true
          n2_clus1   up/up       10.10.0.5/24   n2       e0a
  true
          n2_clus2   up/up       10.10.0.6/24   n2       e0b
  true
          n2_clus3   up/up       10.10.0.7/24   n2       e0c
  true
          n2_clus4   up/up       10.10.0.8/24   n2       e0d
  true

  8 entries were displayed.
```

6. Ping the remote cluster interfaces and then perform a remote procedure call server check:

```
cluster ping-cluster
```

Show example

The following example shows how to ping the remote cluster interfaces:

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1          e0a    10.10.0.1
Cluster n1_clus2 n1          e0b    10.10.0.2
Cluster n1_clus3 n1          e0c    10.10.0.3
Cluster n1_clus4 n1          e0d    10.10.0.4
Cluster n2_clus1 n2          e0a    10.10.0.5
Cluster n2_clus2 n2          e0b    10.10.0.6
Cluster n2_clus3 n2          e0c    10.10.0.7
Cluster n2_clus4 n2          e0d    10.10.0.8

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 16 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 16 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
    Local 10.10.0.3 to Remote 10.10.0.7
    Local 10.10.0.3 to Remote 10.10.0.8
    Local 10.10.0.4 to Remote 10.10.0.5
    Local 10.10.0.4 to Remote 10.10.0.6
    Local 10.10.0.4 to Remote 10.10.0.7
    Local 10.10.0.4 to Remote 10.10.0.8

Larger than PMTU communication succeeds on 16 path(s)
RPC status:
4 paths up, 0 paths down (tcp check)
4 paths up, 0 paths down (udp check)
```

7. Expand the cluster by adding nodes to the Nexus 3132Q-V cluster switches.

8. Display the information about the devices in your configuration:

    ° `network device-discovery show`

    ° `network port show -role cluster`

    ° `network interface show -role cluster`

    ° `system cluster-switch show`

Show example

The following examples show nodes n3 and n4 with 40 GbE cluster ports connected to ports e1/7 and e1/8, respectively on both the Nexus 3132Q-V cluster switches, and both nodes have joined the cluster. The 40 GbE cluster interconnect ports used are e4a and e4e.

```
cluster::*> network device-discovery show

        Local  Discovered
Node    Port   Device         Interface       Platform
------  ------ ------------   --------------  -------------
n1      /cdp
         e0a   C1             Ethernet1/1/1   N3K-C3132Q-V
         e0b   C2             Ethernet1/1/1   N3K-C3132Q-V
         e0c   C2             Ethernet1/1/2   N3K-C3132Q-V
         e0d   C1             Ethernet1/1/2   N3K-C3132Q-V
n2      /cdp
         e0a   C1             Ethernet1/1/3   N3K-C3132Q-V
         e0b   C2             Ethernet1/1/3   N3K-C3132Q-V
         e0c   C2             Ethernet1/1/4   N3K-C3132Q-V
         e0d   C1             Ethernet1/1/4   N3K-C3132Q-V
n3      /cdp
         e4a   C1             Ethernet1/7      N3K-C3132Q-V
         e4e   C2             Ethernet1/7      N3K-C3132Q-V
n4      /cdp
         e4a   C1             Ethernet1/8      N3K-C3132Q-V
         e4e   C2             Ethernet1/8      N3K-C3132Q-V

12 entries were displayed.
```

```
cluster::*> network port show -role cluster
        (network port show)

Node: n1
                Broadcast                     Speed (Mbps) Health
Ignore
Port  IPspace   Domain      Link  MTU   Admin/Open    Status
Health Status
----- --------- ---------- ----- ----- ----------- --------
-------------
e0a   cluster   cluster     up    9000  auto/10000    -         -
e0b   cluster   cluster     up    9000  auto/10000    -         -
e0c   cluster   cluster     up    9000  auto/10000    -         -
e0d   cluster   cluster     up    9000  auto/10000    -         -
```

```
Node: n2
                  Broadcast                   Speed (Mbps) Health
Ignore
Port  IPspace    Domain      Link  MTU   Admin/Open    Status
Health Status
----- ---------  ----------  ----- ----- ------------  --------
------------
e0a   cluster    cluster     up    9000  auto/10000    -          -
e0b   cluster    cluster     up    9000  auto/10000    -          -
e0c   cluster    cluster     up    9000  auto/10000    -          -
e0d   cluster    cluster     up    9000  auto/10000    -          -

Node: n3
                  Broadcast                   Speed (Mbps) Health
Ignore
Port  IPspace    Domain      Link  MTU   Admin/Open    Status
Health Status
----- ---------  ----------  ----- ----- ------------  --------
------------
e4a   cluster    cluster     up    9000  auto/40000    -          -
e4e   cluster    cluster     up    9000  auto/40000    -          -

Node: n4
                  Broadcast                   Speed (Mbps) Health
Ignore
Port  IPspace    Domain      Link  MTU   Admin/Open    Status
Health Status
----- ---------  ----------  ----- ----- ------------  --------
------------
e4a   cluster    cluster     up    9000  auto/40000    -          -
e4e   cluster    cluster     up    9000  auto/40000    -          -

12 entries were displayed.
```

```
cluster::*> network interface show -role Cluster
        (network interface show)

          Logical     Status      Network         Current   Current
Is
Vserver   Interface   Admin/Oper  Address/Mask    Node      Port
Home
--------  ----------  ----------- --------------  --------  --------
-----
Cluster
          n1_clus1    up/up       10.10.0.1/24    n1        e0a
true
          n1_clus2    up/up       10.10.0.2/24    n1        e0b
true
          n1_clus3    up/up       10.10.0.3/24    n1        e0c
true
          n1_clus4    up/up       10.10.0.4/24    n1        e0d
true
          n2_clus1    up/up       10.10.0.5/24    n2        e0a
true
          n2_clus2    up/up       10.10.0.6/24    n2        e0b
true
          n2_clus3    up/up       10.10.0.7/24    n2        e0c
true
          n2_clus4    up/up       10.10.0.8/24    n2        e0d
true
          n3_clus1    up/up       10.10.0.9/24    n3        e4a
true
          n3_clus2    up/up       10.10.0.10/24   n3        e4e
true
          n4_clus1    up/up       10.10.0.11/24   n4        e4a
true
          n4_clus2    up/up       10.10.0.12/24   n4        e4e
true

12 entries were displayed.
```

```
cluster::> system cluster-switch show

Switch                          Type             Address         Model
--------------------------  ----------------  --------------
---------
C1                              cluster-network   10.10.1.103
NX3132V
       Serial Number: FOX000001
         Is Monitored: true
               Reason:
     Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                        7.0(3)I4(1)
      Version Source: CDP

C2                              cluster-network   10.10.1.104
NX3132V
       Serial Number: FOX000002
         Is Monitored: true
               Reason:
     Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                        7.0(3)I4(1)
      Version Source: CDP

CL1                             cluster-network   10.10.1.101    CN1610
       Serial Number: 01234567
         Is Monitored: true
               Reason:
     Software Version: 1.2.0.7
      Version Source: ISDP

CL2                             cluster-network   10.10.1.102
CN1610
       Serial Number: 01234568
         Is Monitored: true
               Reason:
     Software Version: 1.2.0.7
      Version Source: ISDP

4 entries were displayed.
```

9. Remove the replaced CN1610 switches if they are not automatically removed:

```
system cluster-switch delete
```

Show example

> The following example shows how to remove the CN1610 switches:
>
> ```
> cluster::> system cluster-switch delete -device CL1
> cluster::> system cluster-switch delete -device CL2
> ```

10. Configure clusters clus1 and clus4 to `-auto-revert` on each node and confirm:

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node1 -lif clus4 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus4 -auto
-revert true
```

11. Verify that the proper cluster switches are monitored:

```
system cluster-switch show
```

Show example

```
cluster::> system cluster-switch show

Switch                          Type              Address
Model
--------------------------- ----------------- ----------------
--------------
C1                              cluster-network   10.10.1.103
NX3132V
      Serial Number: FOX000001
        Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                    7.0(3)I4(1)
     Version Source: CDP

C2                              cluster-network   10.10.1.104
NX3132V
      Serial Number: FOX000002
        Is Monitored: true
             Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                    7.0(3)I4(1)
     Version Source: CDP

2 entries were displayed.
```

12. Enable the cluster switch health monitor log collection feature for collecting switch-related log files:

```
system cluster-switch log setup-password
```

```
system cluster-switch log enable-collection
```

Show example

```
cluster::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
C1
C2

cluster::*> system cluster-switch log setup-password

Enter the switch name: C1
RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log setup-password

Enter the switch name: C2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster::*>
```

ⓘ   If any of these commands return an error, contact NetApp support.

13. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Migrate from a switchless cluster to a two-node switched cluster

If you have a two-node switchless cluster, you can follow this procedure to migrate to a two-node switched cluster that includes Cisco Nexus 3132Q-V cluster network switches. The replacement procedure is a nondisruptive procedure (NDO).

Review requirements

Ports and node connections

Make sure you understand the port and node connections and cabling requirements when you migrate to a two-node switched cluster with Cisco Nexus 3132Q-V cluster switches.

- The cluster switches use the Inter-Switch Link (ISL) ports e1/31-32.

- The Hardware Universe contains information about supported cabling to Nexus 3132Q-V switches:

  - The nodes with 10 GbE cluster connections require QSFP optical modules with breakout fiber cables or QSFP to SFP+ copper break-out cables.

  - The nodes with 40/100 GbE cluster connections require supported QSFP/QSFP28 optical modules with fiber cables or QSFP/QSFP28 copper direct-attach cables.

  - The cluster switches use the appropriate ISL cabling: 2x QSFP28 fiber or copper direct-attach cables.

- On Nexus 3132Q-V, you can operate QSFP ports as either 40/100 Gb Ethernet or 4 x10 Gb Ethernet modes.

  By default, there are 32 ports in the 40/100 Gb Ethernet mode. These 40 Gb Ethernet ports are numbered in a 2-tuple naming convention. For example, the second 40 Gb Ethernet port is numbered as 1/2. The process of changing the configuration from 40 Gb Ethernet to 10 Gb Ethernet is called *breakout* and the process of changing the configuration from 10 Gb Ethernet to 40 Gb Ethernet is called *breakin*. When you break out a 40/100 Gb Ethernet port into 10 Gb Ethernet ports, the resulting ports are numbered using a 3-tuple naming convention. For example, the breakout ports of the second 40/100 Gb Ethernet port are numbered as 1/2/1, 1/2/2, 1/2/3, 1/2/4.

- On the left side of Nexus 3132Q-V is a set of four SFP+ ports multiplexed to the first QSFP port.

  By default, the RCF is structured to use the first QSFP port.

  You can make four SFP+ ports active instead of a QSFP port for Nexus 3132Q-V by using the `hardware profile front portmode sfp-plus` command. Similarly, you can reset Nexus 3132Q-V to use a QSFP port instead of four SFP+ ports by using the `hardware profile front portmode qsfp` command.

- Make sure you configured some of the ports on Nexus 3132Q-V to run at 10 GbE or 40/100 GbE.

  You can break-out the first six ports into 4x10 GbE mode by using the `interface breakout module 1 port 1-6 map 10g-4x` command. Similarly, you can regroup the first six QSFP+ ports from breakout configuration by using the `no interface breakout module 1 port 1-6 map 10g-4x` command.

- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco ® Cluster Network Switch Reference Configuration File Download page.

What you'll need

- Configurations properly set up and functioning.

- Nodes running ONTAP 9.4 or later.

- All cluster ports in the `up` state.

- The Cisco Nexus 3132Q-V cluster switch is supported.

- The existing cluster network configuration has:

    ◦ The Nexus 3132 cluster infrastructure that is redundant and fully functional on both switches.

    ◦ The latest RCF and NX-OS versions on your switches.

    The Cisco Ethernet Switches page has information about the ONTAP and NX-OS versions supported in this procedure.

    ◦ Management connectivity on both switches.

    ◦ Console access to both switches.

    ◦ All cluster logical interfaces (LIFs) in the `up` state without being migrated.

    ◦ Initial customization of the switch.

    ◦ All the ISL ports enabled and cabled.

In addition, you must plan, migrate, and read the required documentation on 10 GbE and 40/100 GbE connectivity from nodes to Nexus 3132Q-V cluster switches.

Migrate the switches

About the examples

The examples in this procedure use the following switch and node nomenclature:

- Nexus 3132Q-V cluster switches, C1 and C2.

- The nodes are n1 and n2.

> ⓘ The examples in this procedure use two nodes, each utilizing two 40/100 GbE cluster interconnect ports e4a and e4e. The Hardware Universe has details about the cluster ports on your platforms.

About this task

This procedure covers the following scenarios:

- n1_clus1 is the first cluster logical interface (LIF) to be connected to cluster switch C1 for node n1.

- n1_clus2 is the first cluster LIF to be connected to cluster switch C2 for node n1.

- n2_clus1 is the first cluster LIF to be connected to cluster switch C1 for node n2.

- n2_clus2 is the second cluster LIF to be connected to cluster switch C2 for node n2.

- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco ® Cluster Network Switch Reference Configuration File Download page.

> ⓘ The procedure requires the use of both ONTAP commands and Cisco Nexus 3000 Series Switches commands; ONTAP commands are used unless otherwise indicated.

- The cluster starts with two nodes connected and functioning in a two-node switchless cluster setting.

- The first cluster port is moved to C1.

- The second cluster port is moved to C2.

- The two-node switchless cluster option is disabled.

Step 1: Prepare for migration

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

   ```
   system node autosupport invoke -node * -type all - message MAINT=xh
   ```

   x is the duration of the maintenance window in hours.

   > (i) The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Determine the administrative or operational status for each cluster interface:

   a. Display the network port attributes:

      ```
      network port show
      ```

Show example

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1

Ignore

                                              Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
-

Node: n2

Ignore

                                              Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
-
4 entries were displayed.
```

b. Display information about the logical interfaces:

```
network interface show
```

Show example

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical    Status      Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node
Port     Home
----------- ---------- ---------- ------------------
------------- ------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24          n1
e4a      true
            n1_clus2   up/up      10.10.0.2/24          n1
e4e      true
            n2_clus1   up/up      10.10.0.3/24          n2
e4a      true
            n2_clus2   up/up      10.10.0.4/24          n2
e4e      true
4 entries were displayed.
```

3. Verify that the appropriate RCFs and image are installed on the new 3132Q-V switches as necessary for your requirements, and make any essential site customizations, such as users and passwords, network addresses, and so on.

   You must prepare both switches at this time. If you need to upgrade the RCF and image software, you must follow these steps:

   a. Go to the Cisco Ethernet Switches page on the NetApp Support Site.

   b. Note your switch and the required software versions in the table on that page.

   c. Download the appropriate version of RCF.

   d. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.

   e. Download the appropriate version of the image software.

4. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.

Step 2: Move first cluster port to C1

1. On Nexus 3132Q-V switches C1 and C2, disable all node-facing ports C1 and C2, but do not disable the ISL ports.

Show example

The following example shows ports 1 through 30 being disabled on Nexus 3132Q-V cluster switches C1 and C2 using a configuration supported in RCF `NX3132_RCF_v1.1_24p10g_26p40g.txt`:

```
C1# copy running-config startup-config
[########################################] 100%
Copy complete.
C1# configure
C1(config)# int e1/1/1-4,e1/2/1-4,e1/3/1-4,e1/4/1-4,e1/5/1-4,e1/6/1-
4,e1/7-30
C1(config-if-range)# shutdown
C1(config-if-range)# exit
C1(config)# exit

C2# copy running-config startup-config
[########################################] 100%
Copy complete.
C2# configure
C2(config)# int e1/1/1-4,e1/2/1-4,e1/3/1-4,e1/4/1-4,e1/5/1-4,e1/6/1-
4,e1/7-30
C2(config-if-range)# shutdown
C2(config-if-range)# exit
C2(config)# exit
```

2. Connect ports 1/31 and 1/32 on C1 to the same ports on C2 using supported cabling.

3. Verify that the ISL ports are operational on C1 and C2:

```
show port-channel summary
```

Show example

```
C1# show port-channel summary
Flags: D - Down         P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-        Type    Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1    Po1(SU)       Eth     LACP      Eth1/31(P)   Eth1/32(P)

C2# show port-channel summary
Flags: D - Down         P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-        Type    Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1    Po1(SU)       Eth     LACP      Eth1/31(P)   Eth1/32(P)
```

4. Display the list of neighboring devices on the switch:

```
show cdp neighbors
```

Show example

```
C1# show cdp neighbors
Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-
Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute

Device-ID          Local Intrfce  Hldtme Capability  Platform
Port ID
C2                 Eth1/31        174    R S I s     N3K-C3132Q-V
Eth1/31
C2                 Eth1/32        174    R S I s     N3K-C3132Q-V
Eth1/32

Total entries displayed: 2

C2# show cdp neighbors
Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-
Bridge
                  S - Switch, H - Host, I - IGMP, r - Repeater,
                  V - VoIP-Phone, D - Remotely-Managed-Device,
                  s - Supports-STP-Dispute

Device-ID          Local Intrfce  Hldtme Capability  Platform
Port ID
C1                 Eth1/31        178    R S I s     N3K-C3132Q-V
Eth1/31
C1                 Eth1/32        178    R S I s     N3K-C3132Q-V
Eth1/32

Total entries displayed: 2
```

5. Display the cluster port connectivity on each node:

```
network device-discovery show
```

Show example

The following example shows a two-node switchless cluster configuration.

```
cluster::*> network device-discovery show
            Local  Discovered
Node        Port   Device                  Interface        Platform
----------- ------ ------------------- ----------------
----------------
n1          /cdp
            e4a    n2                      e4a              FAS9000
            e4e    n2                      e4e              FAS9000
n2          /cdp
            e4a    n1                      e4a              FAS9000
            e4e    n1                      e4e              FAS9000
```

6. Migrate the clus1 interface to the physical port hosting clus2:

`network interface migrate`

Execute this command from each local node.

Show example

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus1
-source-node n1
-destination-node n1 -destination-port e4e
cluster::*> network interface migrate -vserver Cluster -lif n2_clus1
-source-node n2
-destination-node n2 -destination-port e4e
```

7. Verify the cluster interfaces migration:

`network interface show`

Show example

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status      Network                 Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask            Node
Port     Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
            n1_clus1   up/up       10.10.0.1/24           n1
e4e      false
            n1_clus2   up/up       10.10.0.2/24           n1
e4e      true
            n2_clus1   up/up       10.10.0.3/24           n2
e4e      false
            n2_clus2   up/up       10.10.0.4/24           n2
e4e      true
4 entries were displayed.
```

8. Shut down cluster ports clus1 LIF on both nodes:

   `network port modify`

   ```
   cluster::*> network port modify -node n1 -port e4a -up-admin false
   cluster::*> network port modify -node n2 -port e4a -up-admin false
   ```

9. Ping the remote cluster interfaces and perform an RPC server check:

   `cluster ping-cluster`

Show example

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e4a 10.10.0.1
Cluster n1_clus2 n1      e4e 10.10.0.2
Cluster n2_clus1 n2      e4a 10.10.0.3
Cluster n2_clus2 n2      e4e 10.10.0.4

Local = 10.10.0.1 10.10.0.2
Remote = 10.10.0.3 10.10.0.4
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
...............
Detected 1500 byte MTU on 32 path(s):
    Local 10.10.0.1 to Remote 10.10.0.3
    Local 10.10.0.1 to Remote 10.10.0.4
    Local 10.10.0.2 to Remote 10.10.0.3
    Local 10.10.0.2 to Remote 10.10.0.4
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
1 paths up, 0 paths down (tcp check)
1 paths up, 0 paths down (ucp check)
```

10. Disconnect the cable from e4a on node n1.

    You can refer to the running configuration and connect the first 40 GbE port on the switch C1 (port 1/7 in this example) to e4a on n1 using supported cabling on Nexus 3132Q-V.

    (i)   When reconnecting any cables to a new Cisco cluster switch, the cables used must be either fiber or cabling supported by Cisco.

11. Disconnect the cable from e4a on node n2.

    You can refer to the running configuration and connect e4a to the next available 40 GbE port on C1, port 1/8, using supported cabling.

12. Enable all node-facing ports on C1.

The following example shows ports 1 through 30 being enabled on Nexus 3132Q-V cluster switches C1 and C2 using the configuration supported in RCF `NX3132_RCF_v1.1_24p10g_26p40g.txt`:

```
C1# configure
C1(config)# int e1/1/1-4,e1/2/1-4,e1/3/1-4,e1/4/1-4,e1/5/1-4,e1/6/1-
4,e1/7-30
C1(config-if-range)# no shutdown
C1(config-if-range)# exit
C1(config)# exit
```

13. Enable the first cluster port, e4a, on each node:

`network port modify`

```
cluster::*> network port modify -node n1 -port e4a -up-admin true
cluster::*> network port modify -node n2 -port e4a -up-admin true
```

14. Verify that the clusters are up on both nodes:

`network port show`

Show example

```
cluster::*> network port show -role cluster
  (network port show)
Node: n1

Ignore
                                             Speed(Mbps) Health
Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster       Cluster          up   9000 auto/40000  -
-
e4e       Cluster       Cluster          up   9000 auto/40000  -
-

Node: n2

Ignore
                                             Speed(Mbps) Health
Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster       Cluster          up   9000 auto/40000  -
-
e4e       Cluster       Cluster          up   9000 auto/40000  -
-
4 entries were displayed.
```

15. For each node, revert all of the migrated cluster interconnect LIFs:

`network interface revert`

Show example

The following example shows the migrated LIFs being reverted to their home ports.

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus1
cluster::*> network interface revert -vserver Cluster -lif n2_clus1
```

16. Verify that all of the cluster interconnect ports are now reverted to their home ports:

```
network interface show
```

The `Is Home` column should display a value of `true` for all of the ports listed in the `Current Port` column. If the displayed value is `false`, the port has not been reverted.

Show example

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status     Network                 Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask        Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24        n1
e4a     true
            n1_clus2   up/up      10.10.0.2/24        n1
e4e     true
            n2_clus1   up/up      10.10.0.3/24        n2
e4a     true
            n2_clus2   up/up      10.10.0.4/24        n2
e4e     true
4 entries were displayed.
```

Step 3: Move second cluster port to C2

1. Display the cluster port connectivity on each node:

```
network device-discovery show
```

Show example

```
cluster::*> network device-discovery show
            Local  Discovered
Node        Port   Device                 Interface        Platform
----------- ------ ------------------     ----------------
----------------
n1          /cdp
            e4a    C1                     Ethernet1/7      N3K-C3132Q-V
            e4e    n2                     e4e              FAS9000
n2          /cdp
            e4a    C1                     Ethernet1/8      N3K-C3132Q-V
            e4e    n1                     e4e              FAS9000
```

2. On the console of each node, migrate clus2 to port e4a:

```
network interface migrate
```

Show example

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus2
-source-node n1
-destination-node n1 -destination-port e4a
cluster::*> network interface migrate -vserver Cluster -lif n2_clus2
-source-node n2
-destination-node n2 -destination-port e4a
```

3. Shut down cluster ports clus2 LIF on both nodes:

```
network port modify
```

The following example shows the specified ports being shut down on both nodes:

```
cluster::*> network port modify -node n1 -port e4e -up-admin false
cluster::*> network port modify -node n2 -port e4e -up-admin false
```

4. Verify the cluster LIF status:

```
network interface show
```

Show example

```
cluster::*> network interface show -role cluster
  (network interface show)
           Logical    Status      Network                 Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask         Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
           n1_clus1   up/up       10.10.0.1/24        n1
e4a      true
           n1_clus2   up/up       10.10.0.2/24        n1
e4a      false
           n2_clus1   up/up       10.10.0.3/24        n2
e4a      true
           n2_clus2   up/up       10.10.0.4/24        n2
e4a      false
4 entries were displayed.
```

5. Disconnect the cable from e4e on node n1.

   You can refer to the running configuration and connect the first 40 GbE port on the switch C2 (port 1/7 in this example) to e4e on n1 using supported cabling on Nexus 3132Q-V.

6. Disconnect the cable from e4e on node n2.

   You can refer to the running configuration and connect e4e to the next available 40 GbE port on C2, port 1/8, using supported cabling.

7. Enable all node-facing ports on C2.

Show example

The following example shows ports 1 through 30 being enabled on Nexus 3132Q-V cluster switches C1 and C2 using a configuration supported in RCF `NX3132_RCF_v1.1_24p10g_26p40g.txt`:

```
C2# configure
C2(config)# int e1/1/1-4,e1/2/1-4,e1/3/1-4,e1/4/1-4,e1/5/1-4,e1/6/1-
4,e1/7-30
C2(config-if-range)# no shutdown
C2(config-if-range)# exit
C2(config)# exit
```

8. Enable the second cluster port, e4e, on each node:

`network port modify`

The following example shows the specified ports being brought up:

```
cluster::*> network port modify -node n1 -port e4e -up-admin true
cluster::*> network port modify -node n2 -port e4e -up-admin true
```

9. For each node, revert all of the migrated cluster interconnect LIFs:

`network interface revert`

The following example shows the migrated LIFs being reverted to their home ports.

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus2
cluster::*> network interface revert -vserver Cluster -lif n2_clus2
```

10. Verify that all of the cluster interconnect ports are now reverted to their home ports:

`network interface show`

The `Is Home` column should display a value of `true` for all of the ports listed in the `Current Port` column. If the displayed value is `false`, the port has not been reverted.

Show example

```
cluster::*> network interface show -role cluster
  (network interface show)
           Logical     Status     Network                Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask           Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
           n1_clus1    up/up      10.10.0.1/24           n1
e4a     true
           n1_clus2    up/up      10.10.0.2/24           n1
e4e     true
           n2_clus1    up/up      10.10.0.3/24           n2
e4a     true
           n2_clus2    up/up      10.10.0.4/24           n2
e4e     true
4 entries were displayed.
```

11. Verify that all of the cluster interconnect ports are in the `up` state.

```
network port show -role cluster
```

Show example

```
cluster::*> network port show -role cluster
   (network port show)
Node: n1

Ignore
                                           Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
-


Node: n2

Ignore
                                           Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster          up   9000 auto/40000  -
-
e4e       Cluster      Cluster          up   9000 auto/40000  -
-
4 entries were displayed.
```

Step 4: Disable the two-node switchless cluster option

1. Display the cluster switch port numbers each cluster port is connected to on each node:

```
network device-discovery show
```

Show example

```
    cluster::*> network device-discovery show
            Local  Discovered
Node        Port   Device              Interface        Platform
----------- ------ ------------------- ----------------
----------------
n1          /cdp
            e4a    C1                  Ethernet1/7      N3K-C3132Q-V
            e4e    C2                  Ethernet1/7      N3K-C3132Q-V
n2          /cdp
            e4a    C1                  Ethernet1/8      N3K-C3132Q-V
            e4e    C2                  Ethernet1/8      N3K-C3132Q-V
```

2. Display discovered and monitored cluster switches:

```
system cluster-switch show
```

Show example

```
cluster::*> system cluster-switch show

Switch                          Type              Address
Model
--------------------------- ----------------- ----------------
--------------
C1                              cluster-network    10.10.1.101
NX3132V
      Serial Number: FOX000001
        Is Monitored: true
              Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                     7.0(3)I4(1)
     Version Source: CDP

C2                              cluster-network    10.10.1.102
NX3132V
      Serial Number: FOX000002
        Is Monitored: true
              Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                     7.0(3)I4(1)
     Version Source: CDP

2 entries were displayed.
```

3. Disable the two-node switchless configuration settings on any node:

   `network options switchless-cluster`

   ```
   network options switchless-cluster modify -enabled false
   ```

4. Verify that the `switchless-cluster` option has been disabled.

   ```
   network options switchless-cluster show
   ```

Step 5: Verify the configuration

1. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

Show example

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1     e4a 10.10.0.1
Cluster n1_clus2 n1     e4e 10.10.0.2
Cluster n2_clus1 n2     e4a 10.10.0.3
Cluster n2_clus2 n2     e4e 10.10.0.4

Local = 10.10.0.1 10.10.0.2
Remote = 10.10.0.3 10.10.0.4
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 32 path(s):
    Local 10.10.0.1 to Remote 10.10.0.3
    Local 10.10.0.1 to Remote 10.10.0.4
    Local 10.10.0.2 to Remote 10.10.0.3
    Local 10.10.0.2 to Remote 10.10.0.4
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
1 paths up, 0 paths down (tcp check)
1 paths up, 0 paths down (ucp check)
```

2. Enable the cluster switch health monitor log collection feature for collecting switch-related log files:

```
system cluster-switch log setup-password
```

```
system cluster-switch log enable-collection
```

Show example

```
cluster::*> **system cluster-switch log setup-password**
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
C1
C2

cluster::*> system cluster-switch log setup-password

Enter the switch name: C1
RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log setup-password

Enter the switch name: C2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster::*>
```

ⓘ   If any of these commands return an error, contact NetApp support.

3. If you suppressed automatic case creation, re-enable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Replace switches

Requirements for replacing Cisco Nexus 3132Q-V cluster switches

Make sure you understand the configuration requirements, port connections, and cabling requirements when you replace cluster switches.

Cisco Nexus 3132Q-V requirements

- The Cisco Nexus 3132Q-V cluster switch is supported.

- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco® Cluster Network Switch Reference Configuration File Download page.

- The cluster switches use the Inter-Switch Link (ISL) ports e1/31-32.

- The Hardware Universe contains information about supported cabling to Nexus 3132Q-V switches:

    - The nodes with 10 GbE cluster connections require QSFP optical modules with breakout fiber cables or QSFP to SFP+ copper break-out cables.

    - The nodes with 40/100 GbE cluster connections require supported QSFP/QSFP28 optical modules with fiber cables or QSFP/QSFP28 copper direct-attach cables.

    - The cluster switches use the appropriate ISL cabling: 2x QSFP28 fiber or copper direct-attach cables.

- On Nexus 3132Q-V, you can operate QSFP ports as either 40/100 Gb Ethernet or 4 x10 Gb Ethernet modes.

    By default, there are 32 ports in the 40/100 Gb Ethernet mode. These 40 Gb Ethernet ports are numbered in a 2-tuple naming convention. For example, the second 40 Gb Ethernet port is numbered as 1/2. The process of changing the configuration from 40 Gb Ethernet to 10 Gb Ethernet is called *breakout* and the process of changing the configuration from 10 Gb Ethernet to 40 Gb Ethernet is called *breakin*. When you break out a 40/100 Gb Ethernet port into 10 Gb Ethernet ports, the resulting ports are numbered using a 3-tuple naming convention. For example, the breakout ports of the second 40/100 Gb Ethernet port are numbered as 1/2/1, 1/2/2, 1/2/3, 1/2/4.

- On the left side of Nexus 3132Q-V is a set of four SFP+ ports multiplexed to the first QSFP port.

    By default, the RCF is structured to use the first QSFP port.

    You can make four SFP+ ports active instead of a QSFP port for Nexus 3132Q-V by using the `hardware profile front portmode sfp-plus` command. Similarly, you can reset Nexus 3132Q-V to use a QSFP port instead of four SFP+ ports by using the `hardware profile front portmode qsfp` command.

- You must have configured some of the ports on Nexus 3132Q-V to run at 10 GbE or 40/100 GbE.

    You can break-out the first six ports into 4x10 GbE mode by using the `interface breakout module 1 port 1-6 map 10g-4x` command. Similarly, you can regroup the first six QSFP+ ports from breakout configuration by using the `no interface breakout module 1 port 1-6 map 10g-4x` command.

- You must have done the planning, migration, and read the required documentation on 10 GbE and 40/100 GbE connectivity from nodes to Nexus 3132Q-V cluster switches.

    The Cisco Ethernet Switches page has information about the ONTAP and NX-OS versions supported in this procedure.

Cisco Nexus 5596 requirements

- The following cluster switches are supported:
    - Nexus 5596
    - Nexus 3132Q-V

- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco® Cluster Network Switch Reference Configuration File Download page.

- The cluster switches use the following ports for connections to nodes:
    - Ports e1/1-40 (10 GbE): Nexus 5596
    - Ports e1/1-30 (40/100 GbE): Nexus 3132Q-V

- The cluster switches use the following Inter-Switch Link (ISL) ports:
    - Ports e1/41-48 (10 GbE): Nexus 5596
    - Ports e1/31-32 (40/100 GbE): Nexus 3132Q-V

- The Hardware Universe contains information about supported cabling to Nexus 3132Q-V switches:
    - Nodes with 10 GbE cluster connections require QSFP to SFP+ optical fiber breakout cables or QSFP to SFP+ copper breakout cables.
    - Nodes with 40/100 GbE cluster connections require supported QSFP/QSFP28optical modules with fiber cables or QSFP/QSFP28 copper direct-attach cables.

- The cluster switches use the appropriate ISL cabling:
    - Beginning: Nexus 5596 to Nexus 5596 (SFP+ to SFP+)
        - 8x SFP+ fiber or copper direct-attach cables
    - Interim: Nexus 5596 to Nexus 3132Q-V (QSFP to 4xSFP+ break-out)
        - 1x QSFP to SFP+ fiber break-out or copper break-out cables
    - Final: Nexus 3132Q-V to Nexus 3132Q-V (QSFP28 to QSFP28)
        - 2x QSFP28 fiber or copper direct-attach cables

- On Nexus 3132Q-V switches, you can operate QSFP/QSFP28 ports as either 40/100 Gigabit Ethernet or 4 x10 Gigabit Ethernet modes.

    By default, there are 32 ports in the 40/100 Gigabit Ethernet mode. These 40 Gigabit Ethernet ports are numbered in a 2-tuple naming convention. For example, the second 40 Gigabit Ethernet port is numbered as 1/2. The process of changing the configuration from 40 Gigabit Ethernet to 10 Gigabit Ethernet is called *breakout* and the process of changing the configuration from 10 Gigabit Ethernet to 40 Gigabit Ethernet is called *breakin*. When you break out a 40/100 Gigabit Ethernet port into 10 Gigabit Ethernet ports, the resulting ports are numbered using a 3-tuple naming convention. For example, the break-out ports of the second 40 Gigabit Ethernet port are numbered as 1/2/1, 1/2/2, 1/2/3, and 1/2/4.

- On the left side of Nexus 3132Q-V switches is a set of 4 SFP+ ports multiplexed to that QSFP28 port.

    By default, the RCF is structured to use the QSFP28 port.

    > ⓘ You can make 4x SFP+ ports active instead of a QSFP port for Nexus 3132Q-V switches by using the `hardware profile front portmode sfp-plus` command. Similarly, you can reset Nexus 3132Q-V switches to use a QSFP port instead of 4x SFP+ ports by using the `hardware profile front portmode qsfp` command.

- You have configured some of the ports on Nexus 3132Q-V switches to run at 10 GbE or 40/100 GbE.

  > (i) You can break out the first six ports into 4x10 GbE mode by using the `interface breakout module 1 port 1-6 map 10g-4x` command. Similarly, you can regroup the first six QSFP+ ports from breakout configuration by using the `no interface breakout module 1 port 1-6 map 10g-4x` command.

- You have done the planning, migration, and read the required documentation on 10 GbE and 40/100 GbE connectivity from nodes to Nexus 3132Q-V cluster switches.

- The ONTAP and NX-OS versions supported in this procedure are on the Cisco Ethernet Switches page.

NetApp CN1610 requirements

- The following cluster switches are supported:
  - NetApp CN1610
  - Cisco Nexus 3132Q-V

- The cluster switches support the following node connections:
  - NetApp CN1610: ports 0/1 through 0/12 (10 GbE)
  - Cisco Nexus 3132Q-V: ports e1/1-30 (40/100 GbE)

- The cluster switches use the following inter-switch link (ISL) ports:
  - NetApp CN1610: ports 0/13 through 0/16 (10 GbE)
  - Cisco Nexus 3132Q-V: ports e1/31-32 (40/100 GbE)

- The Hardware Universe contains information about supported cabling to Nexus 3132Q-V switches:
  - Nodes with 10 GbE cluster connections require QSFP to SFP+ optical fiber breakout cables or QSFP to SFP+ copper breakout cables
  - Nodes with 40/100 GbE cluster connections require supported QSFP/QSFP28 optical modules with optical fiber cables or QSFP/QSFP28 copper direct-attach cables

- The appropriate ISL cabling is as follows:
  - Beginning: For CN1610 to CN1610 (SFP+ to SFP+), four SFP+ optical fiber or copper direct-attach cables
  - Interim: For CN1610 to Nexus 3132Q-V (QSFP to four SFP+ breakout), one QSFP to SFP+ optical fiber or copper breakout cable
  - Final: For Nexus 3132Q-V to Nexus 3132Q-V (QSFP28 to QSFP28), two QSFP28 optical fiber or copper direct-attach cables

- NetApp twinax cables are not compatible with Cisco Nexus 3132Q-V switches.

  If your current CN1610 configuration uses NetApp twinax cables for cluster-node-to-switch connections or ISL connections and you want to continue using twinax in your environment, you need to procure Cisco twinax cables. Alternatively, you can use optical fiber cables for both the ISL connections and the cluster-node-to-switch connections.

- On Nexus 3132Q-V switches, you can operate QSFP/QSFP28 ports as either 40/100 Gb Ethernet or 4x 10 Gb Ethernet modes.

  By default, there are 32 ports in the 40/100 Gb Ethernet mode. These 40 Gb Ethernet ports are numbered in a 2-tuple naming convention. For example, the second 40 Gb Ethernet port is numbered as 1/2. The

process of changing the configuration from 40 Gb Ethernet to 10 Gb Ethernet is called *breakout* and the process of changing the configuration from 10 Gb Ethernet to 40 Gb Ethernet is called *breakin*. When you break out a 40/100 Gb Ethernet port into 10 Gb Ethernet ports, the resulting ports are numbered using a 3-tuple naming convention. For example, the breakout ports of the second 40 Gb Ethernet port are numbered as 1/2/1, 1/2/2, 1/2/3, and 1/2/4.

- On the left side of Nexus 3132Q-V switches is a set of four SFP+ ports multiplexed to the first QSFP port.

  By default, the reference configuration file (RCF) is structured to use the first QSFP port.

  You can make four SFP+ ports active instead of a QSFP port for Nexus 3132Q-V switches by using the `hardware profile front portmode sfp-plus` command. Similarly, you can reset Nexus 3132Q-V switches to use a QSFP port instead of four SFP+ ports by using the `hardware profile front portmode qsfp` command.

  > ⓘ When you use the first four SFP+ ports, it will disable the first 40GbE QSFP port.

- You must have configured some of the ports on Nexus 3132Q-V switches to run at 10 GbE or 40/100 GbE.

  You can break out the first six ports into 4x 10 GbE mode by using the `interface breakout module 1 port 1-6 map 10g-4x` command. Similarly, you can regroup the first six QSFP+ ports from *breakout* configuration by using the `no interface breakout module 1 port 1-6 map 10g-4x` command.

- You must have done the planning, migration, and read the required documentation on 10 GbE and 40/100 GbE connectivity from nodes to Nexus 3132Q-V cluster switches.

- The ONTAP and NX-OS versions that are supported in this procedure are listed on the Cisco Ethernet Switches page.

- The ONTAP and FASTPATH versions that are supported in this procedure are listed on the NetApp CN1601 and CN1610 Switches page.

Replace Cisco Nexus 3132Q-V cluster switches

Follow this procedure to replace a defective Cisco Nexus 3132Q-V switch in a cluster network. The replacement procedure is a nondisruptive procedure (NDO).

Review requirements

Switch requirements

Review the Requirements for replacing Cisco Nexus 3132Q-V cluster switches.

What you'll need

- The existing cluster and network configuration has:
  - The Nexus 3132Q-V cluster infrastructure is redundant and fully functional on both switches.

    The Cisco Ethernet Switch page has the latest RCF and NX-OS versions on your switches.

  - All cluster ports are in the `up` state.

  - Management connectivity exists on both switches.

  - All cluster logical interfaces (LIFs) are in the `up` state and have been migrated.

- For the Nexus 3132Q-V replacement switch, make sure that:

- Management network connectivity on the replacement switch is functional.

- Console access to the replacement switch is in place.

- The desired RCF and NX-OS operating system image switch is loaded onto the switch.

- Initial customization of the switch is complete.

- Hardware Universe

Replace the switch

This procedure replaces the second Nexus 3132Q-V cluster switch CL2 with new 3132Q-V switch C2.

About the examples

The examples in this procedure use the following switch and node nomenclature:

- n1_clus1 is the first cluster logical interface (LIF) connected to cluster switch C1 for node n1.

- n1_clus2 is the first cluster LIF connected to cluster switch CL2 or C2, for node n1.

- n1_clus3 is the second LIF connected to cluster switch C2, for node n1.

- n1_clus4 is the second LIF connected to cluster switch CL1, for node n1.

- The number of 10 GbE and 40/100 GbE ports are defined in the reference configuration files (RCFs) available on the Cisco® Cluster Network Switch Reference Configuration File Download page.

- The nodes are n1, n2, n3, and n4. - The examples in this procedure use four nodes: Two nodes use four 10 GB cluster interconnect ports: e0a, e0b, e0c, and e0d. The other two nodes use two 40 GB cluster interconnect ports: e4a and e4e. See the Hardware Universe for the actual cluster ports on your platforms.

About this task

This procedure covers the following scenario:

- The cluster starts with four nodes connected to two Nexus 3132Q-V cluster switches, CL1 and CL2.

- Cluster switch CL2 is to be replaced by C2

  - On each node, cluster LIFs connected to CL2 are migrated onto cluster ports connected to CL1.

  - Disconnect cabling from all ports on CL2 and reconnect cabling to the same ports on the replacement switch C2.

  - On each node, its migrated cluster LIFs are reverted.

Step 1: Prepare for replacement

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all - message MAINT=xh
```

x is the duration of the maintenance window in hours.

> ⓘ The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

2. Display information about the devices in your configuration:

```
network device-discovery show
```

Show example

```
cluster::> network device-discovery show
            Local  Discovered
Node       Port   Device               Interface          Platform
---------- ------ ------------------- ----------------
----------------
n1         /cdp
           e0a    CL1                  Ethernet1/1/1    N3K-C3132Q-V
           e0b    CL2                  Ethernet1/1/1    N3K-C3132Q-V
           e0c    CL2                  Ethernet1/1/2    N3K-C3132Q-V
           e0d    CL1                  Ethernet1/1/2    N3K-C3132Q-V
n2         /cdp
           e0a    CL1                  Ethernet1/1/3    N3K-C3132Q-V
           e0b    CL2                  Ethernet1/1/3    N3K-C3132Q-V
           e0c    CL2                  Ethernet1/1/4    N3K-C3132Q-V
           e0d    CL1                  Ethernet1/1/4    N3K-C3132Q-V
n3         /cdp
           e4a    CL1                  Ethernet1/7      N3K-C3132Q-V
           e4e    CL2                  Ethernet1/7      N3K-C3132Q-V
n4         /cdp
           e4a    CL1                  Ethernet1/8      N3K-C3132Q-V
           e4e    CL2                  Ethernet1/8      N3K-C3132Q-V

12 entries were displayed
```

3. Determine the administrative or operational status for each cluster interface:

   a. Display the network port attributes:

   ```
   network port show
   ```

Show example

```
cluster::*> network port show -role cluster
        (network port show)

Node: n1

Ignore
                                            Speed(Mbps)
Health    Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster         up   9000 auto/10000  -
-
e0b       Cluster      Cluster         up   9000 auto/10000  -
-
e0c       Cluster      Cluster         up   9000 auto/10000  -
-
e0d       Cluster      Cluster         up   9000 auto/10000  -
-

Node: n2

Ignore
                                            Speed(Mbps)
Health    Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster         up   9000  auto/10000 -
-
e0b       Cluster      Cluster         up   9000  auto/10000 -
-
e0c       Cluster      Cluster         up   9000  auto/10000 -
-
e0d       Cluster      Cluster         up   9000  auto/10000 -
-

Node: n3

Ignore
                                            Speed(Mbps)
Health    Health
```

```
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster              up   9000 auto/40000   -
-
e4e       Cluster      Cluster              up   9000 auto/40000   -
-


Node: n4

Ignore
                                                    Speed(Mbps)
Health    Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper
Status    Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster              up   9000 auto/40000   -
-
e4e       Cluster      Cluster              up   9000 auto/40000   -
-
12 entries were displayed.
```

b. Display information about the logical interfaces:

```
network interface show
```

Show example

```
cluster::*> network interface show -role cluster
        (network interface show)

              Logical     Status     Network                   Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask       Node
Port     Home
----------- ---------- ---------- ------------------
------------- ------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24          n1
e0a     true
            n1_clus2   up/up      10.10.0.2/24          n1
e0b     true
            n1_clus3   up/up      10.10.0.3/24          n1
e0c     true
            n1_clus4   up/up      10.10.0.4/24          n1
e0d     true
            n2_clus1   up/up      10.10.0.5/24          n2
e0a     true
            n2_clus2   up/up      10.10.0.6/24          n2
e0b     true
            n2_clus3   up/up      10.10.0.7/24          n2
e0c     true
            n2_clus4   up/up      10.10.0.8/24          n2
e0d     true
            n3_clus1   up/up      10.10.0.9/24          n3
e0a     true
            n3_clus2   up/up      10.10.0.10/24         n3
e0e     true
            n4_clus1   up/up      10.10.0.11/24         n4
e0a     true
            n4_clus2   up/up      10.10.0.12/24         n4
e0e     true

12 entries were displayed.
```

c. Display the information on the discovered cluster switches:

```
system cluster-switch show
```

Show example

```
cluster::> system cluster-switch show

Switch                          Type              Address
Model
--------------------------- ----------------- ----------------
---------------
CL1                               cluster-network   10.10.1.101
NX3132V
        Serial Number: FOX000001
          Is Monitored: true
                Reason:
    Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                       7.0(3)I4(1)
      Version Source: CDP

CL2                               cluster-network   10.10.1.102
NX3132V
        Serial Number: FOX000002
          Is Monitored: true
                Reason:
    Software Version: Cisco Nexus Operating System (NX-OS)
Software, Version
                       7.0(3)I4(1)
      Version Source: CDP

2 entries were displayed.
```

4. Verify that the appropriate RCF and image are installed on the new Nexus 3132Q-V switch as necessary for your requirements, and make any essential site customizations.

   You must prepare the replacement switch at this time. If you need to upgrade the RCF and image, you must follow these steps:

   a. On the NetApp Support Site, go to the Cisco Ethernet Switch page.

   b. Note your switch and the required software versions in the table on that page.

   c. Download the appropriate version of the RCF.

   d. Click **CONTINUE** on the **Description** page, accept the license agreement, and then follow the instructions on the **Download** page to download the RCF.

   e. Download the appropriate version of the image software.

5. Migrate the LIFs associated to the cluster ports connected to switch C2:

   ```
   network interface migrate
   ```

Show example

This example shows that the LIF migration is done on all the nodes:

```
cluster::*> network interface migrate -vserver Cluster -lif n1_clus2
-source-node n1 –destination-node n1 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n1_clus3
-source-node n1 –destination-node n1 -destination-port e0d
cluster::*> network interface migrate -vserver Cluster -lif n2_clus2
-source-node n2 –destination-node n2 -destination-port e0a
cluster::*> network interface migrate -vserver Cluster -lif n2_clus3
-source-node n2 –destination-node n2 -destination-port e0d
cluster::*> network interface migrate -vserver Cluster -lif n3_clus2
-source-node n3 –destination-node n3 -destination-port e4a
cluster::*> network interface migrate -vserver Cluster -lif n4_clus2
-source-node n4 –destination-node n4 -destination-port e4a
```

6. Verify cluster's health:

```
network interface show
```

Show example

```
cluster::*> network interface show -role cluster
        (network interface show)

            Logical     Status      Network             Current
Current Is
Vserver     Interface   Admin/Oper  Address/Mask        Node
Port      Home
----------- ----------  ----------  ------------------  -------------
------- ----
Cluster
            n1_clus1    up/up       10.10.0.1/24        n1
e0a     true
            n1_clus2    up/up       10.10.0.2/24        n1
e0a     false
            n1_clus3    up/up       10.10.0.3/24        n1
e0d     false
            n1_clus4    up/up       10.10.0.4/24        n1
e0d     true
            n2_clus1    up/up       10.10.0.5/24        n2
e0a     true
            n2_clus2    up/up       10.10.0.6/24        n2
e0a     false
            n2_clus3    up/up       10.10.0.7/24        n2
e0d     false
            n2_clus4    up/up       10.10.0.8/24        n2
e0d     true
            n3_clus1    up/up       10.10.0.9/24        n3
e4a     true
            n3_clus2    up/up       10.10.0.10/24       n3
e4a     false
            n4_clus1    up/up       10.10.0.11/24       n4
e4a     true
            n4_clus2    up/up       10.10.0.12/24       n4
e4a     false
12 entries were displayed.
```

7. Shut down the cluster interconnect ports that are physically connected to switch CL2:

```
network port modify
```

Show example

This example shows the specified ports being shut down on all nodes:

```
cluster::*> network port modify -node n1 -port e0b -up-admin false
cluster::*> network port modify -node n1 -port e0c -up-admin false
cluster::*> network port modify -node n2 -port e0b -up-admin false
cluster::*> network port modify -node n2 -port e0c -up-admin false
cluster::*> network port modify -node n3 -port e4e -up-admin false
cluster::*> network port modify -node n4 -port e4e -up-admin false
```

8. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

Show example

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8
Cluster n3_clus1 n4      e0a 10.10.0.9
Cluster n3_clus2 n3      e0e 10.10.0.10
Cluster n4_clus1 n4      e0a 10.10.0.11
Cluster n4_clus2 n4      e0e 10.10.0.12

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8 10.10.0.9
10.10.0.10 10.10.0.11 10.10.0.12
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 32 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 32 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.1 to Remote 10.10.0.9
    Local 10.10.0.1 to Remote 10.10.0.10
    Local 10.10.0.1 to Remote 10.10.0.11
    Local 10.10.0.1 to Remote 10.10.0.12
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.9
    Local 10.10.0.2 to Remote 10.10.0.10
    Local 10.10.0.2 to Remote 10.10.0.11
    Local 10.10.0.2 to Remote 10.10.0.12
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
```

```
      Local 10.10.0.3 to Remote 10.10.0.7
      Local 10.10.0.3 to Remote 10.10.0.8
      Local 10.10.0.3 to Remote 10.10.0.9
      Local 10.10.0.3 to Remote 10.10.0.10
      Local 10.10.0.3 to Remote 10.10.0.11
      Local 10.10.0.3 to Remote 10.10.0.12
      Local 10.10.0.4 to Remote 10.10.0.5
      Local 10.10.0.4 to Remote 10.10.0.6
      Local 10.10.0.4 to Remote 10.10.0.7
      Local 10.10.0.4 to Remote 10.10.0.8
      Local 10.10.0.4 to Remote 10.10.0.9
      Local 10.10.0.4 to Remote 10.10.0.10
      Local 10.10.0.4 to Remote 10.10.0.11
      Local 10.10.0.4 to Remote 10.10.0.12

  Larger than PMTU communication succeeds on 32 path(s)
  RPC status:
  8 paths up, 0 paths down (tcp check)
  8 paths up, 0 paths down (udp check)
```

9. Shut down the ports 1/31 and 1/32 on CL1, and the active Nexus 3132Q-V switch:

   `shutdown`

   Show example

   This example shows the ISL ports 1/31 and 1/32 being shut down on switch CL1:

   ```
   (CL1)# configure
   (CL1)(Config)# interface e1/31-32
   (CL1(config-if-range)# shutdown
   (CL1(config-if-range)# exit
   (CL1)(Config)# exit
   (CL1)#
   ```

Step 2: Configure ports

1. Remove all the cables attached to the Nexus 3132Q-V switch CL2 and reconnect them to the replacement switch C2 on all nodes.

2. Remove the ISL cables from ports e1/31 and e1/32 on CL2 and reconnect them to the same ports on the replacement switch C2.

3. Bring up ISLs ports 1/31 and 1/32 on the Nexus 3132Q-V switch CL1:

```
(CL1)# configure
(CL1)(Config)# interface e1/31-32
(CL1(config-if-range)# no shutdown
(CL1(config-if-range)# exit
(CL1)(Config)# exit
(CL1)#
```

4. Verify that the ISLs are up on CL1:

show port-channel

Ports Eth1/31 and Eth1/32 should indicate (P), which means that the ISL ports are up in the port-channel.

Show example

```
CL1# show port-channel summary
Flags: D - Down          P - Up in port-channel (members)
       I - Individual    H - Hot-standby (LACP only)
       s - Suspended     r - Module-removed
       S - Switched      R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
------------------------------------------------------------------------
------------
Group Port-         Type    Protocol  Member
Ports
       Channel
------------------------------------------------------------------------
------------
1     Po1(SU)       Eth     LACP      Eth1/31(P)    Eth1/32(P)
```

5. Verify that the ISLs are up on C2:

show port-channel summary

Ports Eth1/31 and Eth1/32 should indicate (P), which means that both ISL ports are up in the port-channel.

Show example

```
C2# show port-channel summary
Flags: D - Down          P - Up in port-channel (members)
       I - Individual   H - Hot-standby (LACP only)
       s - Suspended    r - Module-removed
       S - Switched     R - Routed
       U - Up (port-channel)
       M - Not in use. Min-links not met
--------------------------------------------------------------------
------------
Group Port-          Type    Protocol  Member Ports
      Channel
--------------------------------------------------------------------
------------
1     Po1(SU)        Eth     LACP      Eth1/31(P)   Eth1/32(P)
```

6. On all nodes, bring up all the cluster interconnect ports connected to the Nexus 3132Q-V switch C2:

`network port modify`

Show example

```
cluster::*> network port modify -node n1 -port e0b -up-admin true
cluster::*> network port modify -node n1 -port e0c -up-admin true
cluster::*> network port modify -node n2 -port e0b -up-admin true
cluster::*> network port modify -node n2 -port e0c -up-admin true
cluster::*> network port modify -node n3 -port e4e -up-admin true
cluster::*> network port modify -node n4 -port e4e -up-admin true
```

7. For all nodes, revert all of the migrated cluster interconnect LIFs:

`network interface revert`

Show example

```
cluster::*> network interface revert -vserver Cluster -lif n1_clus2
cluster::*> network interface revert -vserver Cluster -lif n1_clus3
cluster::*> network interface revert -vserver Cluster -lif n2_clus2
cluster::*> network interface revert -vserver Cluster -lif n2_clus3
Cluster::*> network interface revert –vserver Cluster –lif n3_clus2
Cluster::*> network interface revert –vserver Cluster –lif n4_clus2
```

8. Verify that the cluster interconnect ports are now reverted to their home:

```
network interface show
```

Show example

This example shows that all the LIFs are successfully reverted because the ports listed under the `Current Port` column have a status of `true` in the `Is Home` column. If the `Is Home` column value is `false`, the LIF has not been reverted.

```
cluster::*> network interface show -role cluster
  (network interface show)
            Logical     Status     Network             Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask         Node
Port     Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
            n1_clus1   up/up      10.10.0.1/24        n1
e0a     true
            n1_clus2   up/up      10.10.0.2/24        n1
e0b     true
            n1_clus3   up/up      10.10.0.3/24        n1
e0c     true
            n1_clus4   up/up      10.10.0.4/24        n1
e0d     true
            n2_clus1   up/up      10.10.0.5/24        n2
e0a     true
            n2_clus2   up/up      10.10.0.6/24        n2
e0b     true
            n2_clus3   up/up      10.10.0.7/24        n2
e0c     true
            n2_clus4   up/up      10.10.0.8/24        n2
e0d     true
            n3_clus1   up/up      10.10.0.9/24        n3
e4a     true
            n3_clus2   up/up      10.10.0.10/24       n3
e4e     true
            n4_clus1   up/up      10.10.0.11/24       n4
e4a     true
            n4_clus2   up/up      10.10.0.12/24       n4
e4e     true
12 entries were displayed.
```

9. Verify that the cluster ports are connected:

```
network port show
```

Show example

```
cluster::*> network port show -role cluster
  (network port show)
Node: n1

Ignore
                                        Speed(Mbps) Health
Health
Port      IPspace     Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ----------- ---------------- ---- ---- -----------
-------- ------
e0a       Cluster     Cluster          up   9000 auto/10000  -
-
e0b       Cluster     Cluster          up   9000 auto/10000  -
-
e0c       Cluster     Cluster          up   9000 auto/10000  -
-
e0d       Cluster     Cluster          up   9000 auto/10000  -
-

Node: n2

Ignore
                                        Speed(Mbps) Health
Health
Port      IPspace     Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ----------- ---------------- ---- ---- -----------
-------- ------
e0a       Cluster     Cluster          up   9000  auto/10000 -
-
e0b       Cluster     Cluster          up   9000  auto/10000 -
-
e0c       Cluster     Cluster          up   9000  auto/10000 -
-
e0d       Cluster     Cluster          up   9000  auto/10000 -
-

Node: n3

Ignore
                                        Speed(Mbps) Health
Health
Port      IPspace     Broadcast Domain Link MTU  Admin/Oper  Status
```

```
   Status
   --------- ------------ --------------- ---- ---- -----------
   -------- ------
   e4a         Cluster      Cluster          up    9000 auto/40000  -
   -
   e4e         Cluster      Cluster          up    9000 auto/40000  -
   -


   Node: n4

   Ignore
                                                    Speed(Mbps) Health
   Health
   Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
   Status
   --------- ------------ --------------- ---- ---- -----------
   -------- ------
   e4a         Cluster      Cluster          up    9000 auto/40000  -
   -
   e4e         Cluster      Cluster          up    9000 auto/40000  -
   -
   12 entries were displayed.
```

10. Ping the remote cluster interfaces and perform an RPC server check:

```
cluster ping-cluster
```

Show example

```
cluster::*> cluster ping-cluster -node n1
Host is n1
Getting addresses from network interface table...
Cluster n1_clus1 n1      e0a 10.10.0.1
Cluster n1_clus2 n1      e0b 10.10.0.2
Cluster n1_clus3 n1      e0c 10.10.0.3
Cluster n1_clus4 n1      e0d 10.10.0.4
Cluster n2_clus1 n2      e0a 10.10.0.5
Cluster n2_clus2 n2      e0b 10.10.0.6
Cluster n2_clus3 n2      e0c 10.10.0.7
Cluster n2_clus4 n2      e0d 10.10.0.8
Cluster n3_clus1 n3      e0a 10.10.0.9
Cluster n3_clus2 n3      e0e 10.10.0.10
Cluster n4_clus1 n4      e0a 10.10.0.11
Cluster n4_clus2 n4      e0e 10.10.0.12

Local = 10.10.0.1 10.10.0.2 10.10.0.3 10.10.0.4
Remote = 10.10.0.5 10.10.0.6 10.10.0.7 10.10.0.8 10.10.0.9
10.10.0.10 10.10.0.11 10.10.0.12
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 32 path(s)
Basic connectivity fails on 0 path(s)
................
Detected 1500 byte MTU on 32 path(s):
    Local 10.10.0.1 to Remote 10.10.0.5
    Local 10.10.0.1 to Remote 10.10.0.6
    Local 10.10.0.1 to Remote 10.10.0.7
    Local 10.10.0.1 to Remote 10.10.0.8
    Local 10.10.0.1 to Remote 10.10.0.9
    Local 10.10.0.1 to Remote 10.10.0.10
    Local 10.10.0.1 to Remote 10.10.0.11
    Local 10.10.0.1 to Remote 10.10.0.12
    Local 10.10.0.2 to Remote 10.10.0.5
    Local 10.10.0.2 to Remote 10.10.0.6
    Local 10.10.0.2 to Remote 10.10.0.7
    Local 10.10.0.2 to Remote 10.10.0.8
    Local 10.10.0.2 to Remote 10.10.0.9
    Local 10.10.0.2 to Remote 10.10.0.10
    Local 10.10.0.2 to Remote 10.10.0.11
    Local 10.10.0.2 to Remote 10.10.0.12
    Local 10.10.0.3 to Remote 10.10.0.5
    Local 10.10.0.3 to Remote 10.10.0.6
```

```
      Local 10.10.0.3 to Remote 10.10.0.7
      Local 10.10.0.3 to Remote 10.10.0.8
      Local 10.10.0.3 to Remote 10.10.0.9
      Local 10.10.0.3 to Remote 10.10.0.10
      Local 10.10.0.3 to Remote 10.10.0.11
      Local 10.10.0.3 to Remote 10.10.0.12
      Local 10.10.0.4 to Remote 10.10.0.5
      Local 10.10.0.4 to Remote 10.10.0.6
      Local 10.10.0.4 to Remote 10.10.0.7
      Local 10.10.0.4 to Remote 10.10.0.8
      Local 10.10.0.4 to Remote 10.10.0.9
      Local 10.10.0.4 to Remote 10.10.0.10
      Local 10.10.0.4 to Remote 10.10.0.11
      Local 10.10.0.4 to Remote 10.10.0.12

Larger than PMTU communication succeeds on 32 path(s)
RPC status:
8 paths up, 0 paths down (tcp check)
8 paths up, 0 paths down (udp check)
```

Step 3: Verify the configuration

1. Display the information about the devices in your configuration:

   ° `network device-discovery show`

   ° `network port show -role cluster`

   ° `network interface show -role cluster`

   ° `system cluster-switch show`

Show example

```
cluster::> network device-discovery show
          Local  Discovered
Node       Port   Device             Interface       Platform
---------- ------ ------------------ ---------------
----------------
n1         /cdp
          e0a    C1                 Ethernet1/1/1   N3K-C3132Q-V
          e0b    C2                 Ethernet1/1/1   N3K-C3132Q-V
          e0c    C2                 Ethernet1/1/2   N3K-C3132Q-V
          e0d    C1                 Ethernet1/1/2   N3K-C3132Q-V
n2         /cdp
          e0a    C1                 Ethernet1/1/3   N3K-C3132Q-V
          e0b    C2                 Ethernet1/1/3   N3K-C3132Q-V
          e0c    C2                 Ethernet1/1/4   N3K-C3132Q-V
          e0d    C1                 Ethernet1/1/4   N3K-C3132Q-V
n3         /cdp
          e4a    C1                 Ethernet1/7     N3K-C3132Q-V
          e4e    C2                 Ethernet1/7     N3K-C3132Q-V
n4         /cdp
          e4a    C1                 Ethernet1/8     N3K-C3132Q-V
          e4e    C2                 Ethernet1/8     N3K-C3132Q-V
12 entries were displayed.
```

```
cluster::*> network port show -role cluster
  (network port show)
Node: n1

Ignore
                                           Speed(Mbps) Health
Health
Port      IPspace       Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ ---------------- ---- ---- -----------
-------- ------
e0a       Cluster       Cluster          up   9000 auto/10000  -
-
e0b       Cluster       Cluster          up   9000 auto/10000  -
-
e0c       Cluster       Cluster          up   9000 auto/10000  -
-
e0d       Cluster       Cluster          up   9000 auto/10000  -
-
```

```
Node: n2

Ignore
                                                   Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e0a       Cluster      Cluster         up   9000 auto/10000  -
-
e0b       Cluster      Cluster         up   9000 auto/10000  -
-
e0c       Cluster      Cluster         up   9000 auto/10000  -
-
e0d       Cluster      Cluster         up   9000 auto/10000  -
-

Node: n3

Ignore
                                                   Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster         up   9000 auto/40000  -
-
e4e       Cluster      Cluster         up   9000 auto/40000  -
-

Node: n4

Ignore
                                                   Speed(Mbps) Health
Health
Port      IPspace      Broadcast Domain Link MTU  Admin/Oper  Status
Status
--------- ------------ --------------- ---- ---- -----------
-------- ------
e4a       Cluster      Cluster         up   9000 auto/40000  -
-
e4e       Cluster      Cluster         up   9000 auto/40000  -
-
```

12 entries were displayed.

```
cluster::*> network interface show -role cluster
 (network interface show)
          Logical    Status     Network            Current
Current Is
Vserver     Interface  Admin/Oper Address/Mask       Node
Port    Home
----------- ---------- ---------- ------------------ -------------
------- ----
Cluster
          n1_clus1   up/up      10.10.0.1/24       n1
e0a     true
          n1_clus2   up/up      10.10.0.2/24       n1
e0b     true
          n1_clus3   up/up      10.10.0.3/24       n1
e0c     true
          n1_clus4   up/up      10.10.0.4/24       n1
e0d     true
          n2_clus1   up/up      10.10.0.5/24       n2
e0a     true
          n2_clus2   up/up      10.10.0.6/24       n2
e0b     true
          n2_clus3   up/up      10.10.0.7/24       n2
e0c     true
          n2_clus4   up/up      10.10.0.8/24       n2
e0d     true
          n3_clus1   up/up      10.10.0.9/24       n3
e4a     true
          n3_clus2   up/up      10.10.0.10/24      n3
e4e     true
          n4_clus1   up/up      10.10.0.11/24      n4
e4a     true
          n4_clus2   up/up      10.10.0.12/24      n4
e4e     true
12 entries were displayed.
```

```
cluster::*> system cluster-switch show

Switch                          Type              Address
Model
-------------------------- ----------------- ----------------
--------------
CL1                             cluster-network   10.10.1.101
NX3132V
        Serial Number: FOX000001
         Is Monitored: true
               Reason:
     Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                       7.0(3)I4(1)
       Version Source: CDP

CL2                             cluster-network   10.10.1.102
NX3132V
        Serial Number: FOX000002
         Is Monitored: true
               Reason:
     Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                       7.0(3)I4(1)
       Version Source: CDP
C2                              cluster-network   10.10.1.103
NX3132V
        Serial Number: FOX000003
         Is Monitored: true
               Reason:
     Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                       7.0(3)I4(1)
       Version Source: CDP

3 entries were displayed.
```

2. Remove the replaced Nexus 3132Q-V switch, if it is not already removed automatically:

```
system cluster-switch delete
```

```
cluster::*> system cluster-switch delete -device CL2
```

3. Verify that the proper cluster switches are monitored:

```
system cluster-switch show
```

Show example

```
cluster::> system cluster-switch show

Switch                          Type              Address
Model
-------------------------- ----------------- -----------------
--------------
CL1                             cluster-network    10.10.1.101
NX3132V
      Serial Number: FOX000001
        Is Monitored: true
              Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                    7.0(3)I4(1)
    Version Source: CDP

C2                              cluster-network    10.10.1.103
NX3132V
      Serial Number: FOX000002
        Is Monitored: true
              Reason:
   Software Version: Cisco Nexus Operating System (NX-OS) Software,
Version
                    7.0(3)I4(1)
    Version Source: CDP

2 entries were displayed.
```

4. Enable the cluster switch health monitor log collection feature for collecting switch-related log files:

```
system cluster-switch log setup-password
```

```
system cluster-switch log enable-collection
```

```
cluster::*> system cluster-switch log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
C1
C2

cluster::*> system cluster-switch log setup-password

Enter the switch name: C1
RSA key fingerprint is
e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log setup-password

Enter the switch name: C2
RSA key fingerprint is
57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y

Enter the password: <enter switch password>
Enter the password again: <enter switch password>

cluster::*> system cluster-switch log enable-collection

Do you want to enable cluster log collection for all nodes in the
cluster?
{y|n}: [n] y

Enabling cluster switch log collection.

cluster::*>
```

ⓘ   If any of these commands return an error, contact NetApp support.

5. If you suppressed automatic case creation, re-enable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Replace Cisco Nexus 3132Q-V cluster switches with switchless connections

You can migrate from a cluster with a switched cluster network to one where two nodes are directly connected for ONTAP 9.3 and later.

Review requirements

Guidelines

Review the following guidelines:

- Migrating to a two-node switchless cluster configuration is a nondisruptive operation. Most systems have two dedicated cluster interconnect ports on each node, but you can also use this procedure for systems with a larger number of dedicated cluster interconnect ports on each node, such as four, six or eight.
- You cannot use the switchless cluster interconnect feature with more than two nodes.
- If you have an existing two-node cluster that uses cluster interconnect switches and is running ONTAP 9.3 or later, you can replace the switches with direct, back-to-back connections between the nodes.

What you'll need

- A healthy cluster that consists of two nodes connected by cluster switches. The nodes must be running the same ONTAP release.
- Each node with the required number of dedicated cluster ports, which provide redundant cluster interconnect connections to support your system configuration. For example, there are two redundant ports for a system with two dedicated cluster interconnect ports on each node.

Migrate the switches

About this task

The following procedure removes the cluster switches in a two-node cluster and replaces each connection to the switch with a direct connection to the partner node.

About the examples

The examples in the following procedure show nodes that are using "e0a" and "e0b" as cluster ports. Your nodes might be using different cluster ports as they vary by system.

Step 1: Prepare for migration

1. Change the privilege level to advanced, entering `y` when prompted to continue:

   ```
   set -privilege advanced
   ```

   The advanced prompt `*>` appears.

2. ONTAP 9.3 and later supports automatic detection of switchless clusters, which is enabled by default.

   You can verify that detection of switchless clusters is enabled by running the advanced privilege command:

   ```
   network options detect-switchless-cluster show
   ```

   Show example

   > The following example output shows if the option is enabled.
   >
   > ```
   > cluster::*> network options detect-switchless-cluster show
   >     (network options detect-switchless-cluster show)
   > Enable Switchless Cluster Detection: true
   > ```

   If "Enable Switchless Cluster Detection" is `false`, contact NetApp support.

3. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

   ```
   system node autosupport invoke -node * -type all -message
   MAINT=<number_of_hours>h
   ```

   where `h` is the duration of the maintenance window in hours. The message notifies technical support of this maintenance task so that they can suppress automatic case creation during the maintenance window.

   In the following example, the command suppresses automatic case creation for two hours:

   Show example

   > ```
   > cluster::*> system node autosupport invoke -node * -type all
   > -message MAINT=2h
   > ```

Step 2: Configure ports and cabling

1. Organize the cluster ports on each switch into groups so that the cluster ports in group1 go to cluster switch1 and the cluster ports in group2 go to cluster switch2. These groups are required later in the procedure.

2. Identify the cluster ports and verify link status and health:

```
network port show -ipspace Cluster
```

In the following example for nodes with cluster ports "e0a" and "e0b", one group is identified as "node1:e0a" and "node2:e0a" and the other group as "node1:e0b" and "node2:e0b". Your nodes might be using different cluster ports because they vary by system.
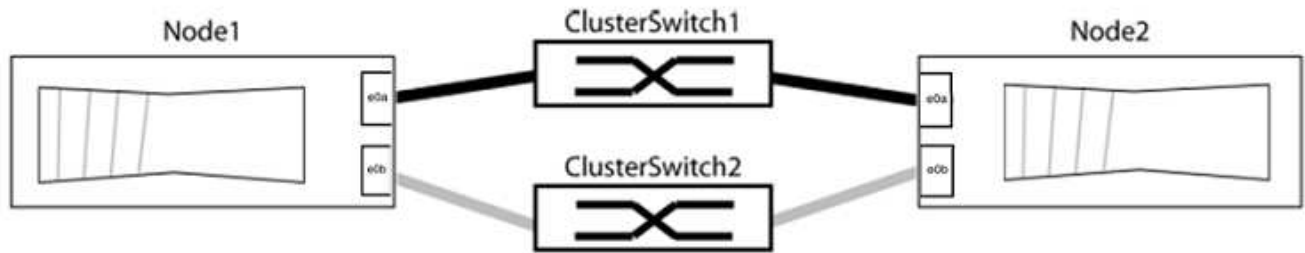
Verify that the ports have a value of `up` for the "Link" column and a value of `healthy` for the "Health Status" column.

Show example

```
cluster::> network port show -ipspace Cluster
Node: node1

Ignore
                                      Speed(Mbps) Health
Health
Port  IPspace    Broadcast Domain Link  MTU   Admin/Oper  Status
Status
----- ---------- ---------------- ----- ----- ----------- -------
-------
e0a   Cluster    Cluster          up    9000  auto/10000  healthy
false
e0b   Cluster    Cluster          up    9000  auto/10000  healthy
false

Node: node2

Ignore
                                      Speed(Mbps) Health
Health
Port  IPspace    Broadcast Domain Link  MTU   Admin/Oper  Status
Status
----- ---------- ---------------- ----- ----- ----------- -------
-------
e0a   Cluster    Cluster          up    9000  auto/10000  healthy
false
e0b   Cluster    Cluster          up    9000  auto/10000  healthy
false
4 entries were displayed.
```

3. Confirm that all the cluster LIFs are on their home ports.

   Verify that the "is-home" column is `true` for each of the cluster LIFs:

   ```
   network interface show -vserver Cluster -fields is-home
   ```

Show example

```
cluster::*> net int show -vserver Cluster -fields is-home
  (network interface show)
vserver   lif          is-home
--------  ------------ --------
Cluster   node1_clus1  true
Cluster   node1_clus2  true
Cluster   node2_clus1  true
Cluster   node2_clus2  true
4 entries were displayed.
```

If there are cluster LIFs that are not on their home ports, revert those LIFs to their home ports:

```
network interface revert -vserver Cluster -lif *
```

4. Disable auto-revert for the cluster LIFs:

```
network interface modify -vserver Cluster -lif * -auto-revert false
```

5. Verify that all ports listed in the previous step are connected to a network switch:

```
network device-discovery show -port cluster_port
```

The "Discovered Device" column should be the name of the cluster switch that the port is connected to.

Show example

The following example shows that cluster ports "e0a" and "e0b" are correctly connected to cluster switches "cs1" and "cs2".

```
cluster::> network device-discovery show -port e0a|e0b
  (network device-discovery show)
Node/     Local  Discovered
Protocol  Port   Device (LLDP: ChassisID)  Interface  Platform
--------- ------ ------------------------- ---------- ----------
node1/cdp
          e0a    cs1                        0/11       BES-53248
          e0b    cs2                        0/12       BES-53248
node2/cdp
          e0a    cs1                        0/9        BES-53248
          e0b    cs2                        0/9        BES-53248
4 entries were displayed.
```

6. Verify the cluster connectivity:

```
cluster ping-cluster -node local
```

7. Verify that the cluster is healthy:

```
cluster ring show
```

All units must be either master or secondary.

8. Set up the switchless configuration for the ports in group 1.

> ⓘ To avoid potential networking issues, you must disconnect the ports from group1 and
> reconnect them back-to-back as quickly as possible, for example, **in less than 20 seconds**.

a. Disconnect all the cables from the ports in group1 at the same time.

In the following example, the cables are disconnected from port "e0a" on each node, and cluster traffic
continues through the switch and port "e0b" on each node:
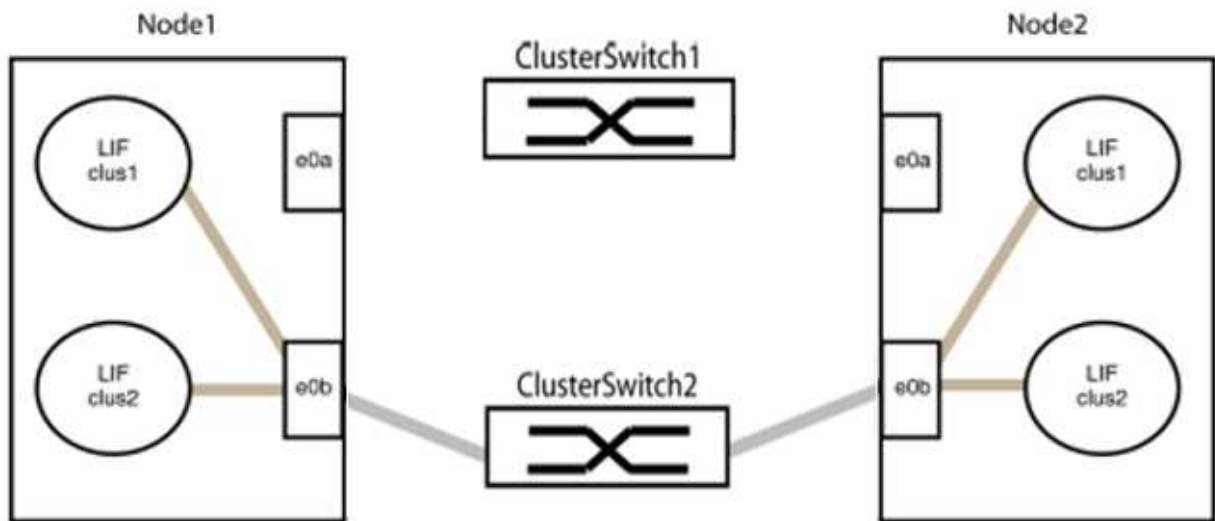
b. Cable the ports in group1 back-to-back.

In the following example, "e0a" on node1 is connected to "e0a" on node2:

9. The switchless cluster network option transitions from `false` to `true`. This might take up to 45 seconds. Confirm that the switchless option is set to `true`:

```
network options switchless-cluster show
```

The following example shows that the switchless cluster is enabled:

```
cluster::*> network options switchless-cluster show
Enable Switchless Cluster: true
```

10. Verify that the cluster network is not disrupted:

```
cluster ping-cluster -node local
```

> ⓘ Before proceeding to the next step, you must wait at least two minutes to confirm a working back-to-back connection on group 1.

11. Set up the switchless configuration for the ports in group 2.

> ⓘ To avoid potential networking issues, you must disconnect the ports from group2 and reconnect them back-to-back as quickly as possible, for example, **in less than 20 seconds**.

a. Disconnect all the cables from the ports in group2 at the same time.

In the following example, the cables are disconnected from port "e0b" on each node, and cluster traffic continues through the direct connection between the "e0a" ports:

b. Cable the ports in group2 back-to-back.

In the following example, "e0a" on node1 is connected to "e0a" on node2 and "e0b" on node1 is connected to "e0b" on node2:

Step 3: Verify the configuration

1. Verify that the ports on both nodes are correctly connected:

   `network device-discovery show -port` *`cluster_port`*

Show example

The following example shows that cluster ports "e0a" and "e0b" are correctly connected to the corresponding port on the cluster partner:

```
cluster::> net device-discovery show -port e0a|e0b
  (network device-discovery show)
Node/      Local  Discovered
Protocol   Port   Device (LLDP: ChassisID)  Interface  Platform
---------- ------ ------------------------- ---------- ----------
node1/cdp
           e0a    node2                     e0a        AFF-A300
           e0b    node2                     e0b        AFF-A300
node1/lldp
           e0a    node2 (00:a0:98:da:16:44) e0a        -
           e0b    node2 (00:a0:98:da:16:44) e0b        -
node2/cdp
           e0a    node1                     e0a        AFF-A300
           e0b    node1                     e0b        AFF-A300
node2/lldp
           e0a    node1 (00:a0:98:da:87:49) e0a        -
           e0b    node1 (00:a0:98:da:87:49) e0b        -
8 entries were displayed.
```

2. Re-enable auto-revert for the cluster LIFs:

```
network interface modify -vserver Cluster -lif * -auto-revert true
```

3. Verify that all LIFs are home. This might take a few seconds.

```
network interface show -vserver Cluster -lif lif_name
```

Show example

> The LIFs have been reverted if the "Is Home" column is `true`, as shown for `node1_clus2` and `node2_clus2` in the following example:
>
> ```
> cluster::> network interface show -vserver Cluster -fields curr-
> port,is-home
> vserver  lif           curr-port is-home
> -------- ------------- --------- -------
> Cluster  node1_clus1   e0a       true
> Cluster  node1_clus2   e0b       true
> Cluster  node2_clus1   e0a       true
> Cluster  node2_clus2   e0b       true
> 4 entries were displayed.
> ```

If any cluster LIFS have not returned to their home ports, revert them manually from the local node:

```
network interface revert -vserver Cluster -lif lif_name
```

4. Check the cluster status of the nodes from the system console of either node:

```
cluster show
```

Show example

> The following example shows epsilon on both nodes to be `false`:
>
> ```
> Node   Health  Eligibility Epsilon
> -----  ------- ----------- --------
> node1 true     true        false
> node2 true     true        false
> 2 entries were displayed.
> ```

5. Confirm connectivity between the cluster ports:

```
cluster ping-cluster local
```

6. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

For more information, see [NetApp KB Article 1010449: How to suppress automatic case creation during scheduled maintenance windows](#).

7. Change the privilege level back to admin:

```
set -privilege admin
```