



Migrate switches

Cluster and storage switches

NetApp
May 08, 2024

Table of Contents

Migrate switches. 1

 Migrate from a switchless cluster environment to a switched NetApp CN1610 cluster environment 1

Migrate switches

Migrate from a switchless cluster environment to a switched NetApp CN1610 cluster environment

If you have an existing two-node switchless cluster environment, you can migrate to a two-node switched cluster environment using CN1610 cluster network switches that enables you to scale beyond two nodes.

Review requirements

What you'll need

For a two-node switchless configuration, ensure that:

- The two-node switchless configuration is properly set up and functioning.
- The nodes are running ONTAP 8.2 or later.
- All cluster ports are in the `up` state.
- All cluster logical interfaces (LIFs) are in the `up` state and on their home ports.

For the CN1610 cluster switch configuration:

- The CN1610 cluster switch infrastructure are fully functional on both switches.
- Both switches have management network connectivity.
- There is console access to the cluster switches.
- CN1610 node-to-node switch and switch-to-switch connections use twinax or fiber cables.

The [Hardware Universe](#) contains more information about cabling.

- Inter-Switch Link (ISL) cables are connected to ports 13 through 16 on both CN1610 switches.
- Initial customization of both the CN1610 switches are completed.

Any previous site customization, such as SMTP, SNMP, and SSH should be copied to the new switches.

Related information

- [Hardware Universe](#)
- [NetApp CN1601 and CN1610 description page](#)
- [CN1601 and CN1610 Switch Setup and Configuration Guide](#)
- [NetApp KB Article 1010449: How to suppress automatic case creation during scheduled maintenance windows](#)

Migrate the switches

About the examples

The examples in this procedure use the following cluster switch and node nomenclature:

- The names of the CN1610 switches are cs1 and cs2.
- The names of the LIFs are clus1 and clus2.
- The names of the nodes are node1 and node2.
- The `cluster::*>` prompt indicates the name of the cluster.
- The cluster ports used in this procedure are e1a and e2a.

The [Hardware Universe](#) contains the latest information about the actual cluster ports for your platforms.

Step 1: Prepare for migration

1. Change the privilege level to advanced, entering `y` when prompted to continue:

```
set -privilege advanced
```

The advanced prompt (`*>`) appears.

2. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=xh
```

`x` is the duration of the maintenance window in hours.



The AutoSupport message notifies technical support of this maintenance task so that automatic case creation is suppressed during the maintenance window.

Show example

The following command suppresses automatic case creation for two hours:

```
cluster::*> system node autosupport invoke -node * -type all
-message MAINT=2h
```

Step 2: Configure ports

1. Disable all of the node-facing ports (not ISL ports) on both the new cluster switches cs1 and cs2.

You must not disable the ISL ports.

Show example

The following example shows that node-facing ports 1 through 12 are disabled on switch cs1:

```
(cs1)> enable
(cs1)# configure
(cs1)(Config)# interface 0/1-0/12
(cs1)(Interface 0/1-0/12)# shutdown
(cs1)(Interface 0/1-0/12)# exit
(cs1)(Config)# exit
```

The following example shows that node-facing ports 1 through 12 are disabled on switch cs2:

```
(c2)> enable
(cs2)# configure
(cs2)(Config)# interface 0/1-0/12
(cs2)(Interface 0/1-0/12)# shutdown
(cs2)(Interface 0/1-0/12)# exit
(cs2)(Config)# exit
```

2. Verify that the ISL and the physical ports on the ISL between the two CN1610 cluster switches cs1 and cs2 are up:

```
show port-channel
```

Show example

The following example shows that the ISL ports are up on switch cs1:

```
(cs1)# show port-channel 3/1
Local Interface..... 3/1
Channel Name..... ISL-LAG
Link State..... Up
Admin Mode..... Enabled
Type..... Static
Load Balance Option..... 7
(Enhanced hashing mode)
```

Mbr Ports	Device/ Timeout	Port Speed	Port Active
-----	-----	-----	-----
0/13	actor/long partner/long	10G Full	True
0/14	actor/long partner/long	10G Full	True
0/15	actor/long partner/long	10G Full	True
0/16	actor/long partner/long	10G Full	True

The following example shows that the ISL ports are up on switch cs2:

```
(cs2)# show port-channel 3/1
Local Interface..... 3/1
Channel Name..... ISL-LAG
Link State..... Up
Admin Mode..... Enabled
Type..... Static
Load Balance Option..... 7
(Enhanced hashing mode)
```

Mbr	Device/ Ports	Port Timeout	Port Speed	Port Active
-----	-----	-----	-----	-----
0/13	actor/long partner/long	10G Full	True	
0/14	actor/long partner/long	10G Full	True	
0/15	actor/long partner/long	10G Full	True	
0/16	actor/long partner/long	10G Full	True	

3. Display the list of neighboring devices:

```
show isdp neighbors
```

This command provides information about the devices that are connected to the system.

Show example

The following example lists the neighboring devices on switch cs1:

```
(cs1)# show isdp neighbors
Capability Codes: R - Router, T - Trans Bridge, B - Source Route
Bridge,
                S - Switch, H - Host, I - IGMP, r - Repeater
Device ID      Intf      Holdtime  Capability  Platform
Port ID
-----
cs2            0/13      11        S           CN1610
0/13
cs2            0/14      11        S           CN1610
0/14
cs2            0/15      11        S           CN1610
0/15
cs2            0/16      11        S           CN1610
0/16
```

The following example lists the neighboring devices on switch cs2:

```
(cs2)# show isdp neighbors
Capability Codes: R - Router, T - Trans Bridge, B - Source Route
Bridge,
                S - Switch, H - Host, I - IGMP, r - Repeater
Device ID      Intf      Holdtime  Capability  Platform
Port ID
-----
cs1            0/13      11        S           CN1610
0/13
cs1            0/14      11        S           CN1610
0/14
cs1            0/15      11        S           CN1610
0/15
cs1            0/16      11        S           CN1610
0/16
```

4. Display the list of cluster ports:

```
network port show
```


Show example

The following example shows the available cluster ports:

```
cluster::~*> network port show -ipspace Cluster
```

```
Node: node1
```

```
Ignore
```

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							
-----	-----	-----	-----	----	----	-----	
-----	-----						
e0a	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0b	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0c	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0d	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e4a	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e4b	Cluster	Cluster		up	9000	auto/10000	
healthy	false						

```
Node: node2
```

```
Ignore
```

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							
-----	-----	-----	-----	----	----	-----	
-----	-----						
e0a	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0b	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0c	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e0d	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e4a	Cluster	Cluster		up	9000	auto/10000	
healthy	false						
e4b	Cluster	Cluster		up	9000	auto/10000	
healthy	false						

```
12 entries were displayed.
```

5. Verify that each cluster port is connected to the corresponding port on its partner cluster node:

```
run * cdpd show-neighbors
```

Show example

The following example shows that cluster ports e1a and e2a are connected to the same port on their cluster partner node:

```
cluster::*> run * cdpd show-neighbors
2 entries were acted on.
```

Node: node1

Local Remote	Remote	Remote	Remote	Hold
Port Device	Interface	Platform	Time	
Capability				

e1a	node2	e1a	FAS3270	137
H				
e2a	node2	e2a	FAS3270	137
H				

Node: node2

Local Remote	Remote	Remote	Remote	Hold
Port Device	Interface	Platform	Time	
Capability				

e1a	node1	e1a	FAS3270	161
H				
e2a	node1	e2a	FAS3270	161
H				

6. Verify that all of the cluster LIFs are up and operational:

```
network interface show -vserver Cluster
```

Each cluster LIF should display `true` in the “Is Home” column.

Show example

```
cluster::*> network interface show -vserver Cluster
```

	Logical	Status	Network	Current	
Current Is					
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port
Home					
-----	-----	-----	-----	-----	-----
node1					
true	clus1	up/up	10.10.10.1/16	node1	e1a
true	clus2	up/up	10.10.10.2/16	node1	e2a
node2					
true	clus1	up/up	10.10.11.1/16	node2	e1a
true	clus2	up/up	10.10.11.2/16	node2	e2a

4 entries were displayed.



The following modification and migration commands in steps 10 through 13 must be done from the local node.

7. Verify that all cluster ports are up:

```
network port show -ipSPACE Cluster
```

Show example

```
cluster::*> network port show -ipspace Cluster
```

					Auto-Negot	Duplex	Speed
(Mbps)							
Node	Port	Role	Link	MTU	Admin/Oper	Admin/Oper	
Admin/Oper							
-----	-----	-----	-----	-----	-----	-----	
node1							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							
node2							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							

4 entries were displayed.

8. Set the `-auto-revert` parameter to `false` on cluster LIFs `clus1` and `clus2` on both nodes:

```
network interface modify
```

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node1 -lif clus2 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert false
cluster::*> network interface modify -vserver node2 -lif clus2 -auto
-revert false
```



For release 8.3 and later, use the following command: `network interface modify -vserver Cluster -lif * -auto-revert false`

9. Ping the cluster ports to verify the cluster connectivity:

```
cluster ping-cluster local
```

The command output shows connectivity between all of the cluster ports.

10. Migrate clus1 to port e2a on the console of each node:

```
network interface migrate
```

Show example

The following example shows the process for migrating clus1 to port e2a on node1 and node2:

```
cluster::*> network interface migrate -vserver node1 -lif clus1  
-source-node node1 -dest-node node1 -dest-port e2a  
cluster::*> network interface migrate -vserver node2 -lif clus1  
-source-node node2 -dest-node node2 -dest-port e2a
```



For release 8.3 and later, use the following command: `network interface migrate -vserver Cluster -lif clus1 -destination-node node1 -destination-port e2a`

11. Verify that the migration took place:

```
network interface show -vserver Cluster
```

Show example

The following example verifies that clus1 is migrated to port e2a on node1 and node2:

```
cluster::*> network interface show -vserver Cluster
```

	Logical	Status	Network	Current	
Current Is					
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port
Home					
-----	-----	-----	-----	-----	-----
node1					
false	clus1	up/up	10.10.10.1/16	node1	e2a
true	clus2	up/up	10.10.10.2/16	node1	e2a
node2					
false	clus1	up/up	10.10.11.1/16	node2	e2a
true	clus2	up/up	10.10.11.2/16	node2	e2a

4 entries were displayed.

12. Shut down cluster port e1a on both nodes:

```
network port modify
```

Show example

The following example shows how to shut down the port e1a on node1 and node2:

```
cluster::*> network port modify -node node1 -port e1a -up-admin  
false  
cluster::*> network port modify -node node2 -port e1a -up-admin  
false
```

13. Verify the port status:

```
network port show
```

Show example

The following example shows that port e1a is down on node1 and node2:

```
cluster::*> network port show -role cluster
```

					Auto-Negot	Duplex	Speed
(Mbps)							
Node	Port	Role	Link	MTU	Admin/Oper	Admin/Oper	
Admin/Oper							
-----	-----	-----	----	-----	-----	-----	-----

node1							
	e1a	clus1	down	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							
node2							
	e1a	clus1	down	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							

4 entries were displayed.

14. Disconnect the cable from cluster port e1a on node1, and then connect e1a to port 1 on cluster switch cs1, using the appropriate cabling supported by the CN1610 switches.

The [Hardware Universe](#) contains more information about cabling.

15. Disconnect the cable from cluster port e1a on node2, and then connect e1a to port 2 on cluster switch cs1, using the appropriate cabling supported by the CN1610 switches.
16. Enable all of the node-facing ports on cluster switch cs1.

Show example

The following example shows that ports 1 through 12 are enabled on switch cs1:

```
(cs1)# configure
(cs1)(Config)# interface 0/1-0/12
(cs1)(Interface 0/1-0/12)# no shutdown
(cs1)(Interface 0/1-0/12)# exit
(cs1)(Config)# exit
```

17. Enable the first cluster port e1a on each node:


```
network port modify
```

Show example

The following example shows how to enable the port e1a on node1 and node2:

```
cluster::*> network port modify -node node1 -port e1a -up-admin true
cluster::*> network port modify -node node2 -port e1a -up-admin true
```

18. Verify that all of the cluster ports are up:

```
network port show -ipspace Cluster
```

Show example

The following example shows that all of the cluster ports are up on node1 and node2:

```
cluster::*> network port show -ipspace Cluster
```

					Auto-Negot	Duplex	Speed
(Mbps)							
Node	Port	Role	Link	MTU	Admin/Oper	Admin/Oper	
Admin/Oper							

node1							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							
node2							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							

4 entries were displayed.

19. Revert clus1 (which was previously migrated) to e1a on both nodes:

```
network interface revert
```

Show example

The following example shows how to revert clus1 to the port e1a on node1 and node2:

```
cluster::*> network interface revert -vserver node1 -lif clus1
cluster::*> network interface revert -vserver node2 -lif clus1
```



For release 8.3 and later, use the following command: `network interface revert -vserver Cluster -lif <nodename_clus<N>>`

20. Verify that all of the cluster LIFs are up, operational, and display as true in the "Is Home" column:

```
network interface show -vserver Cluster
```

Show example

The following example shows that all of the LIFs are up on node1 and node2 and that the "Is Home" column results are true:

```
cluster::*> network interface show -vserver Cluster
      Logical      Status      Network      Current
Current Is
Vserver      Interface  Admin/Oper Address/Mask  Node        Port
Home
-----
node1
      clus1      up/up      10.10.10.1/16  node1        e1a
true
      clus2      up/up      10.10.10.2/16  node1        e2a
true
node2
      clus1      up/up      10.10.11.1/16  node2        e1a
true
      clus2      up/up      10.10.11.2/16  node2        e2a
true

4 entries were displayed.
```

21. Display information about the status of the nodes in the cluster:

```
cluster show
```

Show example

The following example displays information about the health and eligibility of the nodes in the cluster:

```
cluster::*> cluster show
Node           Health Eligibility  Epsilon
-----
node1          true   true       false
node2          true   true       false
```

22. Migrate clus2 to port e1a on the console of each node:

```
network interface migrate
```

Show example

The following example shows the process for migrating clus2 to port e1a on node1 and node2:

```
cluster::*> network interface migrate -vserver node1 -lif clus2
-source-node node1 -dest-node node1 -dest-port e1a
cluster::*> network interface migrate -vserver node2 -lif clus2
-source-node node2 -dest-node node2 -dest-port e1a
```



For release 8.3 and later, use the following command: `network interface migrate -vserver Cluster -lif node1_clus2 -dest-node node1 -dest-port e1a`

23. Verify that the migration took place:

```
network interface show -vserver Cluster
```

Show example

The following example verifies that clus2 is migrated to port e1a on node1 and node2:

```
cluster::*> network interface show -vserver Cluster
```

	Logical	Status	Network	Current	
Current Is					
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port
Home					
-----	-----	-----	-----	-----	-----
node1					
clus1	up/up	10.10.10.1/16	node1	e1a	
clus2	up/up	10.10.10.2/16	node1	e1a	
node2					
clus1	up/up	10.10.11.1/16	node2	e1a	
clus2	up/up	10.10.11.2/16	node2	e1a	

4 entries were displayed.

24. Shut down cluster port e2a on both nodes:

```
network port modify
```

Show example

The following example shows how to shut down the port e2a on node1 and node2:

```
cluster::*> network port modify -node node1 -port e2a -up-admin  
false  
cluster::*> network port modify -node node2 -port e2a -up-admin  
false
```

25. Verify the port status:

```
network port show
```

Show example

The following example shows that port e2a is down on node1 and node2:

```
cluster::*> network port show -role cluster
```

					Auto-Negot	Duplex	Speed
(Mbps)							
Node	Port	Role	Link	MTU	Admin/Oper	Admin/Oper	
Admin/Oper							
-----	-----	-----	----	-----	-----	-----	-----

node1							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	down	9000	true/true	full/full	
auto/10000							
node2							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	down	9000	true/true	full/full	
auto/10000							

4 entries were displayed.

26. Disconnect the cable from cluster port e2a on node1, and then connect e2a to port 1 on cluster switch cs2, using the appropriate cabling supported by the CN1610 switches.
27. Disconnect the cable from cluster port e2a on node2, and then connect e2a to port 2 on cluster switch cs2, using the appropriate cabling supported by the CN1610 switches.
28. Enable all of the node-facing ports on cluster switch cs2.

Show example

The following example shows that ports 1 through 12 are enabled on switch cs2:

```
(cs2)# configure
(cs2)(Config)# interface 0/1-0/12
(cs2)(Interface 0/1-0/12)# no shutdown
(cs2)(Interface 0/1-0/12)# exit
(cs2)(Config)# exit
```

29. Enable the second cluster port e2a on each node.

Show example

The following example shows how to enable the port e2a on node1 and node2:

```
cluster::*> network port modify -node node1 -port e2a -up-admin true
cluster::*> network port modify -node node2 -port e2a -up-admin true
```

30. Verify that all of the cluster ports are up:

```
network port show -ipspace Cluster
```

Show example

The following example shows that all of the cluster ports are up on node1 and node2:

```
cluster::*> network port show -ipspace Cluster
```

					Auto-Negot	Duplex	Speed
(Mbps)							
Node	Port	Role	Link	MTU	Admin/Oper	Admin/Oper	
Admin/Oper							
-----	-----	-----	----	-----	-----	-----	

node1							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							
node2							
	e1a	clus1	up	9000	true/true	full/full	
auto/10000							
	e2a	clus2	up	9000	true/true	full/full	
auto/10000							

4 entries were displayed.

31. Revert clus2 (which was previously migrated) to e2a on both nodes:

```
network interface revert
```

Show example

The following example shows how to revert clus2 to the port e2a on node1 and node2:

```
cluster::*> network interface revert -vserver node1 -lif clus2
cluster::*> network interface revert -vserver node2 -lif clus2
```



For release 8.3 and later, the commands are: `cluster::*> network interface revert -vserver Cluster -lif node1_clus2` and `cluster::*> network interface revert -vserver Cluster -lif node2_clus2`

Step 3: Complete the configuration

1. Verify that all of the interfaces display `true` in the "Is Home" column:

```
network interface show -vserver Cluster
```

Show example

The following example shows that all of the LIFs are up on node1 and node2 and that the "Is Home" column results are true:

```
cluster::*> network interface show -vserver Cluster
```

Current Vserver Port	Is Home	Logical Interface	Status Admin/Oper	Network Address/Mask	Current Node

node1					
e1a	true	clus1	up/up	10.10.10.1/16	node1
e2a	true	clus2	up/up	10.10.10.2/16	node1
node2					
e1a	true	clus1	up/up	10.10.11.1/16	node2
e2a	true	clus2	up/up	10.10.11.2/16	node2

2. Ping the cluster ports to verify the cluster connectivity:

```
cluster ping-cluster local
```

The command output shows connectivity between all of the cluster ports.

3. Verify that both nodes have two connections to each switch:

```
show isdp neighbors
```


Show example

The following example shows the appropriate results for both switches:

```
(cs1)# show isdp neighbors
Capability Codes: R - Router, T - Trans Bridge, B - Source Route
Bridge,
                S - Switch, H - Host, I - IGMP, r - Repeater
Device ID      Intf      Holdtime  Capability  Platform
Port ID
-----
node1          0/1        132       H           FAS3270
e1a
node2          0/2        163       H           FAS3270
e1a
cs2            0/13       11        S           CN1610
0/13
cs2            0/14       11        S           CN1610
0/14
cs2            0/15       11        S           CN1610
0/15
cs2            0/16       11        S           CN1610
0/16

(cs2)# show isdp neighbors
Capability Codes: R - Router, T - Trans Bridge, B - Source Route
Bridge,
                S - Switch, H - Host, I - IGMP, r - Repeater
Device ID      Intf      Holdtime  Capability  Platform
Port ID
-----
node1          0/1        132       H           FAS3270
e2a
node2          0/2        163       H           FAS3270
e2a
cs1            0/13       11        S           CN1610
0/13
cs1            0/14       11        S           CN1610
0/14
cs1            0/15       11        S           CN1610
0/15
cs1            0/16       11        S           CN1610
0/16
```

4. Display information about the devices in your configuration:

```
network device discovery show
```

5. Disable the two-node switchless configuration settings on both nodes using the advanced privilege command:

```
network options detect-switchless modify
```

Show example

The following example shows how to disable the switchless configuration settings:

```
cluster::*> network options detect-switchless modify -enabled false
```



For release 9.2 and later, skip this step since the configuration is automatically converted.

6. Verify that the settings are disabled:

```
network options detect-switchless-cluster show
```

Show example

The false output in the following example shows that the configuration settings are disabled:

```
cluster::*> network options detect-switchless-cluster show
Enable Switchless Cluster Detection: false
```



For release 9.2 and later, wait until Enable Switchless Cluster is set to false. This can take up to three minutes.

7. Configure clusters clus1 and clus2 to auto revert on each node and confirm.

Show example

```
cluster::*> network interface modify -vserver node1 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node1 -lif clus2 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus1 -auto
-revert true
cluster::*> network interface modify -vserver node2 -lif clus2 -auto
-revert true
```



For release 8.3 and later, use the following command: `network interface modify -vserver Cluster -lif * -auto-revert true` to enable auto-revert on all nodes in the cluster.

8. Verify the status of the node members in the cluster:

```
cluster show
```

Show example

The following example shows information about the health and eligibility of the nodes in the cluster:

```
cluster::*> cluster show
Node                Health  Eligibility  Epsilon
-----
node1                true    true         false
node2                true    true         false
```

9. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```

Show example

```
cluster::*> system node autosupport invoke -node * -type all
-mmessage MAINT=END
```

10. Change the privilege level back to admin:

```
set -privilege admin
```

Copyright information

Copyright © 2024 NetApp, Inc. All Rights Reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

LIMITED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (b)(3) of the Rights in Technical Data -Noncommercial Items at DFARS 252.227-7013 (FEB 2014) and FAR 52.227-19 (DEC 2007).

Data contained herein pertains to a commercial product and/or commercial service (as defined in FAR 2.101) and is proprietary to NetApp, Inc. All NetApp technical data and computer software provided under this Agreement is commercial in nature and developed solely at private expense. The U.S. Government has a non-exclusive, non-transferrable, nonsublicensable, worldwide, limited irrevocable license to use the Data only in connection with and in support of the U.S. Government contract under which the Data was delivered. Except as provided herein, the Data may not be used, disclosed, reproduced, modified, performed, or displayed without the prior written approval of NetApp, Inc. United States Government license rights for the Department of Defense are limited to those rights identified in DFARS clause 252.227-7015(b) (FEB 2014).

Trademark information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.